



Computer Systems

Exercise 9



Last Exercise

1.1 Segmentation

Consider the following segment table:

Segment	Base	Length
0	219	600
1	2300	14
2	90	100
3	1327	580
4	1952	96

What are the physical addresses for the following logical addresses?

a) 0,430

1.) $430 < 600$? ✓

b) 1,10

2.) $219 + 430 = 649$

c) 2,500

d) 3,400

e) 4,112



Last Exercise

1.1 Segmentation

Consider the following segment table:

Segment	Base	Length
0	219	600
1	2300	14
2	90	100
3	1327	580
4	1952	96

What are the physical addresses for the following logical addresses?

a) 0,430

1.) $500 < 100$? 

b) 1,10

c) 2,500

d) 3,400

e) 4,112



Last Exercise

1.2 Paging

Consider a paging system with the page table stored in memory.

- a) If a memory reference takes 200 nanoseconds, how long does a paged memory reference take if there is no TLB?
- b) If we add a TLB and 75% of all page-table references are found in the TLB, what is the average memory reference time when a TLB access takes 2 nanosecond?
- c) A typical program has 20% memory instructions. Assume there are 5% data TLB misses, each requiring 100 cycles to handle. Furthermore each instruction requires 1 cycle to execute and each memory operation in the cache takes 1 cycle. Also, 10% of the data accesses are cache misses each of which takes 15 cycles. How long would it take to execute 1000 instructions?

a) No TLB: one access for page, one access for page table: $200 + 200 = 400$
b) $(2+200)*0.75 + (2+200+200)*0.25 = 252\text{ns}$
c) Idea: calculate the overhead


```
# cache misses = 1000 * 20% * 10% = 20
# TLB misses = 1000 * 20% * 5% = 10

total cycles = 1000 cycles + (cache miss overhead) + (TLB miss overhead)
              = 1000 + (20 * 14) + (10 * 99) = 1000 + 280 + 990 = 2270 cycles
```



Last exercise

1.3 Virtual Memory

Consider a paged virtual address space composed of 32 pages of 2 KB each which is mapped into a 1 MB physical memory space.

- a) What is the format of the logical address; i.e., which bits are the offset bits and which are the page number bits? Explain.
 - a) offset: each byte needs to get accessed
 - $2\text{KB} = 2^{11}$ bytes -> **11 offset bits**
 - page number: each page needs to get accessed
 - $32 \text{ pages} = 2^5 \text{ pages}$ -> **5 page number bits**
- b) length of page table: for each page one entry
 - **32 entries**
- width of page table: each entry contains one physical address
 - size of physical address: $1\text{MB} = 2^{20}$ bytes. So **20 bits** to represent address



Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

- a) LRU replacement
- b) FIFO replacement
- c) Optimal replacement

✗

1



Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, **2**, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

- a) LRU replacement
- b) FIFO replacement
- c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

- a) LRU replacement
- b) FIFO replacement
- c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, **4**, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

- a) LRU replacement
- b) FIFO replacement
- c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, **2**, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

- a) LRU replacement
- b) FIFO replacement
- c) Optimal replacement





Last exercise

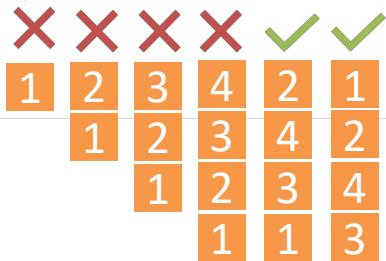
1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

- a) LRU replacement
- b) FIFO replacement
- c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

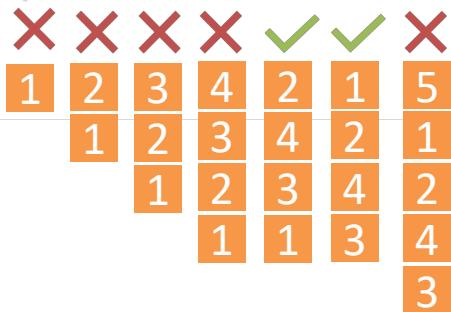
1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

a) LRU replacement

b) FIFO replacement

c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

- a) LRU replacement
 - b) FIFO replacement
 - c) Optimal replacement
- 1



Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

- a) LRU replacement
- b) FIFO replacement
- c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

a) LRU replacement

b) FIFO replacement

c) Optimal replacement

✗ ✗ ✗

1	2	3
1	2	
1		



Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

a) LRU replacement

b) FIFO replacement

c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, **2**, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

a) LRU replacement

b) FIFO replacement

c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

a) LRU replacement

b) FIFO replacement

c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

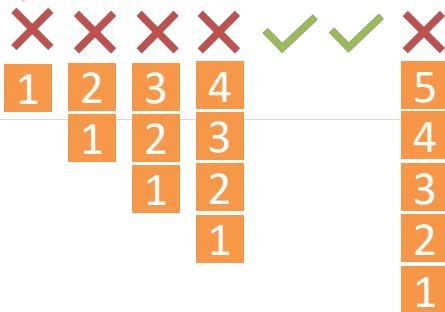
1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

a) LRU replacement

b) FIFO replacement

c) Optimal replacement





Last exercise

1.4 Page Replacement

Consider the following page access pattern:

1, 2, 3, 4, 2, 1, 5, 6, 2, 1, 2, 3, 7, 6, 3, 2, 1, 2, 3, 6

How many page faults would occur for the following replacement algorithms, assuming one, two, three, four, five, six, or seven frames?

a) LRU replacement

b) FIFO replacement

c) Optimal replacement

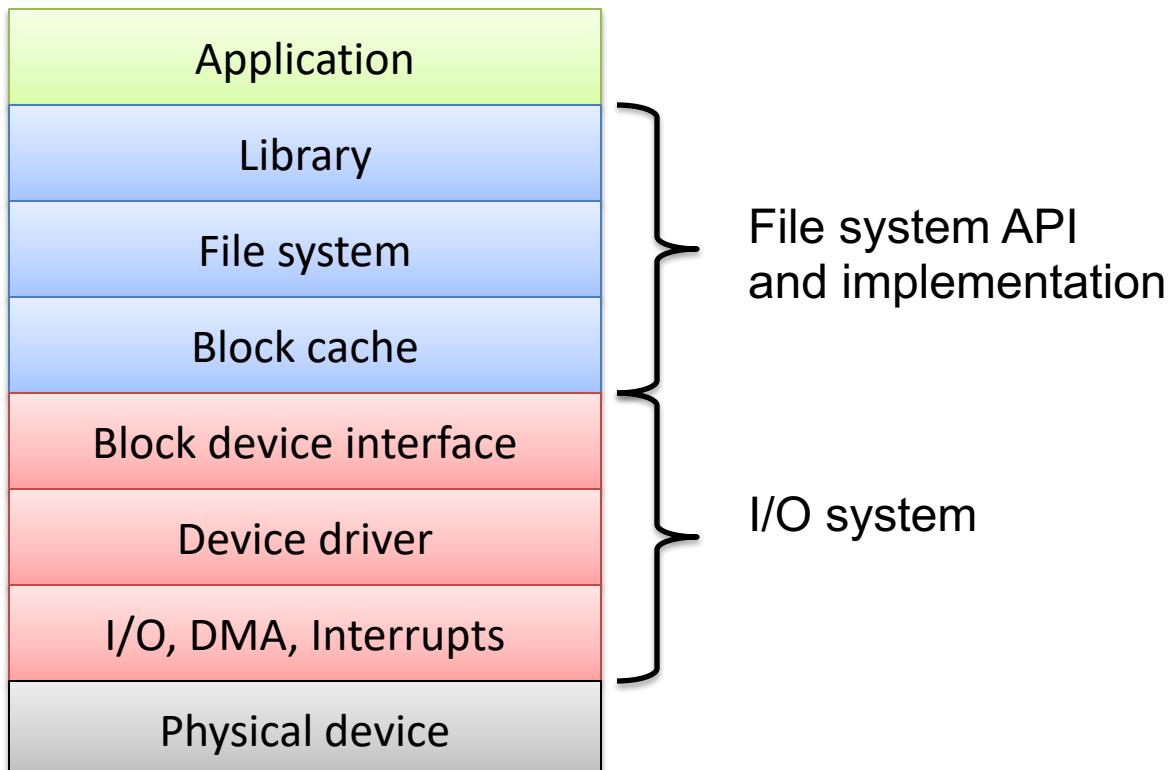




File System Abstractions



File system in context





File systems operations

- Files
 - Read
 - Write
 - Execute
- Directory
 - Link
 - Unlink
 - Rename
 - List



File metadata

- Metadata: important concept!
 - Data *about* an object, not the object *itself*
- File metadata examples:
 - Name
 - Location on disk
 - Times of creation, last change, last access
 - Ownership, access control rights File type, file structure
 - Arbitrary descriptive data (used for searching)



Protection

- File owner/creator should be able to control:
 - what can be done
 - by whom
- Types of access
 - **Read**
 - **Write**
 - **Execute**
 - **Append**
 - **Delete**
 - **List (“read” for directory)**



Access control matrix

For a single file or directory:

	Objects								
	File1	File2	File3	File4	File5	File6	File7	...	
A	rwx	r	x				rw	x	
B	r	rwx			x			x	
C	rw					rw			
D	rw	rw	x	x					
E	r								
F	r					rw			
...									

Problem: how to scalably represent this matrix?



ACLs

- Access Control Lists
 - For each file, list the principals and their rights on the object
 - Store with the file
- Good:
 - Easy to change rights quickly
 - Scales to large numbers of files
- Bad:
 - Doesn't scale to large numbers of principals



File System Implementation



Implementation aspects

- Directories and indexes
 - Where on the disk is the data for each file?
- Index granularity
 - What is the unit of allocation for files?
- Free space maps
 - How to allocate more sectors on the disk?
- Locality optimizations
 - How to make it go fast in the common case



FAT

FAT background

- No access control
- Very little metadata
- Limited volume size
- BUT still extensively used ☹
 - Flash devices, cameras, phones



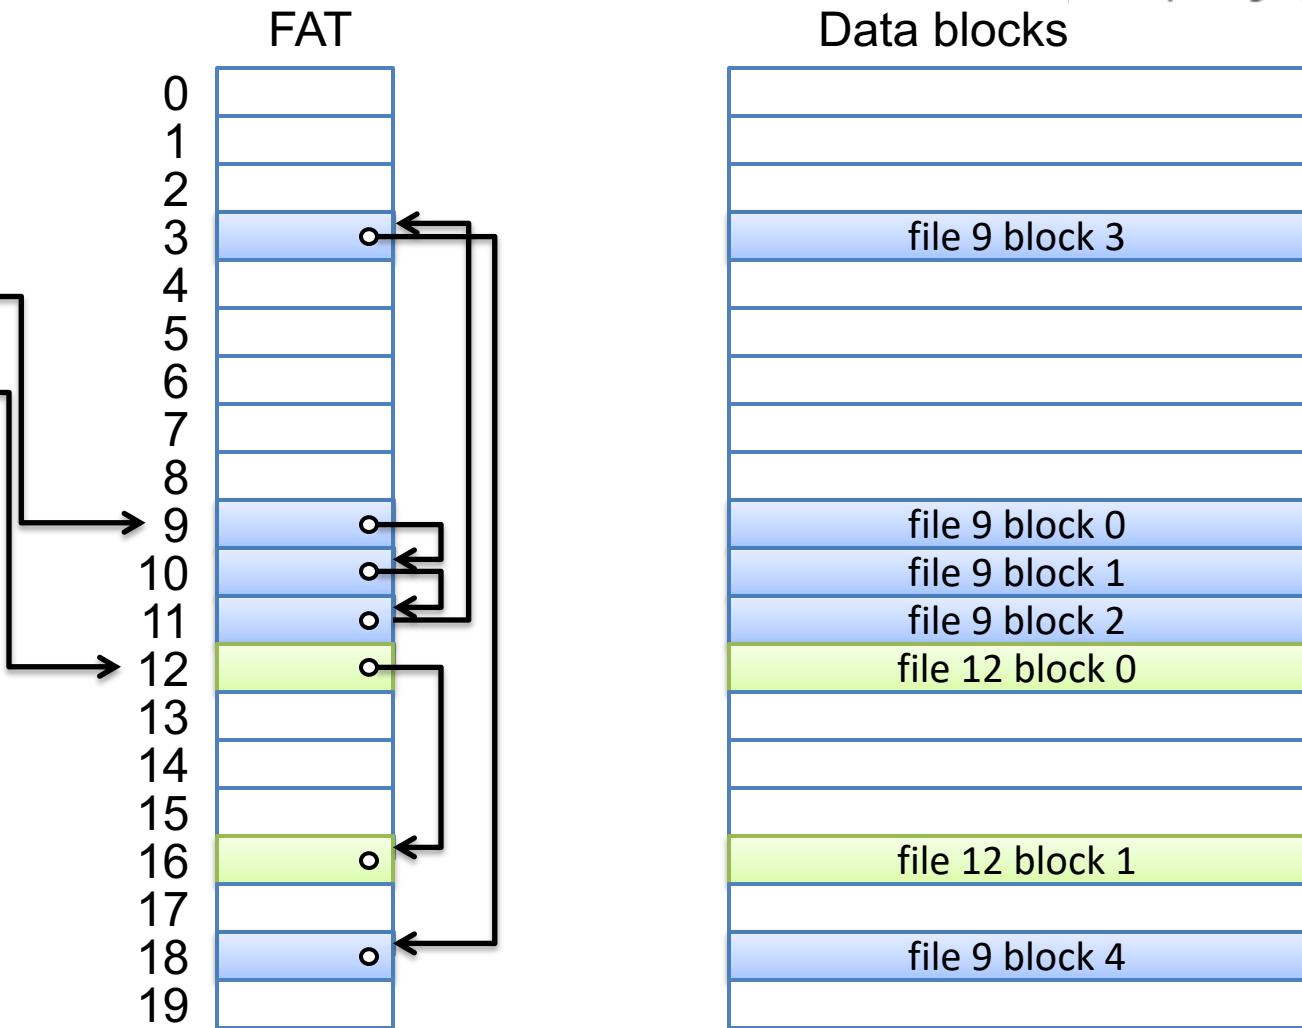
Bundesarchiv, Bild 145-Bild-Z077009-0042
Foto: Rehbein, Biegelbett 16. April 1988



FAT file system

Directory

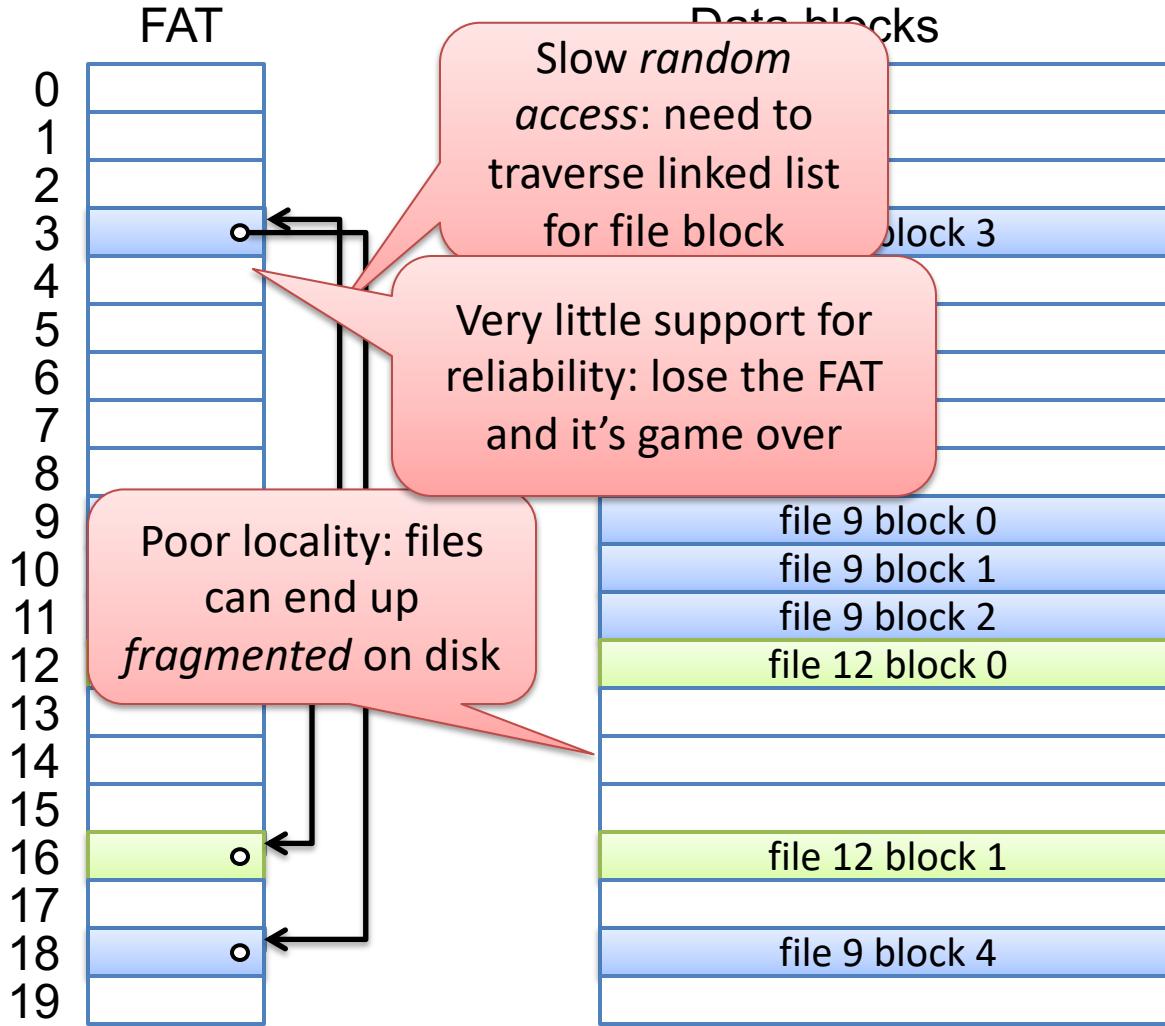
Foo	.exe	9
Bar	.doc	12





FAT file system

Free space:
Linear search
through FAT

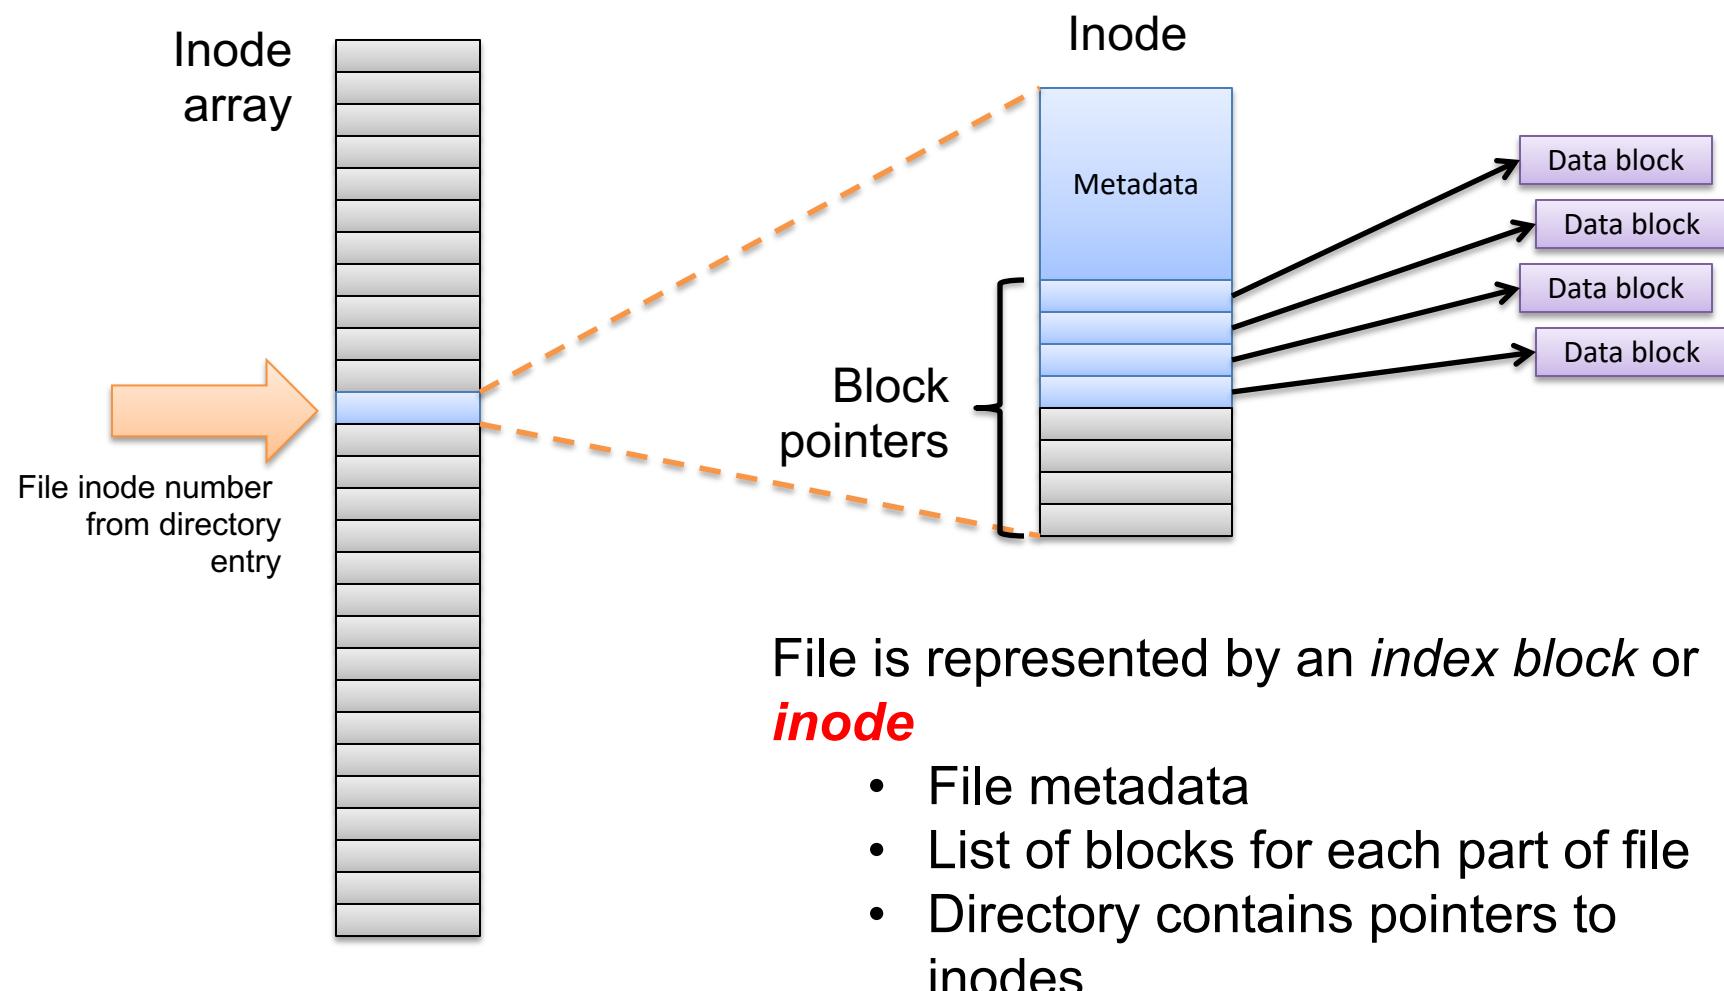




FFS



FFS uses indexed allocation



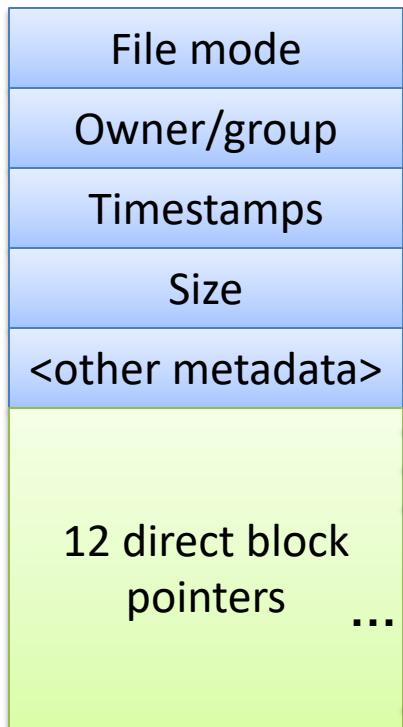


Inode and file size in FFS

- Example:
 - Inode is 1 block = 4096 bytes
 - Inode metadata = 512 bytes
 - Block addresses = 8 bytes
- Hence:
 - $(4096 - 512) / 8 = 448$ block pointers
 - $448 * 4096 = 1792\text{kB}$ max. file size

Unix file system inode format (simplified)

Inode:



(all blocks 4kB)

Data block

Data block

Data block

...

Data block

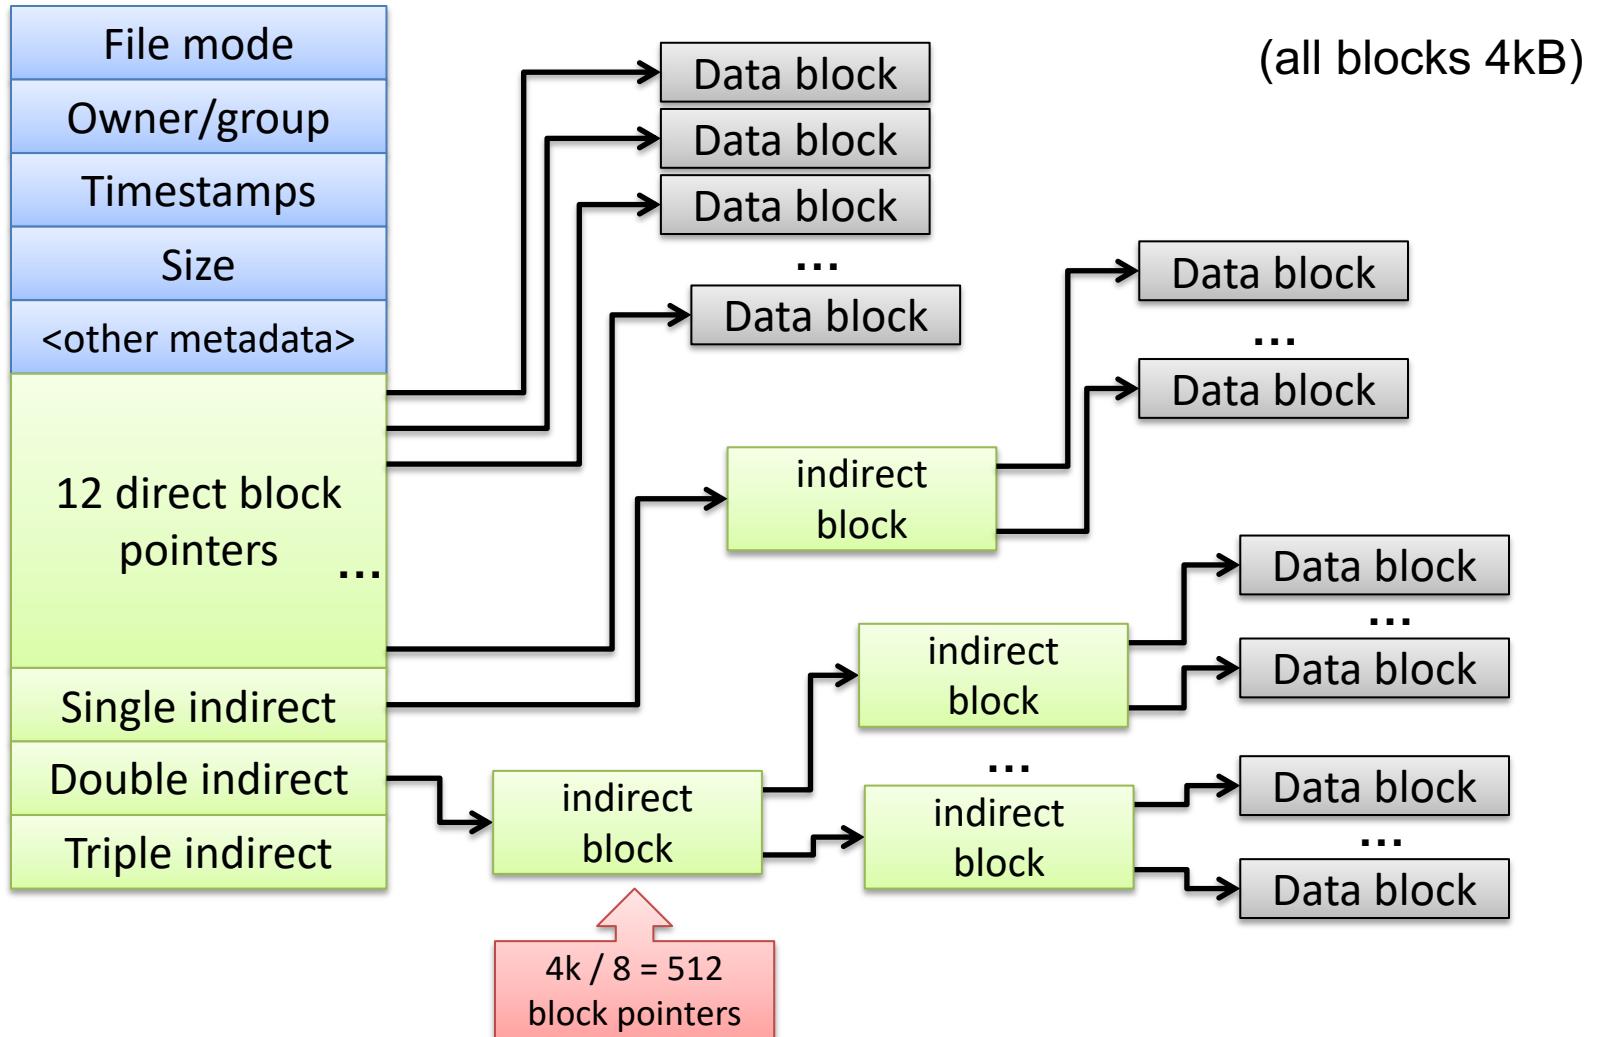
Question:

How to extend file size if there are no more block pointers in the Inode?



Unix file system inode format (simplified)

Inode:

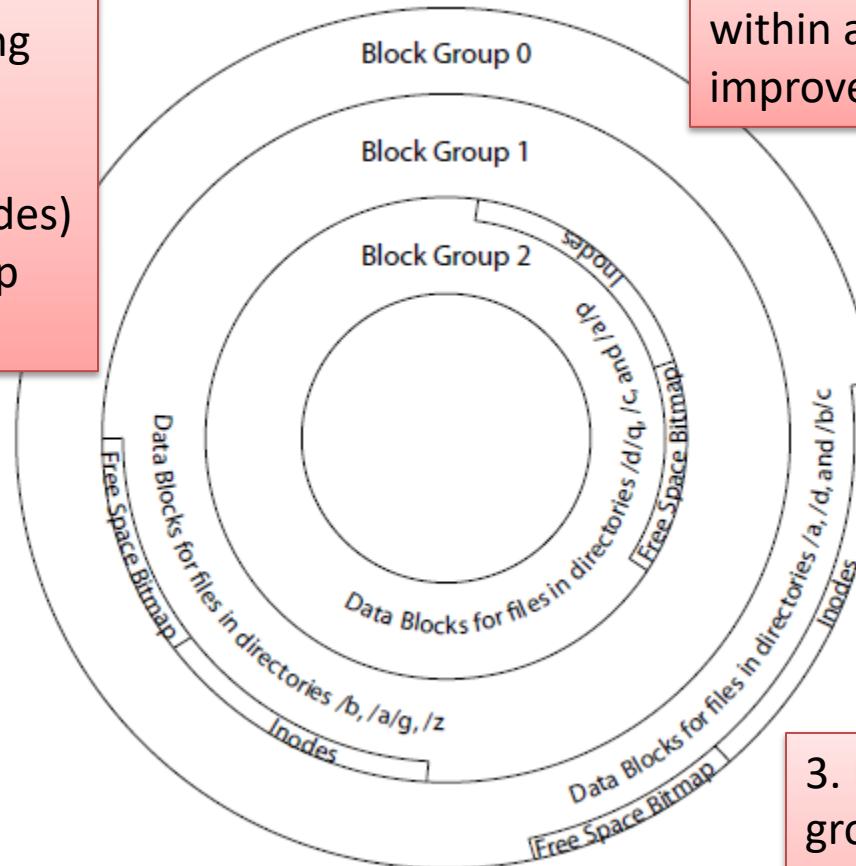


Block groups

1. Optimize disk performance by keeping together related:
- Files
 - Metadata (inodes)
 - Free space map
 - Directories

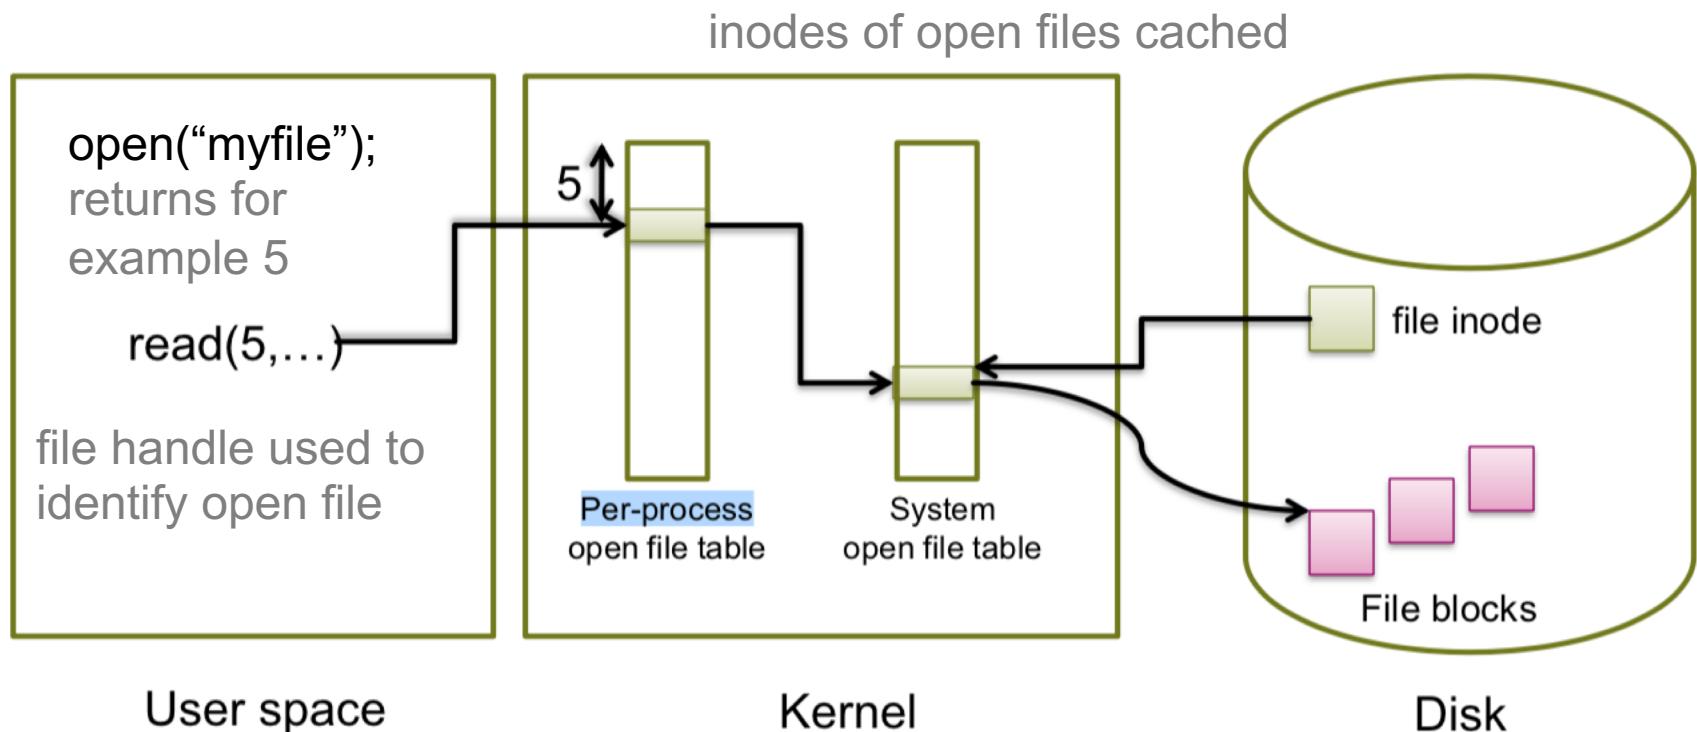
2. Use *first-fit* allocation within a block group to improve disk locality

3. Layout and block groups defined in the *superblock* (not shown); Replicated several times.





Open Files

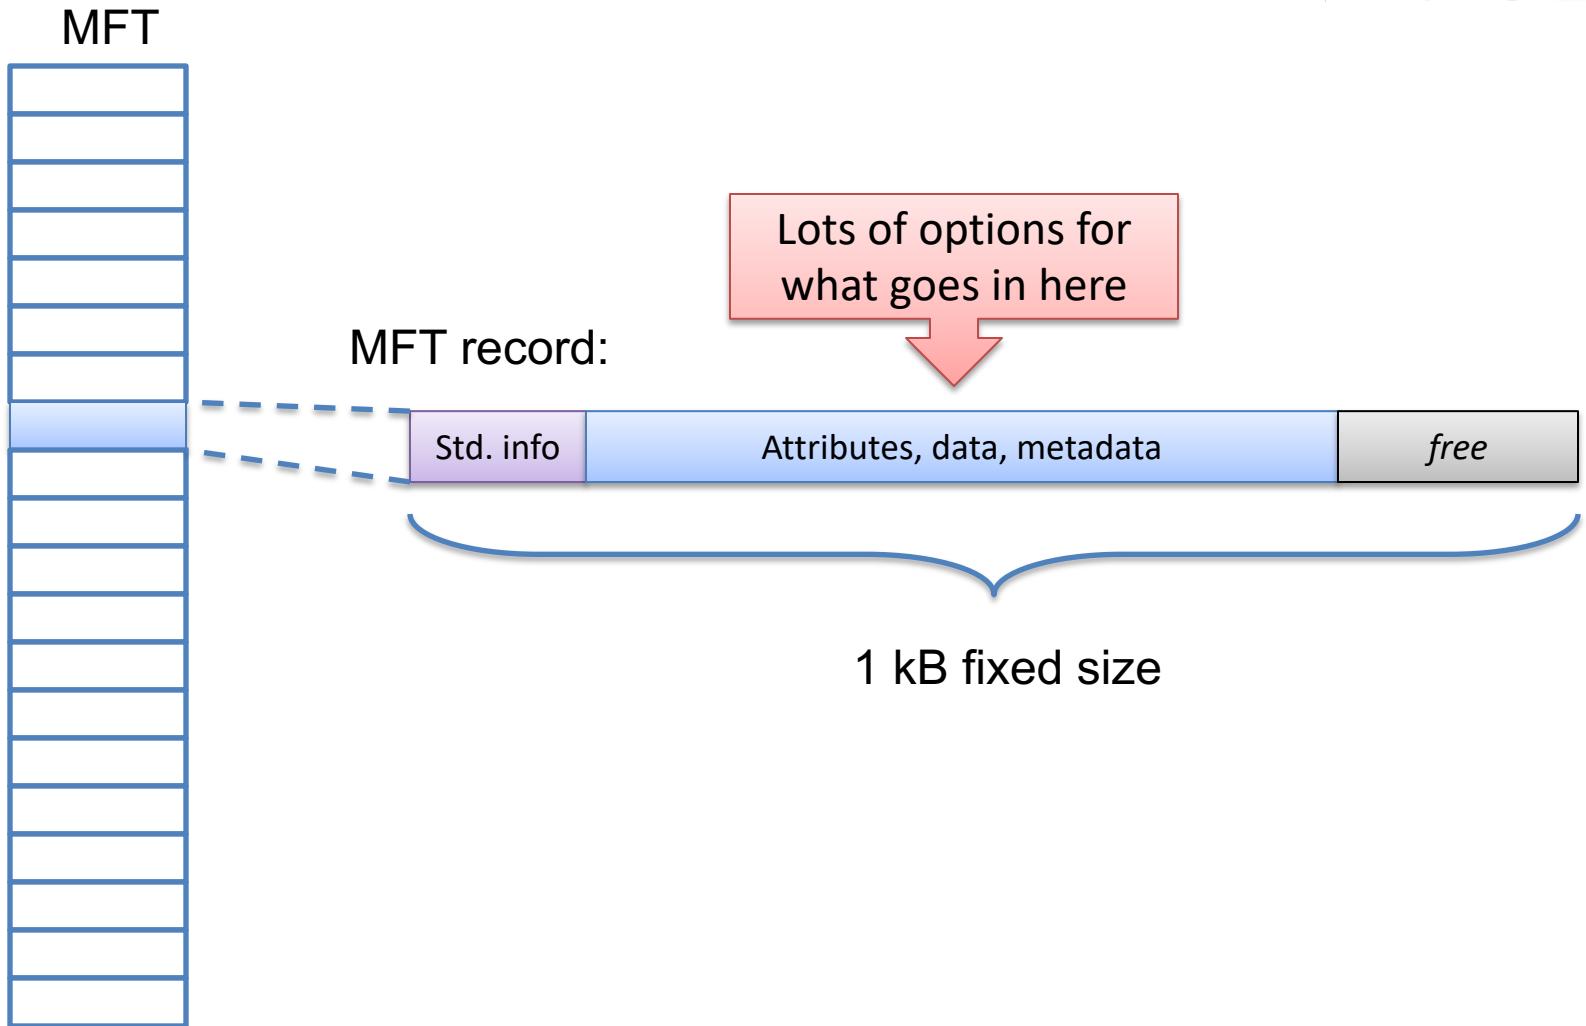




NTFS



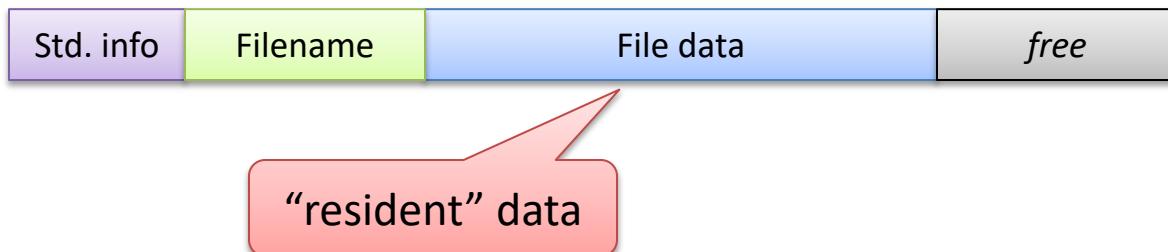
NFTS Master file table





NTFS small files

- Small file fits into MFT record:



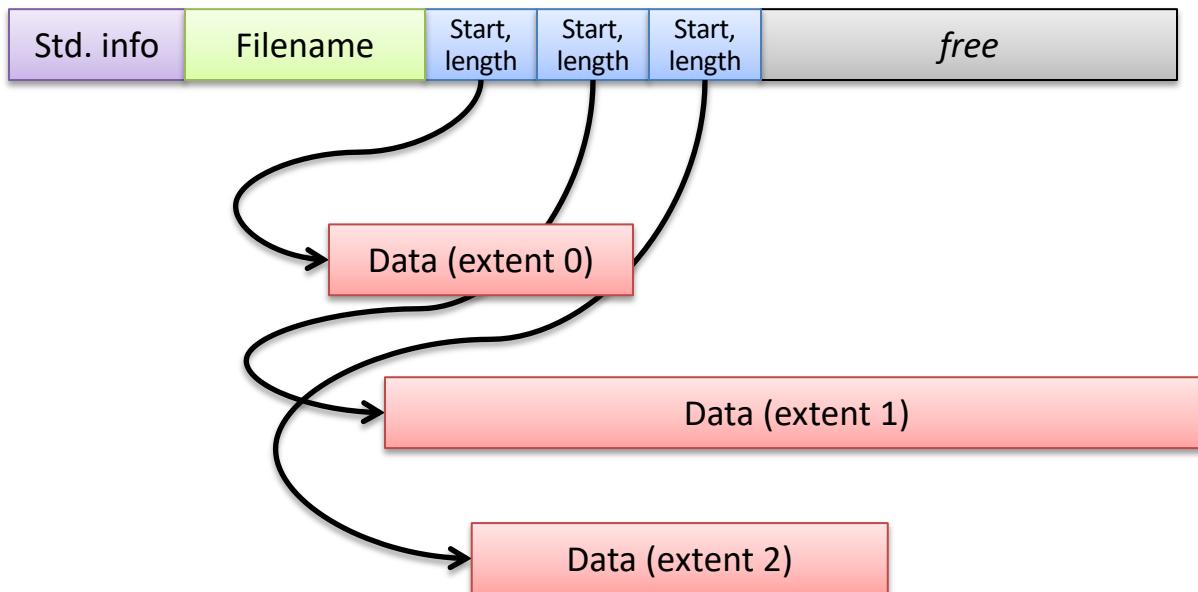
- Hard links (multiple names) stored in MFT:





NTFS normal files

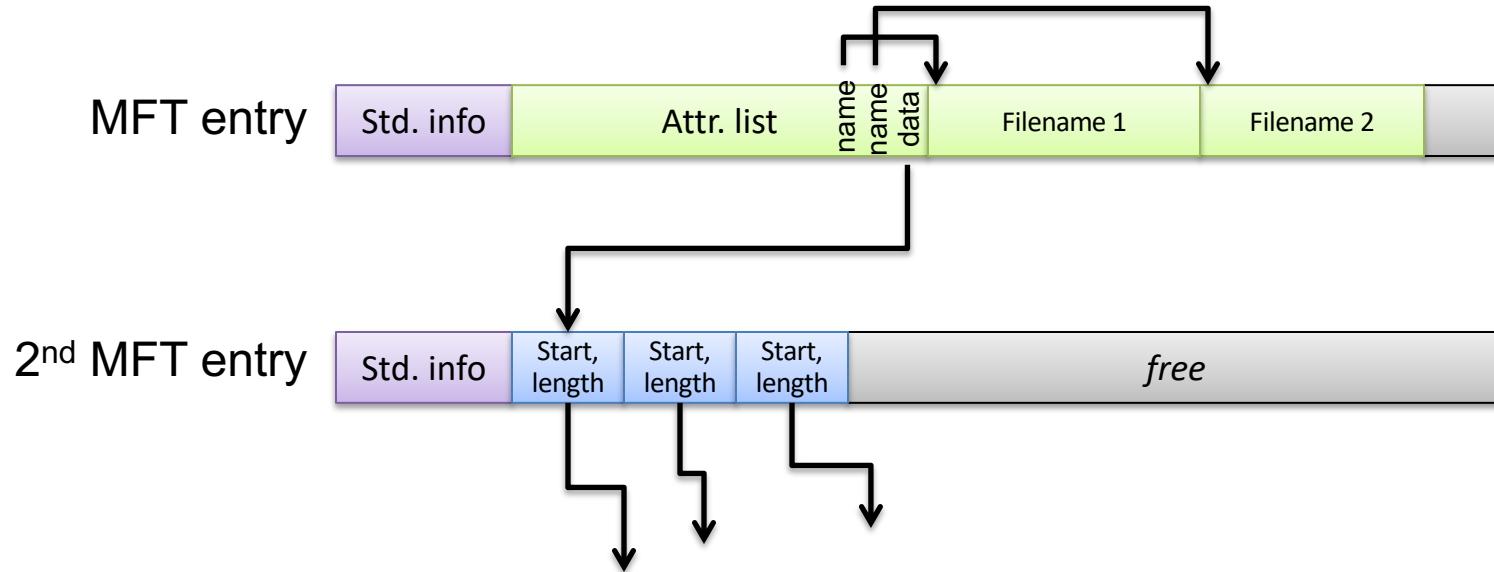
- MFT holds list of *extents*:





Too many attributes?

- Attribute list holds list of attribute locations



In addition, attributes can also be stored in extents \Rightarrow very large scaling



Metadata files

- File system metadata in NTFS is held *in files!*

File num.	Name	Description
0	\$MFT	Master file table
1	\$MFTirr	Copy of first 4 MFT entries
2	\$LogFile	Transaction log of FS changes
3	\$Volume	Volume information & metadata
4	\$AttrDef	Table mapping numeric IDs to attributes
5	.	Root directory
6	\$Bitmap	Free space bitmap
7	\$Boot	Volume boot record
8	\$BadClus	Bad cluster map
9	\$Secure	Access control list database
10	\$UpCase	Filename mappings to DOS
11	\$Extend	Extra file system attributes (e.g. quota)

Question:
Huh?
Where is it then?

Answer:
First sector of
volume points to
first block of MFT



File system implementations

	FAT	FFS	NTFS
Index structure	Linked list	Fixed, assymmetric tree	Dynamic tree
Index granularity:	Block	Block	Extent
Free space management	FAT Array	Fixed bitmap	Bitmap in file
Locality heuristics	Defragmentation	Block groups, Reserve space	Best fit, Defragmentation



Feedback

- <https://bit.ly/2Fuon6W>

