

# Exercise Session 8

## Computer Systems



# Agenda

- Virtual Memory
  - Segmentation
  - Paging
  - TLB
- Demand Paging
  - Page Replacement

# Terminology

- **Physical** address: address as seen by the memory unit
- **Virtual** or **Logical** address: address issued by the processor
  - Loads
  - Stores
  - Instruction fetches
  - Possible others (e.g. TLB fills)...

# Memory management


1. Allocating physical addresses to applications
  2. Managing the name translation of virtual addresses to physical addresses
  3. Performing access control on memory access
- Functions 2 & 3 usually involve the hardware Memory Management Unit (MMU)

# Segmentation

# Segmentation


- Generalize base + limit:
  - Physical memory divided into *segments*
  - Logical address = (segment id, offset)
- Segment identifier supplied by:
  - Explicit instruction reference
  - Explicit processor segment register
  - Implicit process state

# User's View of a Program




logical address

# Segmentation reality



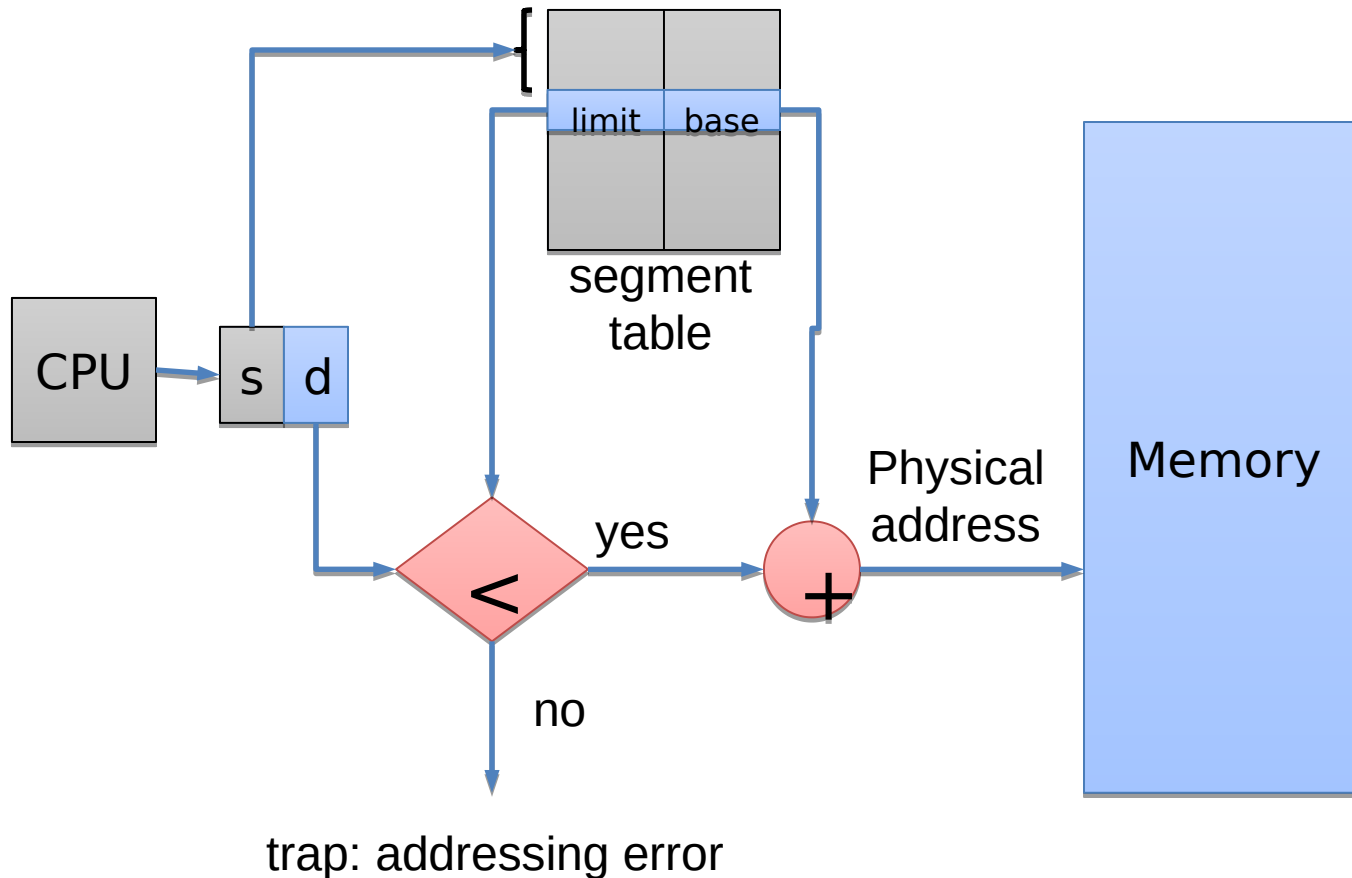
logical address




physical memory



# Segmentation Hardware




# Segmentation reality



logical address

	limit	base
0	300	1500
1	1000	5000
2	400	3400
3	400	4600
4	1000	1800

segment  
table



physical memory

# Segmentation Architecture

- **Segment table** – each entry has:
  - **base** – starting physical address of segment
  - **limit** – length of the segment
- **Segment-table base register (STBR)**
  - Current segment table location in memory
- **Segment-table length register (STLR)**
  - Current size of segment table

segment number  **$s$**  is legal if  **$s < \text{STLR}$**

# Paging

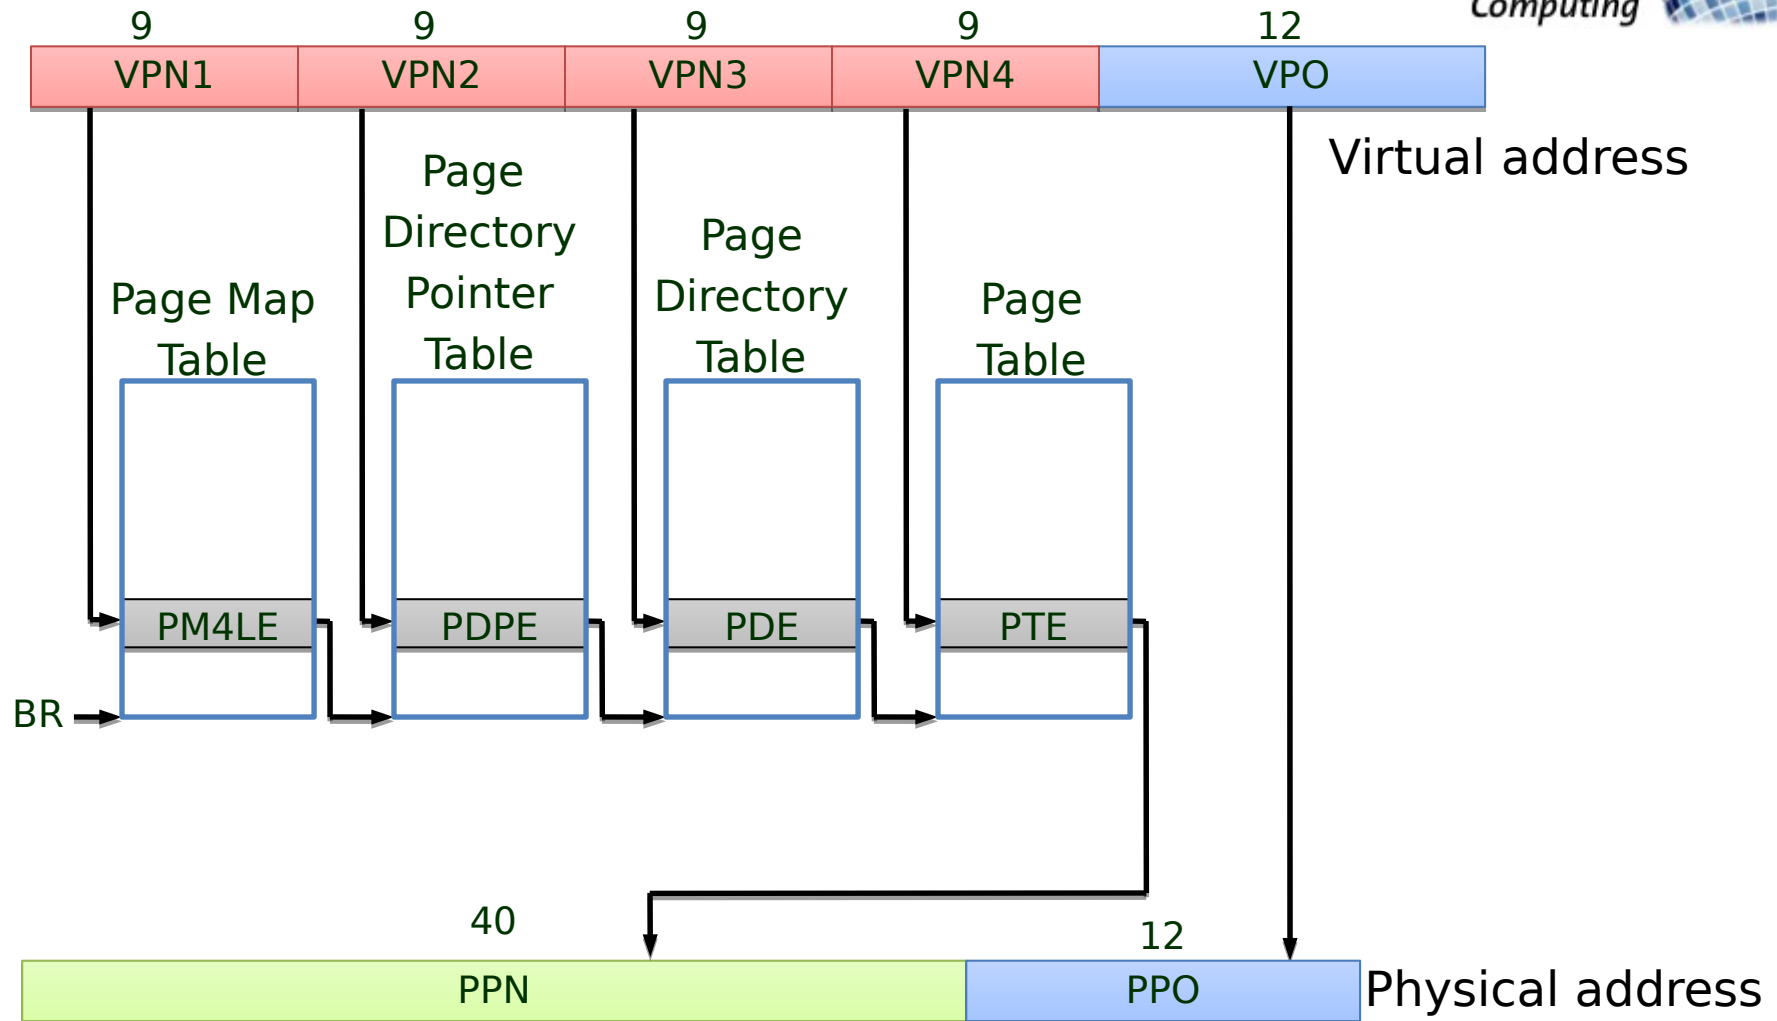
# Paging

- Solves contiguous physical memory problem
  - Process can always fit if there is available free memory
- Divide physical memory into **frames**
  - Size is power of two, e.g. 4096 bytes
- Divide logical memory into **pages** of the same size
- For a program of  $n$  pages in size:
  - Find and allocate  $n$  frames
  - Load program
  - Set up **page table** to translate logical pages to physical frames

# Page table jargon

- Page tables maps VPNs to PFNs
  - Page table entry = PTE
- **VPN** = Virtual Page Number
  - Upper bits of virtual or logical address
- **PFN** = Page Frame Number
  - Upper bits of physical or logical address
- Same number of bits (usually).

# x86-64 Paging



# Translation Lookaside Buffer (TLB)

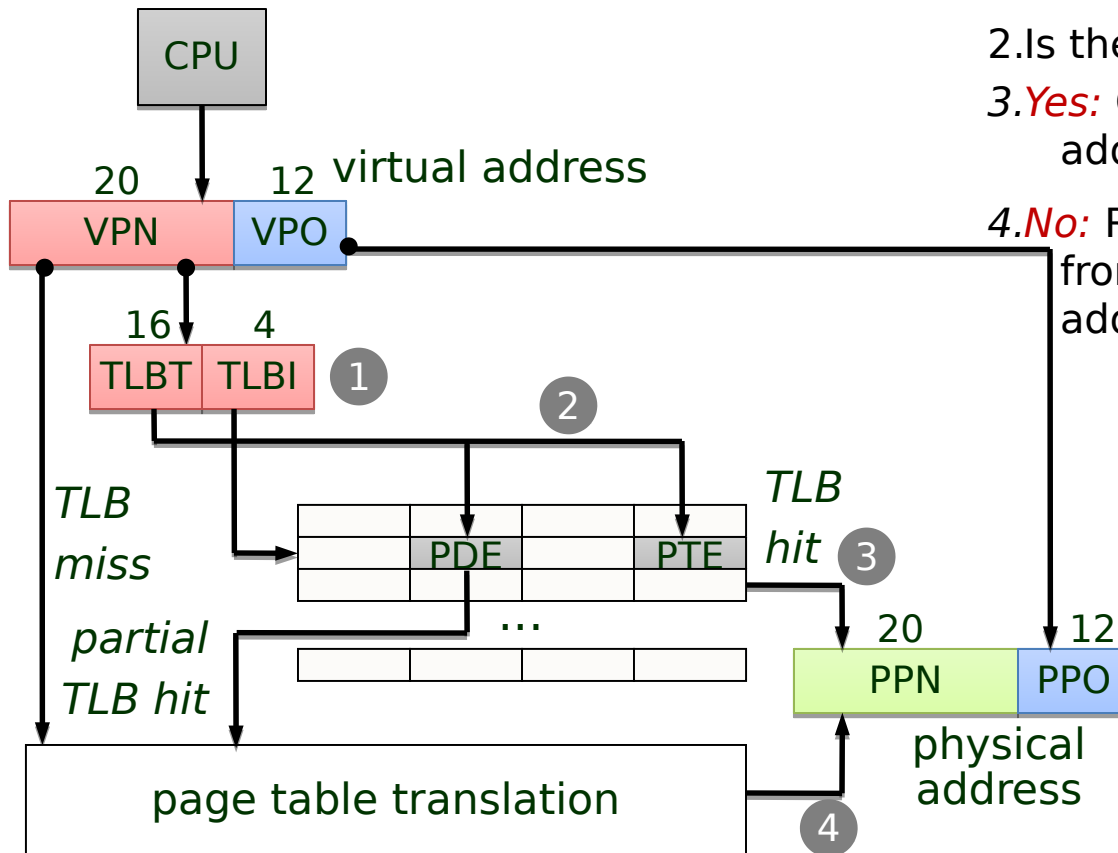


# Problem: performance

- Every logical memory access needs two physical memory accesses
  - Load page table entry
  - Load desired location
- Performance  $\Rightarrow$  half as fast as with no translation
  - Solution: cache page table entries



# Translating with the P6 TLB



1. Partition VPN into TLBT and TLBI.
2. Is the PTE for VPN cached in set TLBI?
3. **Yes:** Check permissions, build physical address
4. **No:** Read PTE (and PDE if not cached) from memory and build physical address

# Caches and Paging

# Cache Management

- Virtual Index, Virtual Tag
  - +No need to wait for translation
  - Might get the wrong data if not flushed
- Physical Index, Physical Tag
  - +Easy to manage
  - Need to wait for address translation

# Cache Management


- Virtual Index, Physical Tag
  - +Address lookup can happen in parallel with cache lookup
  - Either limited size or overhead
- Physical Index, Virtual Tag
  - No

# Virtual Memory really does make your life easier!

# Example: ESP8266

- Small Wifi-enabled microcontroller
- Typically 4MB flash storage for code
- Used in embedded applications #IoT
- Wouldn't it be nice if we could update our smart flower pots without disassembling them?

# OTA updates



Idea: Run from one partition and update the other one

This works great!

Until someone tries to use the PC to address something...

Solution: Link two images and never flash the wrong one!



Startseite > Technik > Handwerkzeuge > [Äxte & Hämmer](#)

🔍 Vollbild

## LUX Fäustel 1.250 g Classic

★★★★★ (0)



- Gewicht: 1'250 g
- Glasfasergriff
- Nach DIN gefertigt

[Alle Artikelinfos](#)

Menge

1

**Fr. 12.90**

Preis inkl. gesetzl. MWST. 7.7%

**Reservieren & im Markt selbst abholen**

Abholbereit ab nächstem Öffnungstag 11 Uhr im OBI Markt Volketswil

Im OBI Markt Volketswil  
weniger als 3 Artikel vorrätig[Verfügbarkeit in anderem Markt prüfen](#)☐ [Artikel vergleichen](#) [Artikel merken](#)

Artikelbeschreibung ▾ Bewertungen (0) ▾

## Artikelbeschreibung

Art.Nr. 3439098

★★★★★ (0) [Bewertungen lesen](#)

Der LUX-TOOLS Fäustel 1.250 g hat einen geschmiedeten Kopf und ist mit einem stabilen, gummierten Griff aus Glasfaser verklebt.

Gewicht: 1.250 g

Material des Kopfes: Werkzeugstahl.

Nach DIN 1193 und DIN 6473.

## Technische Daten

### Produktmerkmale

Art:	Fäustel
Typ:	Sonstige

### Masse und Gewicht

## Ähnliche Produkte



LUX Sappie 600 g Comfort

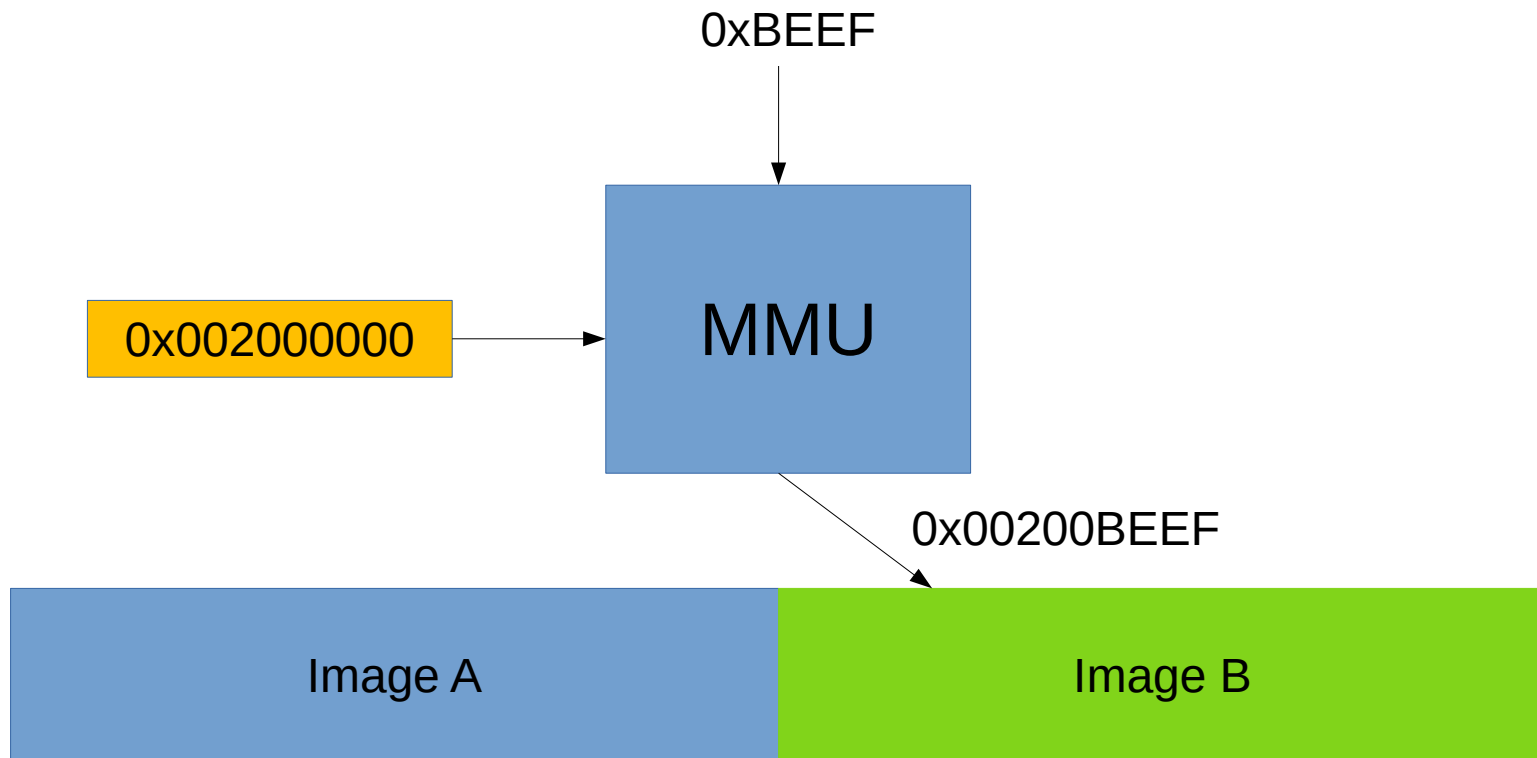
★★★★★

**Fr. 44.-**LUX Latthammer Ganzstahl  
770 g Comfort

★★★★★

**Fr. 43.-**

# Better Solution: Use an MMU!



# But Wait! There's More!

- Protection features can be abused for fun and profit
- E.g. Copy-On-Write:
- Share a page between two processes
- Make it read-only
- On a write, create a copy of that page and write there

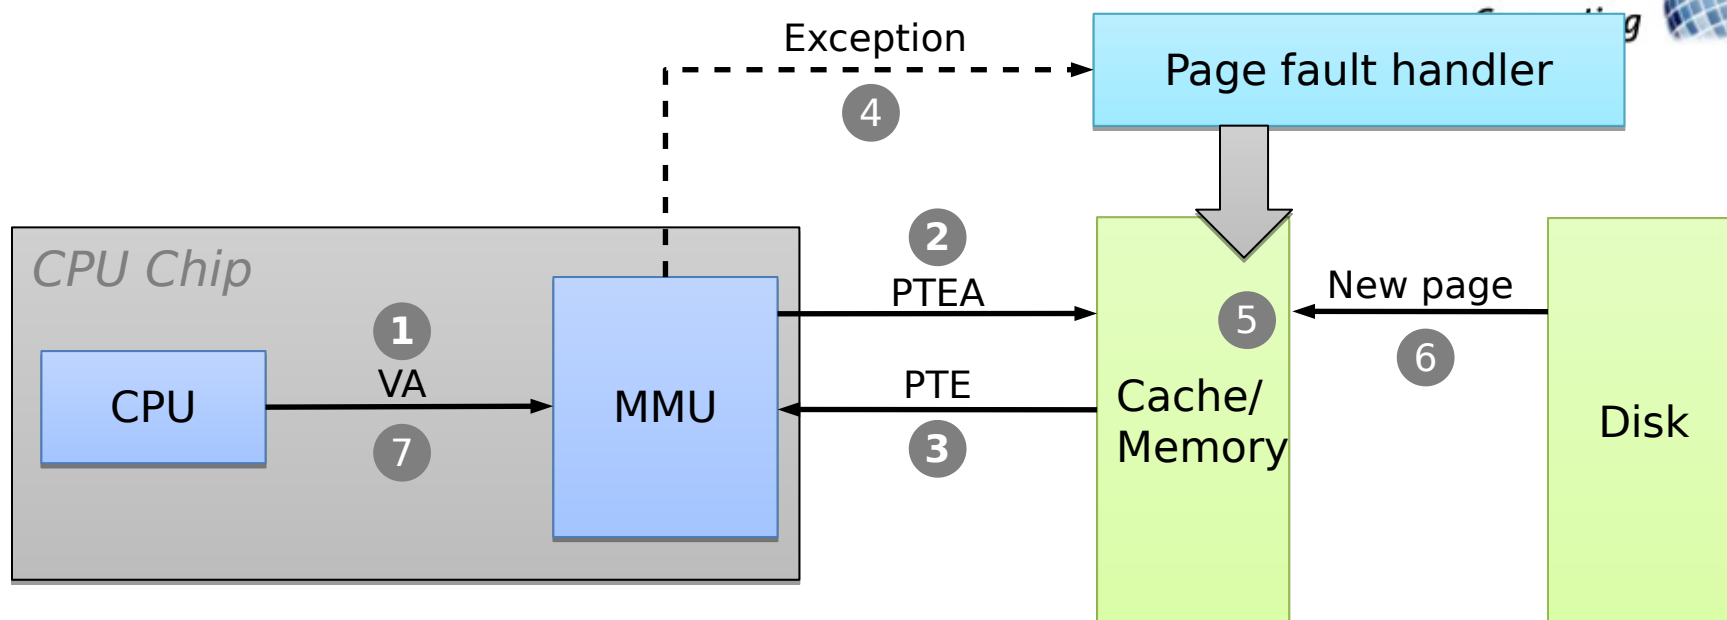
# Demand Paging

# Demand Paging

- Bring a page into memory only when it is needed
  - Less I/O needed
  - Less memory needed
  - Faster response
  - More users
- Turns RAM into a *cache* for processes on *disk*!



# Recall: handling a page fault




- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler finds a frame to use for missing page
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

# Anyone remember this?



# Anyone remember THIS?





# Page Replacement

# What happens if there is no free frame?

- *Page replacement* – find “little used” resident page to discard or write to disk
  - “victim page”
  - algorithm
  - performance – want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times

# FIFO (First-In-First-Out) page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

# FIFO (First-In-First-Out) page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

7

# FIFO (First-In-First-Out) page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

7	7
	0

# FIFO (First-In-First-Out) page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

7	7	7
	0	0
		1

# FIFO (First-In-First-Out) page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

7	7	7	2
	0	0	0
		1	1

# FIFO (First-In-First-Out) page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

7	7	7	2	2
	0	0	0	3
		1	1	1



# FIFO (First-In-First-Out) page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

7	7	7	2
	0	0	0
		1	1

2	2	4	4	4	0
3	3	3	2	2	2
1	0	0	0	3	3

# FIFO (First-In-First-Out) page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

7	7	7	2
	0	0	0
		1	1

2	2	4	4	4	0
3	3	3	2	2	2
1	0	0	0	3	3

0	0
1	1
3	2

# FIFO (First-In-First-Out) page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

7	7	7	2
	0	0	0
		1	1

2	2	4	4	4	0
3	3	3	2	2	2
1	0	0	0	3	3

0	0
1	1
3	2

7	7	7
1	0	0
2	2	1

Here, 15 page faults.

# Least Recently Used (LRU) algorithm

- Reference string: 1 2 3 4 1 2 **5** 1 2 **3** **4** **5**

4 frames

1
2
3
4

1
2
<b>5</b>
4

1	1	<b>5</b>
2	2	2
5	<b>4</b>	4
<b>3</b>	3	3

- Counter implementation
  - Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
  - When a page needs to be changed, look at the counters to determine which are to change

# LRU page replacement

reference string: 7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

3 page  
frames:

7	7	7	2		2		4	4	4	0		1		1		1			
	0	0	0		0		0	0	3	3		3		0		0			
		1	1		3		3	2	2	2		2		2		7			

Here, 12 page faults.

Used for measuring how well your algorithm performs