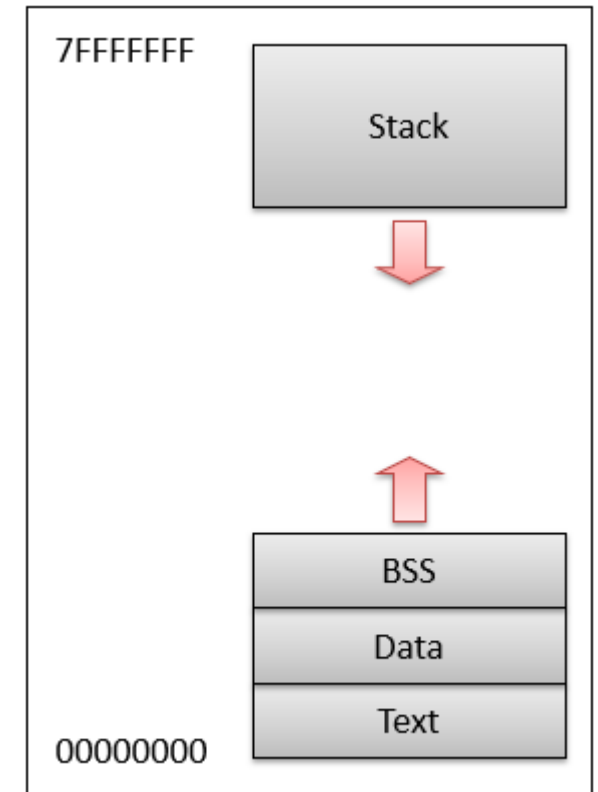


# Computer Systems

## Exercise 2

# Process

- Instance of a program
- Ingredients of a process:
  - Virtual processor (address space, registers, program counter, instruction pointer)
  - Program text (code)
  - Program data (heap, stack)
  - OS stuff (open files, sockets, CPU shares, security rights)
- Executed in an “execution environment” (virtual address space, available system calls, etc.)
- Purpose of process concept
  - Protection of program data
  - Running of program
- Identified by PID (process ID)
- Software = running processes + kernel



# Process creation – spawning new process

Did it work?

```
BOOL CreateProcess (  
    in_opt      LPCTSTR      ApplicationName,   
    inout_opt   LPTSTR       CommandLine,   
    in_opt      LPSECURITY_ATTRIBUTES ProcessAttributes,   
    in_opt      LPSECURITY_ATTRIBUTES ThreadAttributes,   
    in          BOOL          InheritHandles,   
    in          DWORD         CreationFlags,   
    in_opt      LPVOID        Environment,   
    in_opt      LPCTSTR       CurrentDirectory,   
    in          LPSTARTUPINFO  StartupInfo,   
    out         LPPROCESS_INFORMATION ProcessInformation   
);
```

address of binary

What to run?

What rights will it have?

What will it see when it starts up?

The result

Moral: the parameter space is large!

# Process creation – fork() and exec()

- Simplifies creating processes:
    - Fork creates a copy of the current process, same only that child has return value of fork() = 0 and parent has return value of fork() = PID of child
    - Exec() replaces text of calling process with new program.
- Creates tree of processes

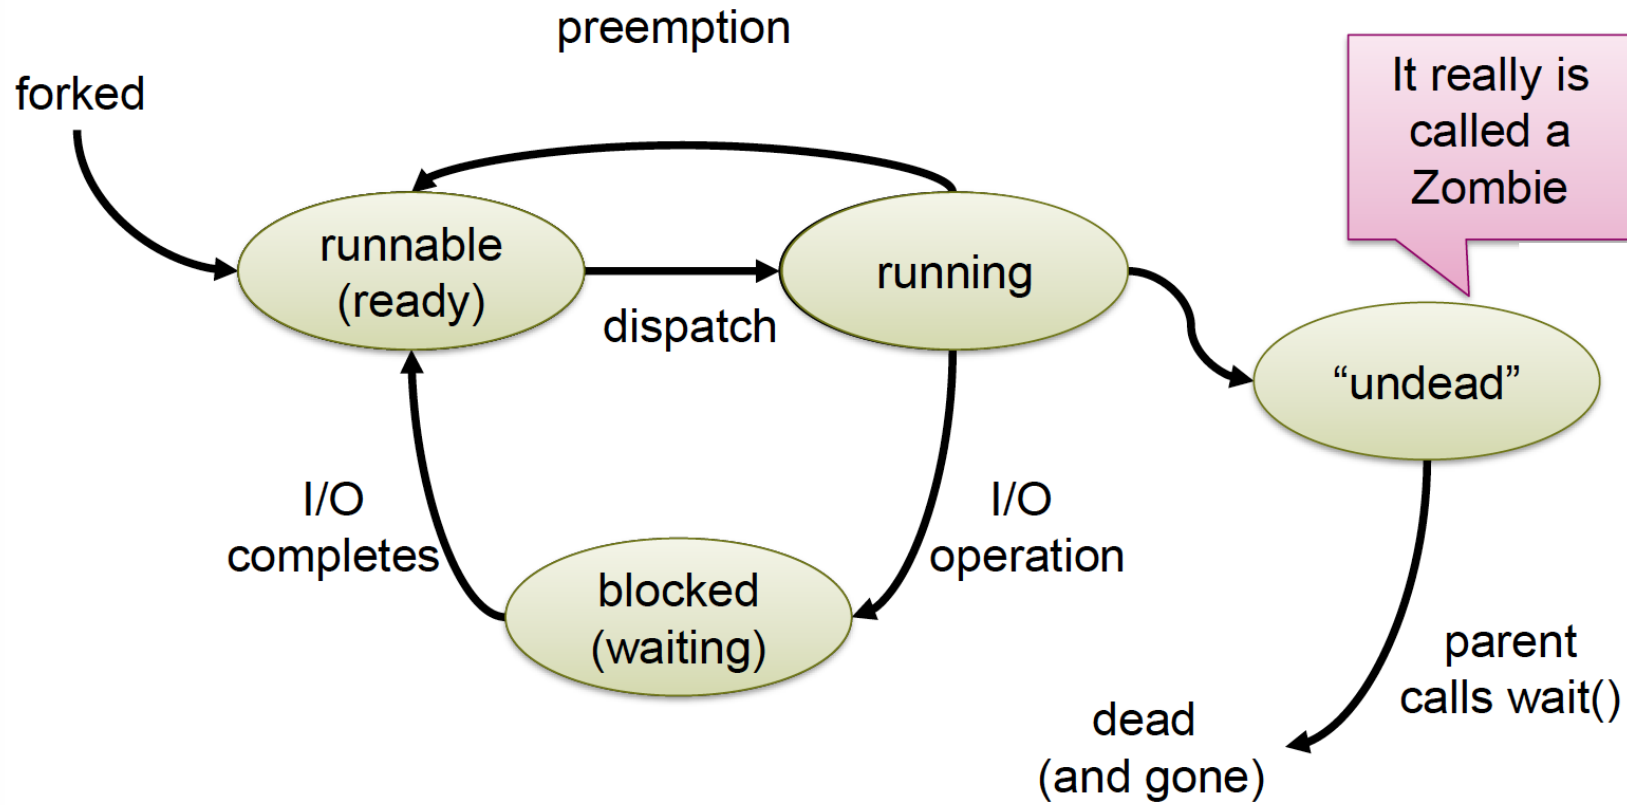
# Example fork(), exec()

```
pid_t p = fork();
if ( p < 0 ) {
    // Error...
    exit(-1);
} else if ( p == 0 ) {
    // We're in the child
    execlp("/bin/ls", "ls", NULL);
} else {
    // We're a parent.
    // p is the pid of the child
    wait(NULL);
    exit(0);
}
```

Return code from  
fork() tells you  
whether you're in the  
parent or child  
(cf. setjmp())

Child process can't  
actually be cleaned  
up until parent  
"waits" for it.

# Process lifecycle



# Zombies & Orphans

- Why Zombies?
  - If no Zombies, child process could fail and nobody would know because parent has no chance to catch exit status
- What happens with child process whose parents have exited?
  - They are called orphans and get “adopted” by the init process (PID = 1)

# Coroutines


- Concurrency without parallelism
- Basically two functions that call each other (symmetric relationship)



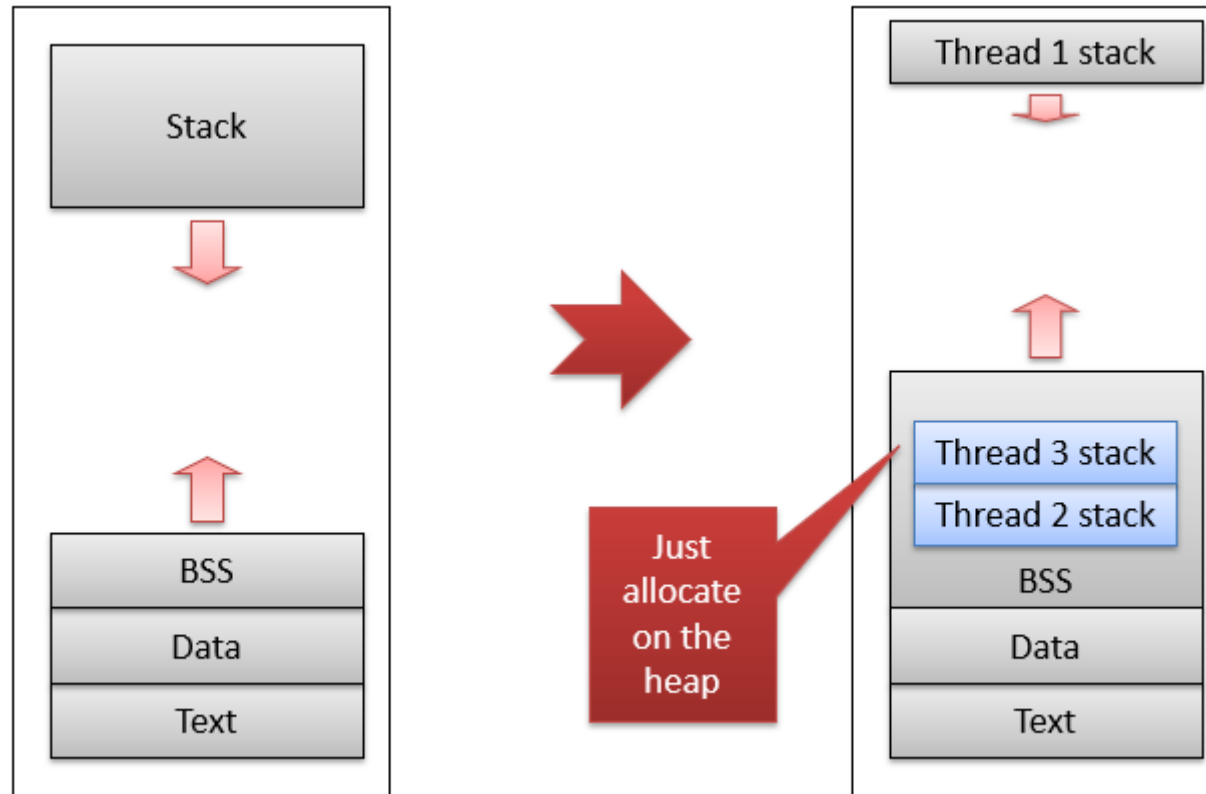
# Threads

- A thread is part of a process. Multiple threads can exist in one process and share resources such as memory.
- Process – unit of resources, thread – unit of scheduling/execution
- User threads (lightweight processes)
  - Kernel is unaware of them, scheduled in user space
  - Fast to create and manage, but can't make use of multithreading
- Kernel threads
  - At least one kernel thread exists for each process
  - One kernel thread can be mapped to each logical core and can be swapped once it gets blocked, but takes long to swap

# User vs. Kernel Threads



# Address space for user-level threads



# How do processes communicate?

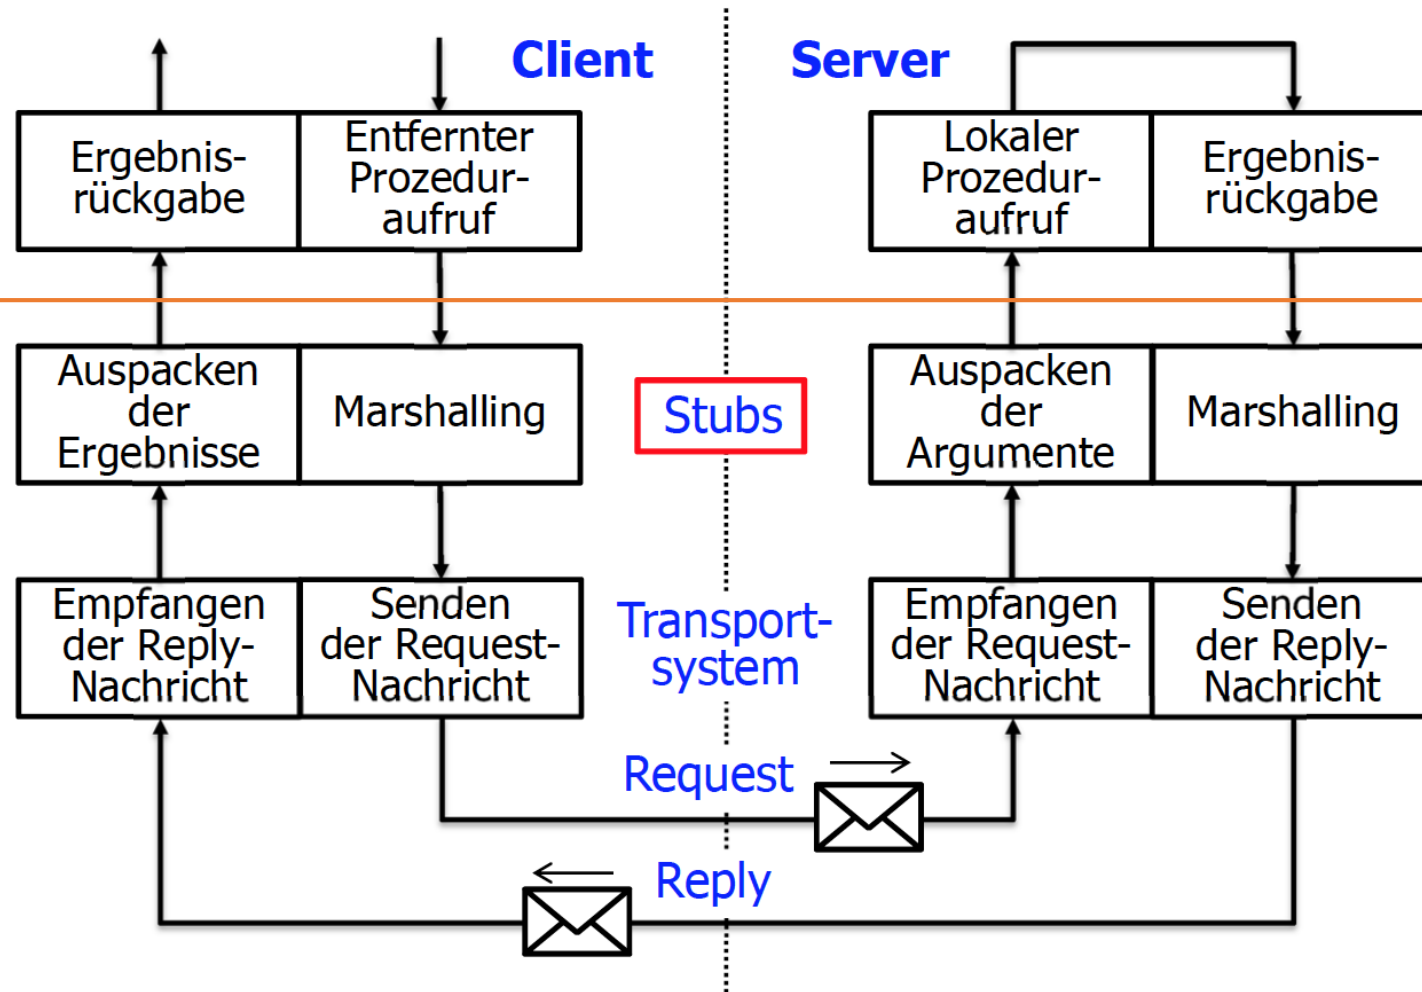
- **Shared Memory**

- Semaphores/Locks
  - Synchronization instructions (CAS, TAS, LL/SC)
- Transactional Memory

- **Message passing**

- Asynchronous/Synchronous
- Blocking/non-blocking
- Pipes
  - Named/unnamed
- Upcalls/Signals
  - SIGSEGV from memory management
  - SIGKILL from other process
- RPC

# RPC



# Quiz

- **What is the relation between a process and a program?**
  - A process is a running instance of a program
- **How are processes identified?**
  - By the process ID (PID)
- **In what state is a process after it exited?**
  - Zombie state
- **Are user threads or kernel threads easier to switch?**
  - User space threads
- **How long should you spin for waiting for a lock?**
  - One context switch time