

Distributed Systems

10. Quorum-Based Consensus: Paxos

Paul Krzyzanowski

Rutgers University

Fall 2018

Consensus Goal

Allow a group of processes to agree on a result

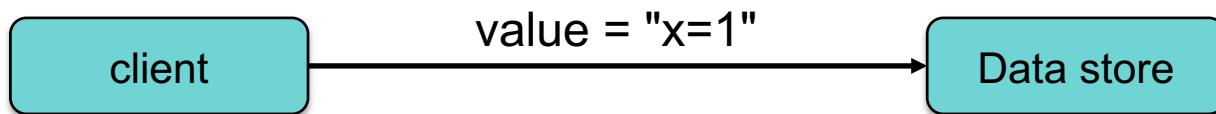
- All processes must agree on the same value
- The value must be one that was submitted by at least one process
(the consensus algorithm cannot just make up a value)

We saw versions of this

- Mutual exclusion
 - Agree on who gets a resource or who becomes a coordinator
- Election algorithms
 - Agree on who is in charge
- Other uses of consensus:
 - Synchronize state to manage replicas: make sure every group member agrees on the message ordering of events
 - Manage group membership
 - Agree on distributed transaction commit
- General consensus problem:
 - *How do we get unanimous agreement on a given value?*

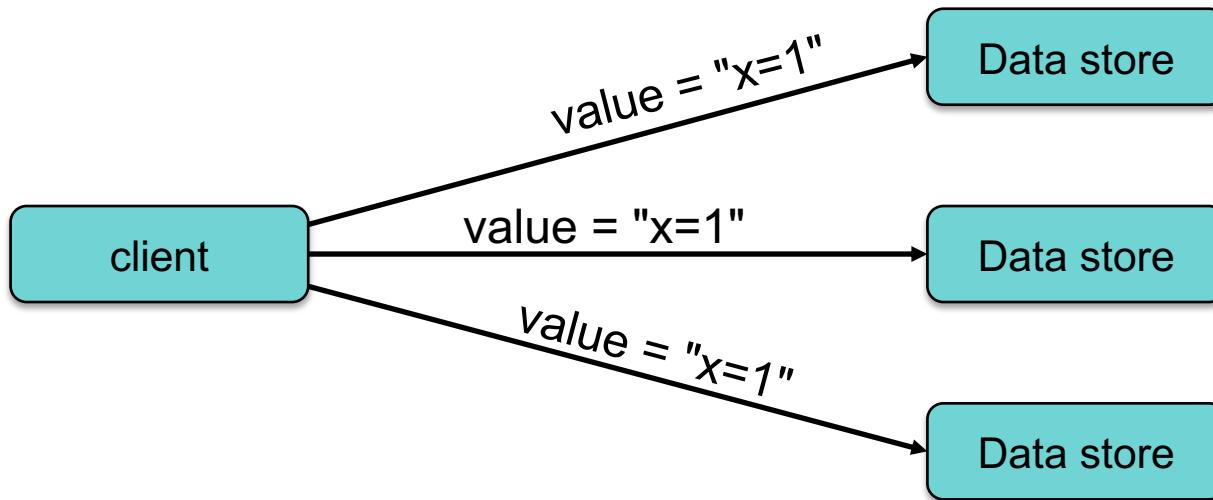
value = sequence number of a message, key=value, operation, whatever...

Achieving consensus seems easy!



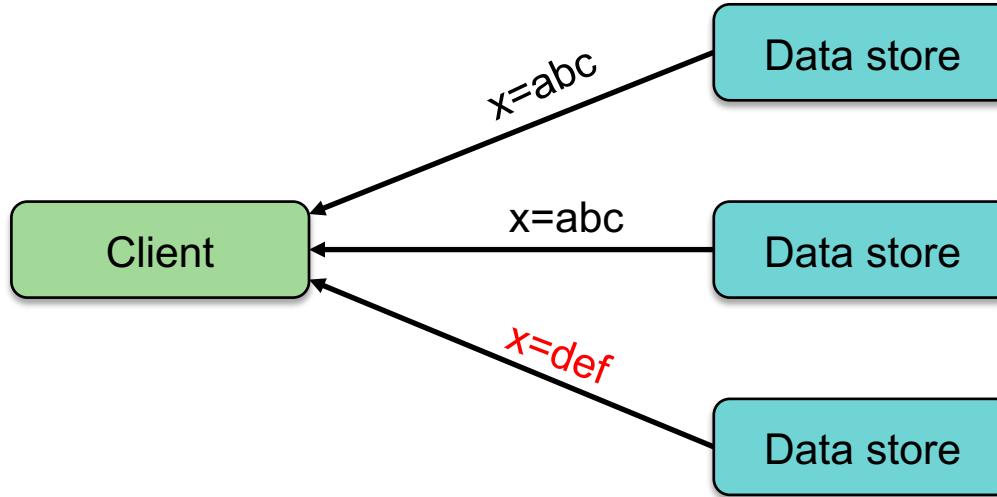
- One request at a time
- Server that never dies

Servers might die – let's add replicas



- One request at a time

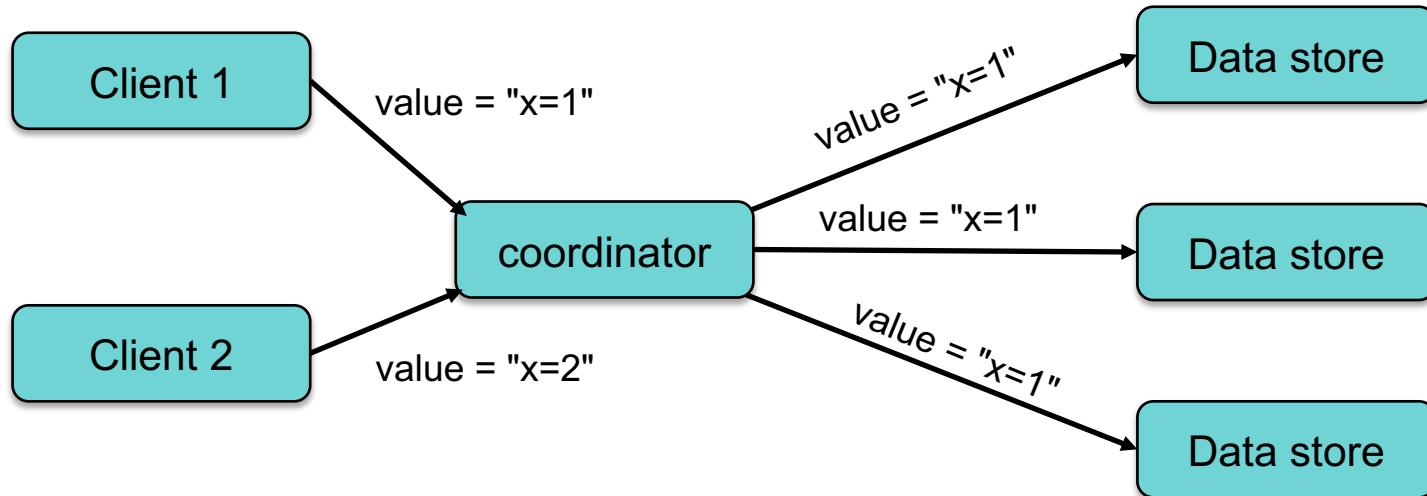
Reading from replicas is easy



We rely on a quorum (majority) to read successfully

No quorum = failed read!

What about concurrent updates?



- Coordinator processes requests one at a time
- But now we have a single point of failure!
- We need something safer

Consensus algorithm goal

- Create a fault-tolerant consensus algorithm that does not block if a **majority of processes** are working
- Goal: agree on one result among a group of participants
 - Processors may fail (some may need stable storage)
 - Messages may be lost, out of order, or duplicated
 - If delivered, messages are not corrupted

Quorum: majority (>50%) agreement is the key part: If a majority of coins show heads, there is no way that a majority will show tails at the same time.

If members die and others come up, **there will be one member in common** with the old group that still holds the information.

Consensus requirements

- Validity
 - Only proposed values may be selected
- Uniform agreement
 - No two nodes may select different values
- Integrity
 - A node can select only a single value
- Termination (Progress)
 - Every node will eventually decide on a value

Consensus algorithm: Paxos

Paxos

Goal: provide a consistent ordering of events from multiple clients

- All machines running the algorithm **agree on a proposed value** from a client
- The *value* will be associated with an event or action
- Paxos ensures that no other machine associates the value with another event

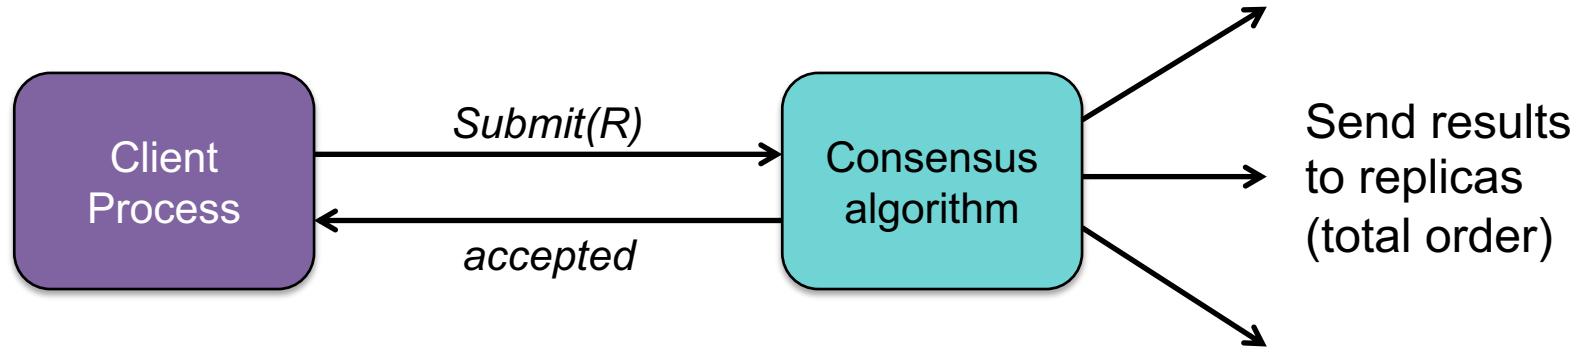
Fault-tolerant distributed consensus algorithm

- Does not block if a majority of processes are working
- The algorithm needs a **majority** ($2P+1$) of processors survive the simultaneous failure of P processors

Paxos provides **abortable consensus**

- A client's request may be rejected
- It then has to re-issue the request

A Programmer's View



If your request is not accepted, you can submit it again later:

```
while (submit_request(R) != ACCEPTED) ;
```

Think of R as a **key:value** pair in a database where multiple clients might want to modify the same key

Paxos players

- **Client:** makes a request
- **Proposers:**
 - Get a request from a client and run the protocol to get everyone in the cluster to agree
 - **Leader:** elected coordinator among the proposers
(not necessary but simplifies message numbering and ensures no contention) – we don't need to rely on the presence of a single leader
- **Acceptors:**
 - Multiple processes that remember the state of the protocol
 - **Quorum** = any majority of acceptors
- **Learners:**
 - When agreement has been reached by acceptors, a Learner executes the request and/or sends a response back to the client

These different roles are usually part of the same system

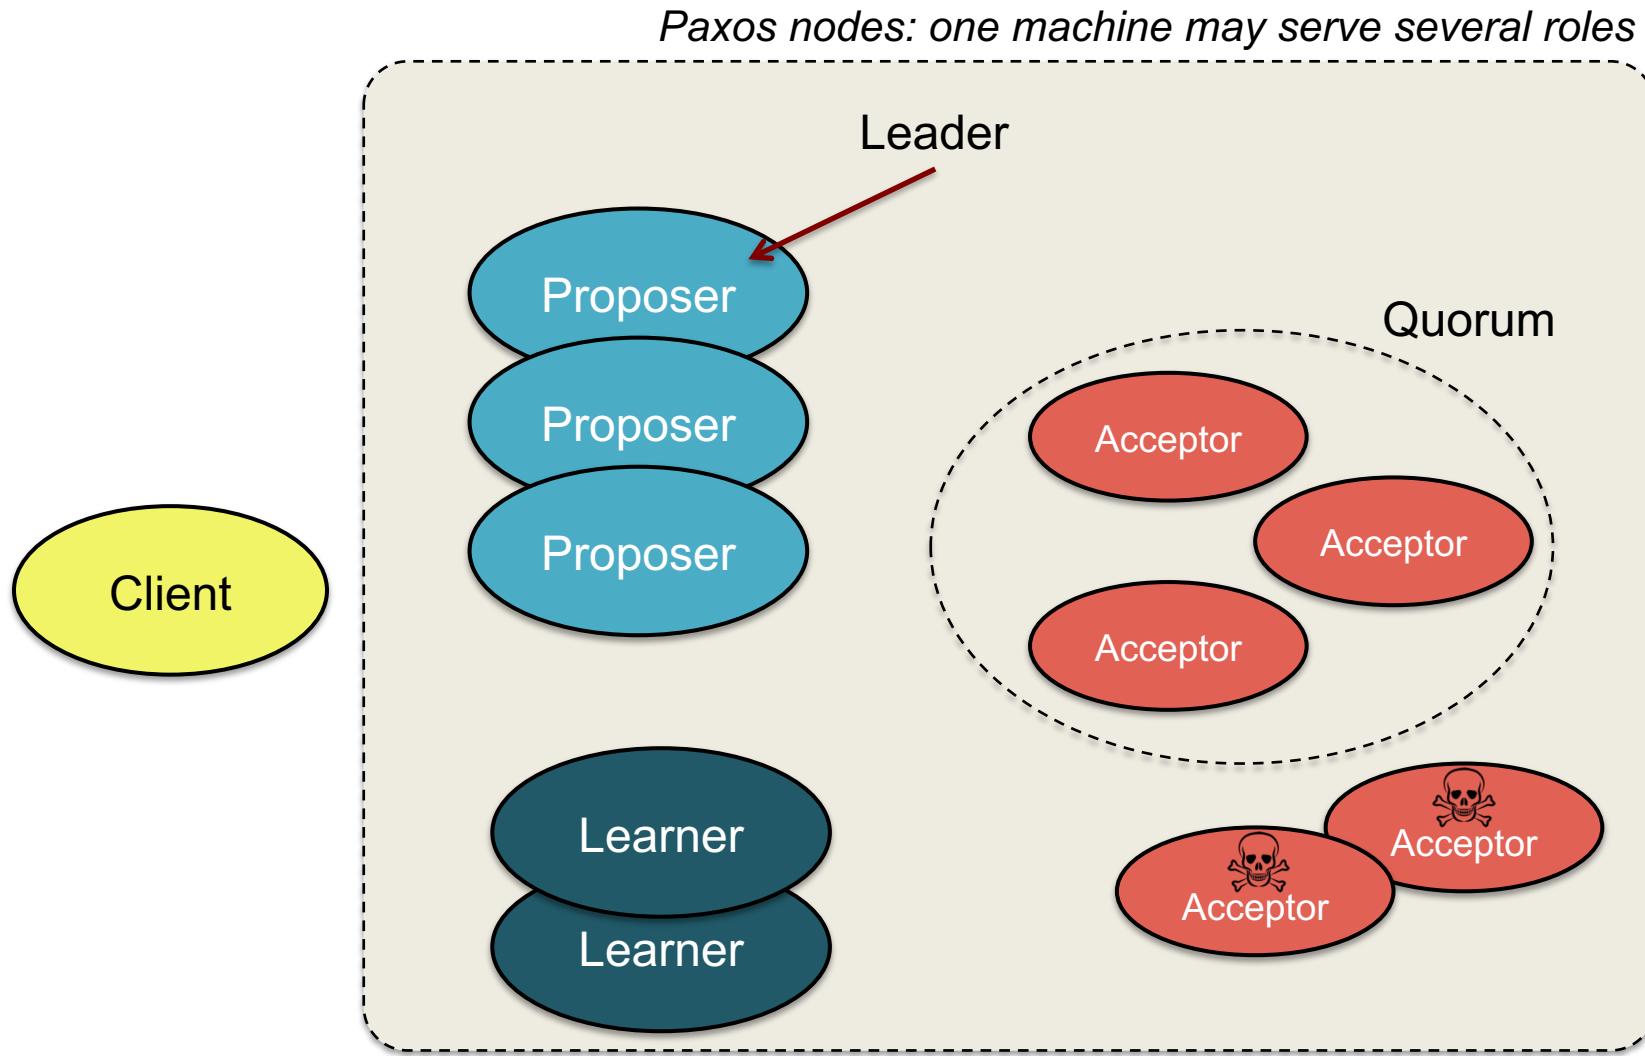
Proposal numbers

- **Paxos ensures a consistent ordering in a cluster of machines**
 - Events are ordered by sequential event IDs (N)
- Client wants to log an event: sends request to a Proposer
 - E.g., $v = \text{"add \$100 to my checking account"}$
- **Proposer**
 - Increments the latest proposal number (event ID) it knows about
 - ID = sequence number
 - Asks all the acceptors to reserve that proposal #
- **Acceptors**
 - A majority of acceptors have to accept the requested proposal #

Proposal Numbers

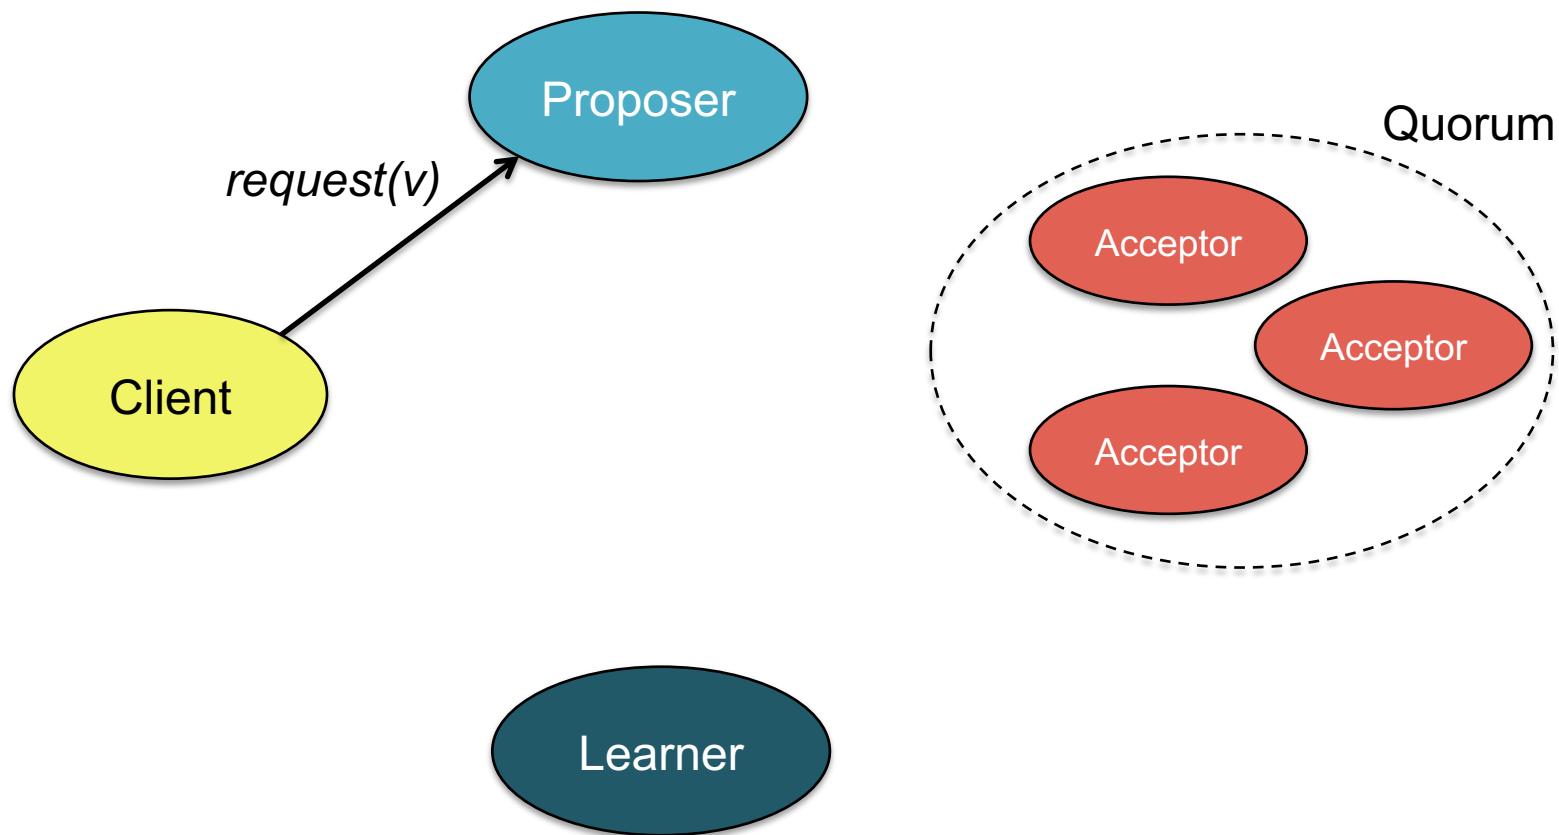
- Each proposal has a **unique** number (created by proposer)
 - Must be unique (e.g., <sequence #>.<process_id>)
- Newer proposals take precedence over older ones
- Each acceptor
 - Keeps track of the largest number it has seen so far
 - Lower proposal numbers get rejected
 - Acceptor sends back the $\{number, value\}$ of the currently accepted proposal
 - Proposer has to “play fair”:
 - It will ask the acceptors to accept the $\{number, value\}$
 - Either its own or the one it got from the acceptor

Paxos in action



Paxos in action: Phase 0

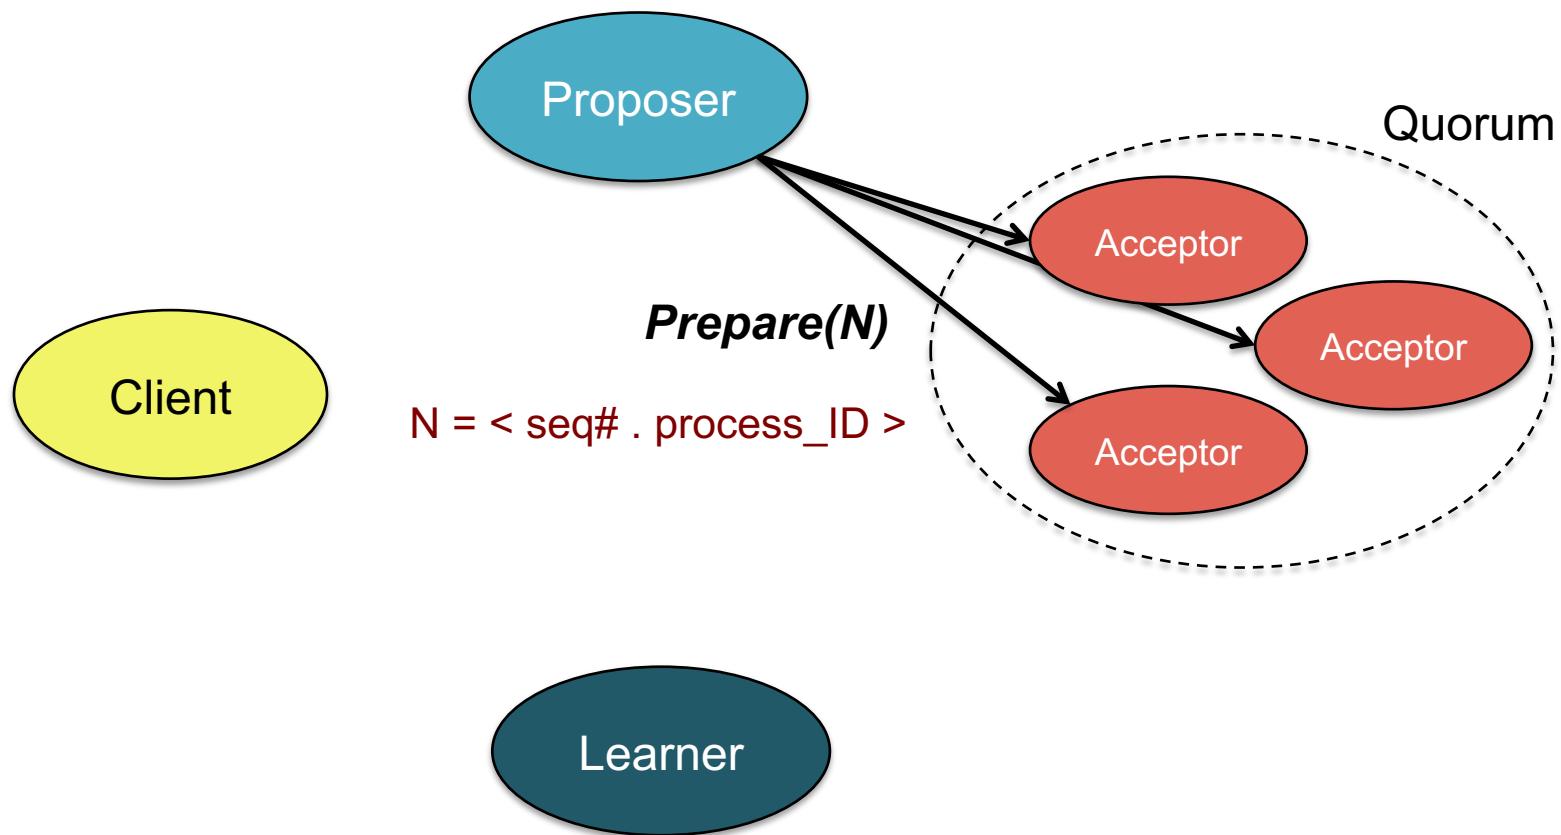
Client sends a request to a proposer



Paxos in action: Phase 1a – *PREPARE*

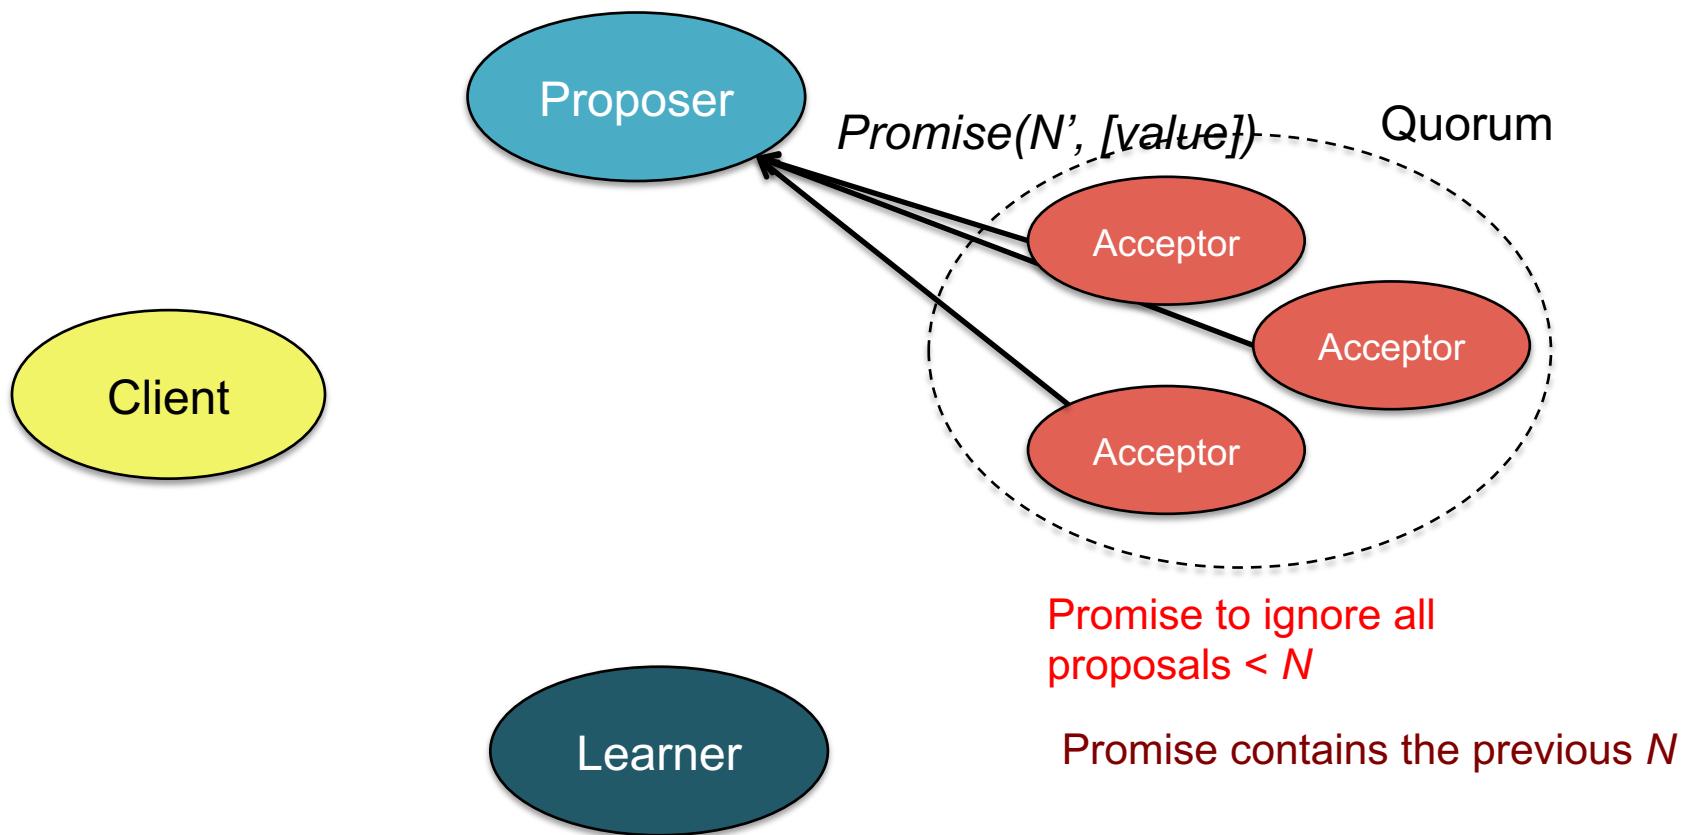
Proposer: creates a *proposal #N* (*N* acts like a *Lamport time stamp*), where *N* is greater than any previous proposal number used by this proposer

Send to Quorum of Acceptors (however many you can reach – but a majority)



Paxos in action: Phase 1b – PROMISE

Acceptor: if proposer's ID > any previous proposal
promise to ignore all requests with IDs < N
reply with info about highest past proposal: { N' , value }



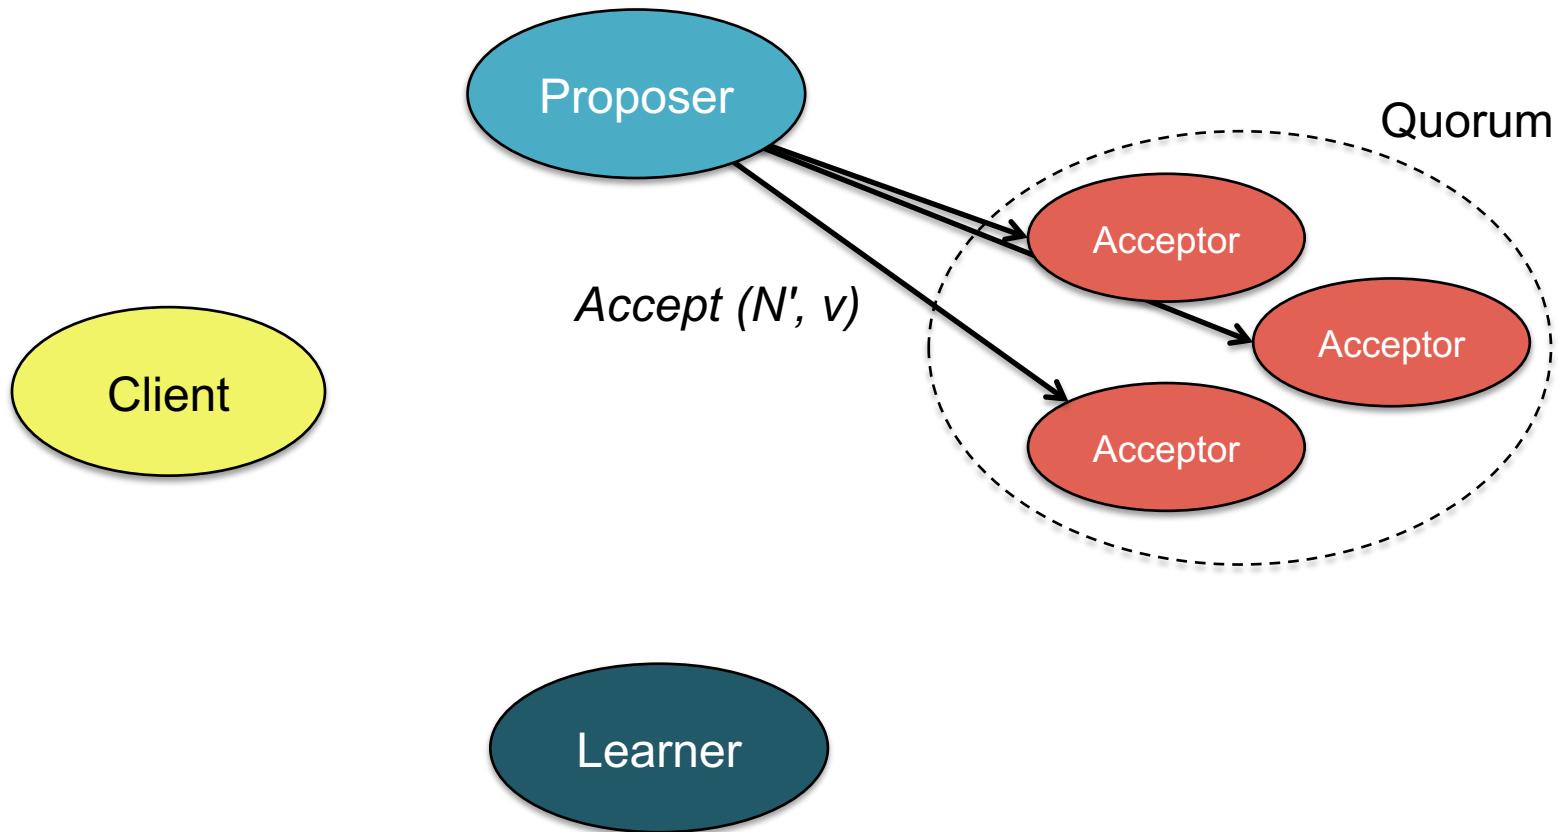
Paxos in action: Phase 2a – ACCEPT

Proposer: if proposer receives promises from the quorum (majority):

 Attach a value v to the proposal (the event).

 Send **Accept** to quorum with the **chosen** value

If promise was for another $\{N', v\}$, proposer MUST accept v for the highest accepted proposal

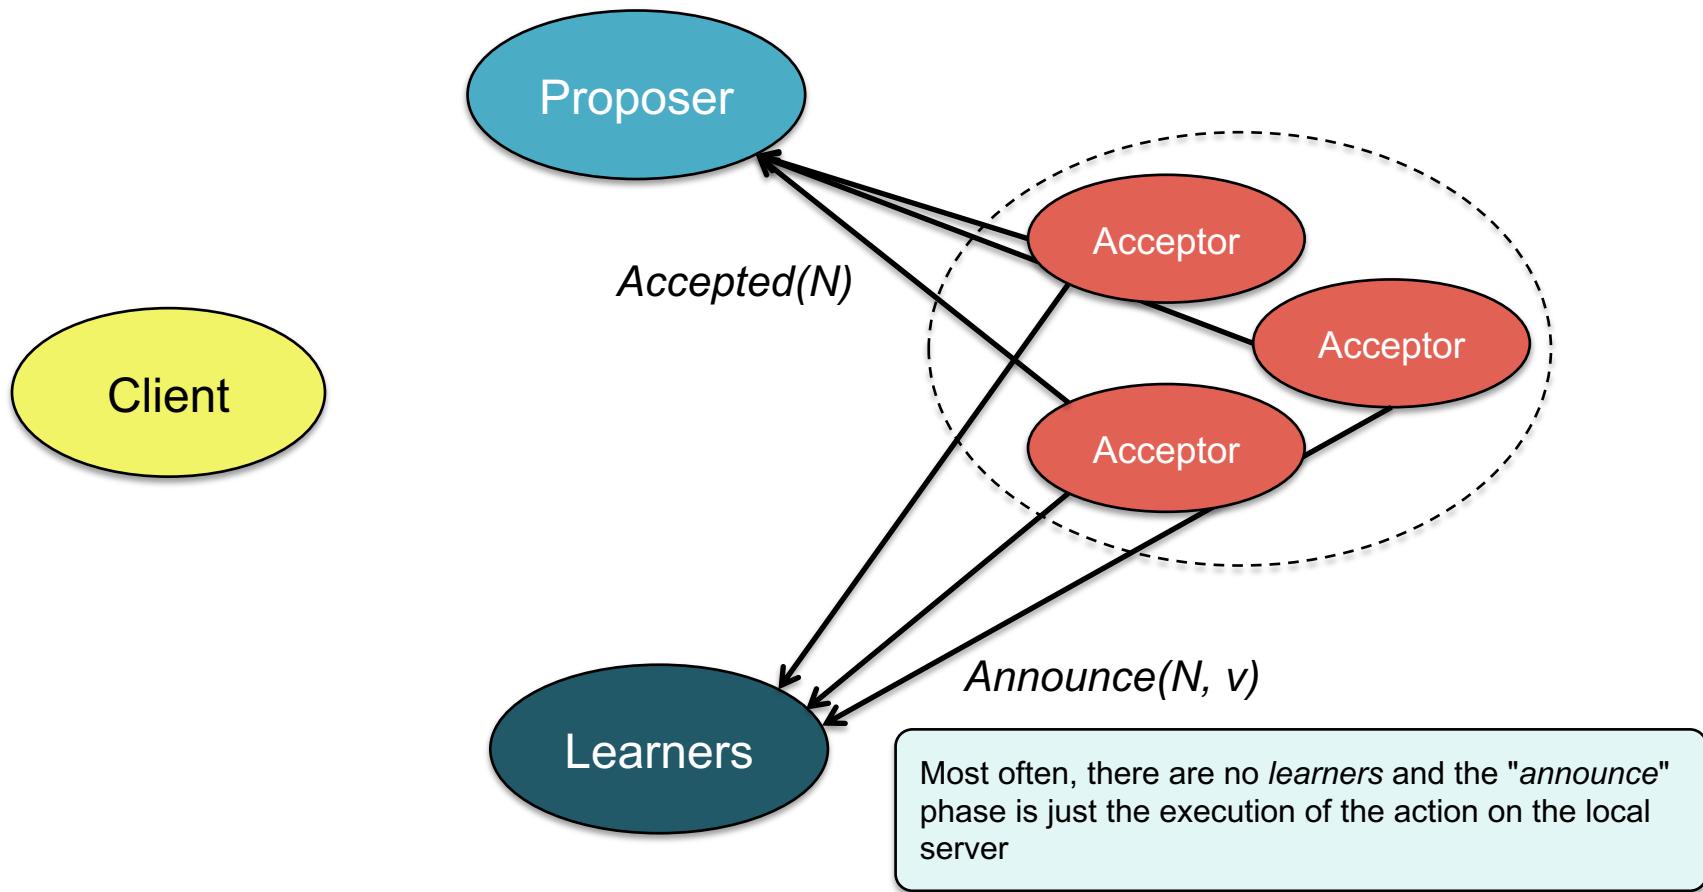


Paxos in action: Phase 2b – ANNOUNCE

Acceptor: if the promise still holds, then announce the value v

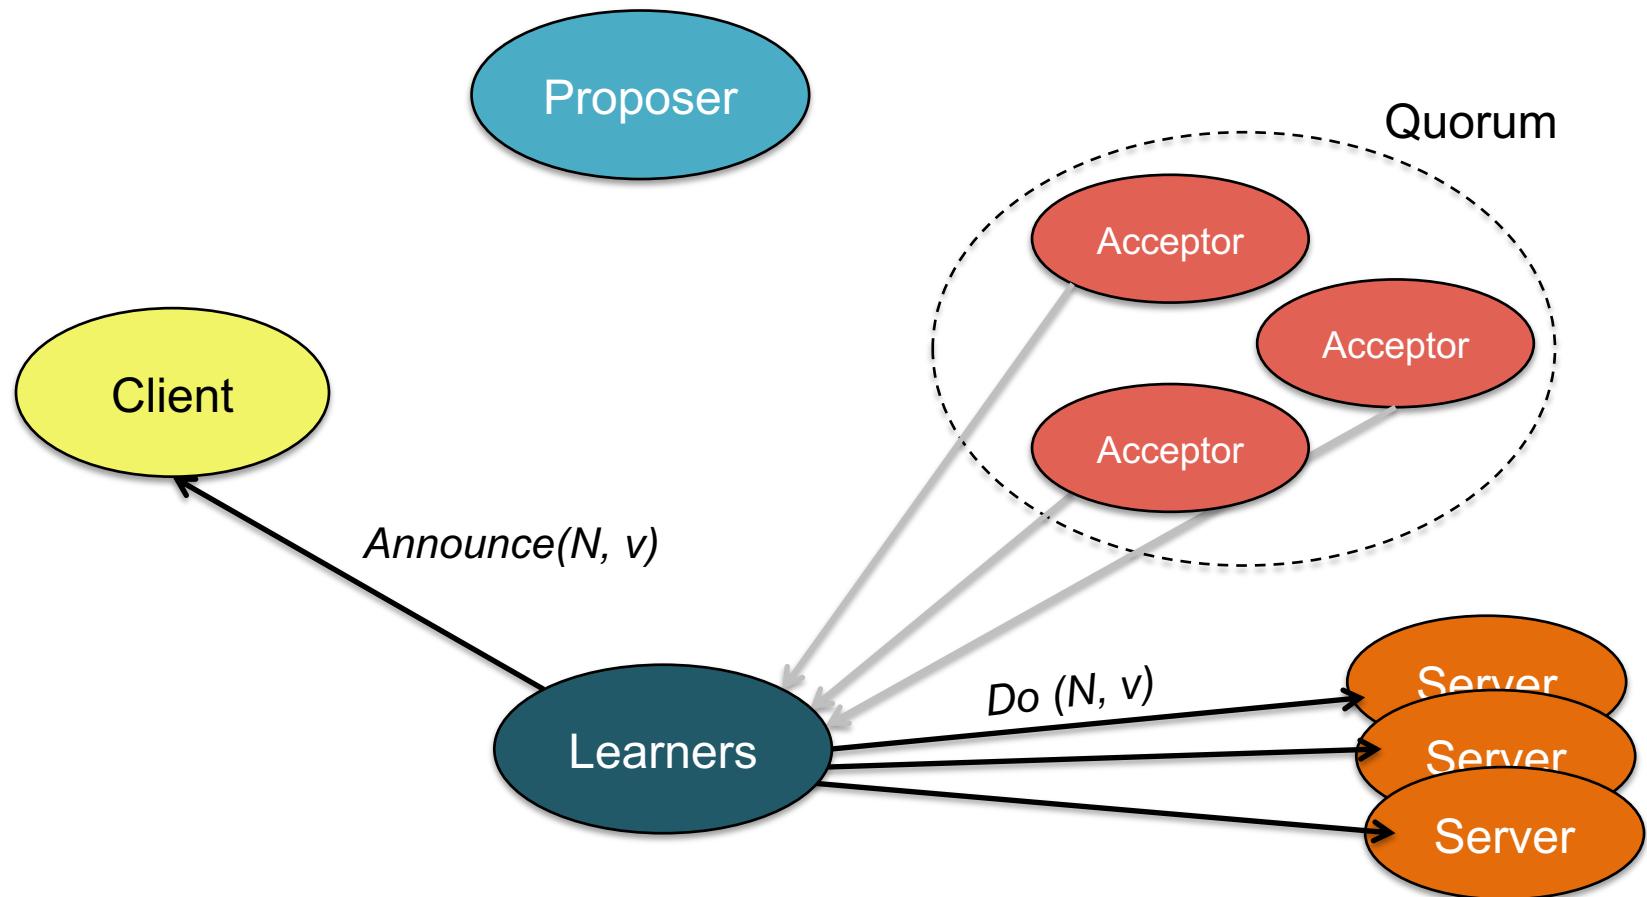
Send **Accepted** message to Proposer and every Learner

BUT: if a higher proposal # may have been received during this time
then send **NACK** to proposer so it can try again



Paxos in action: Phase 3 – ANNOUNCE

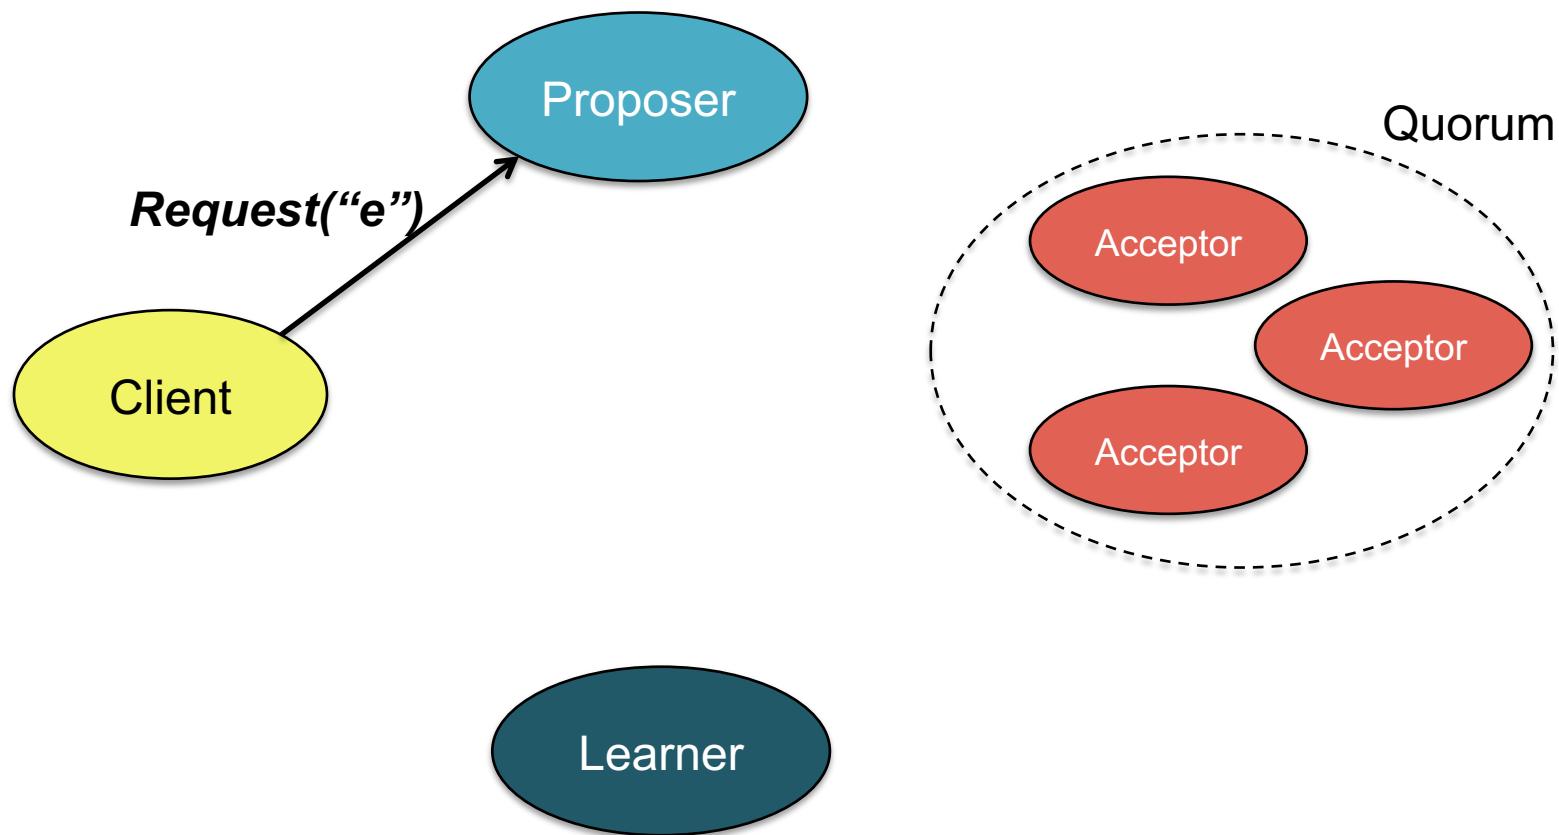
Learner: Respond to client and/or take action on the request



Paxos: A Simple Example – All Good

Paxos in action: Phase 0

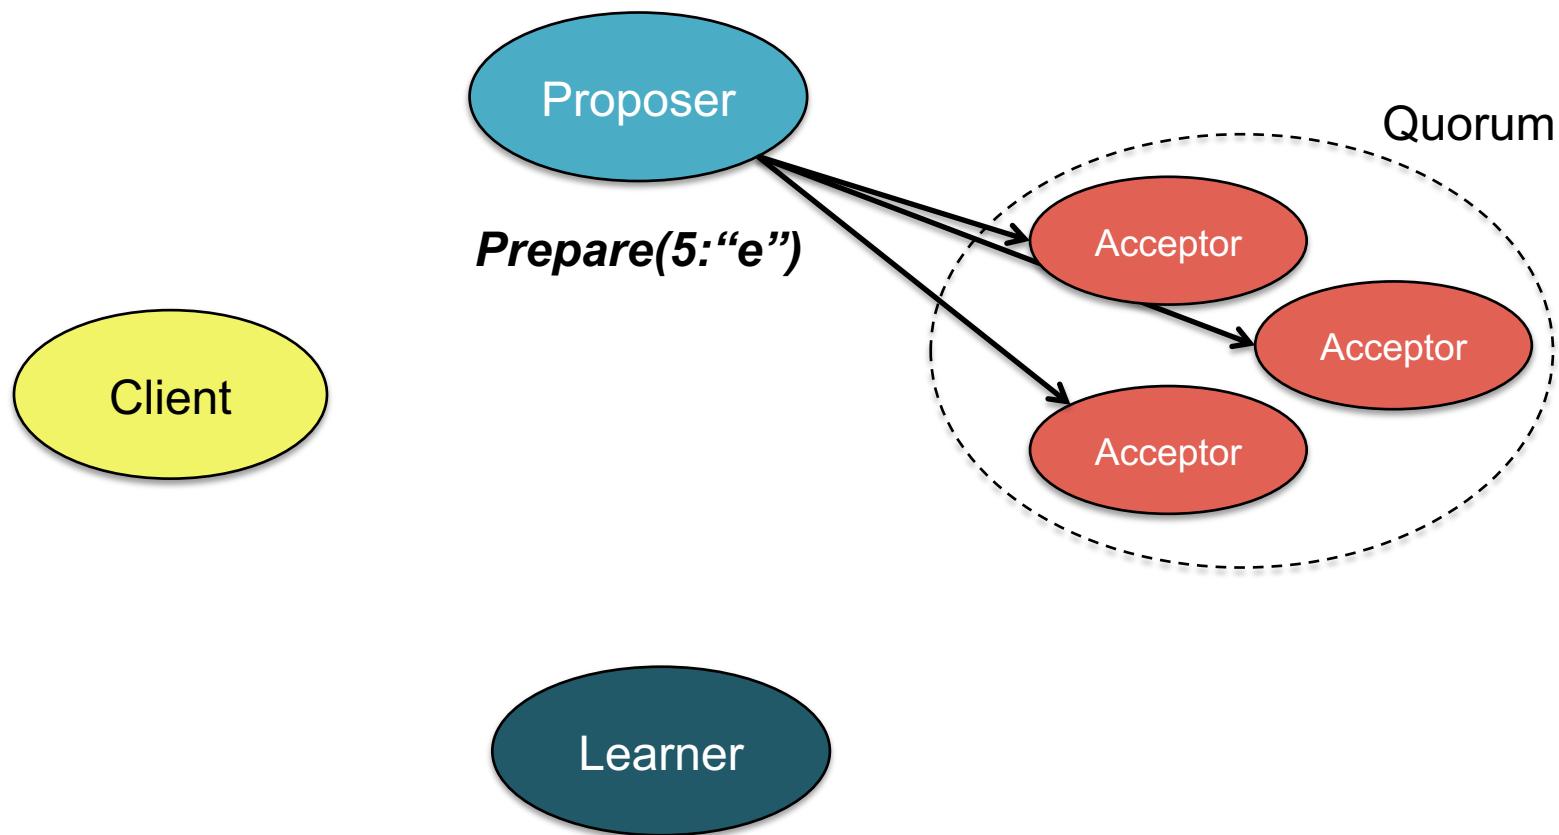
Client sends a request to a proposer



Paxos in action: Phase 1a – *PREPARE*

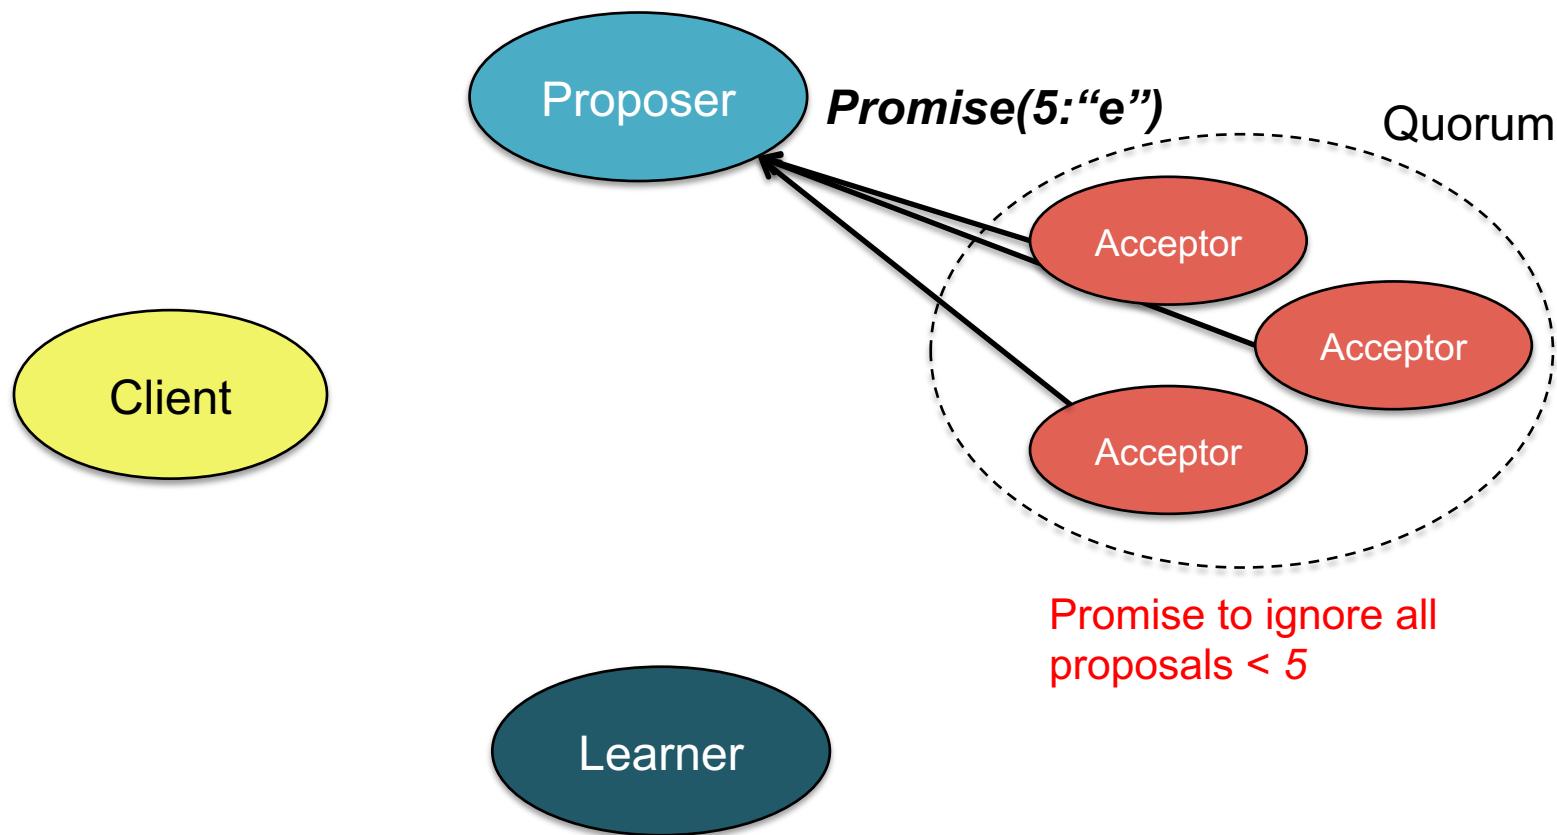
Proposer: picks a sequence number: 5

Send to Quorum of Acceptors



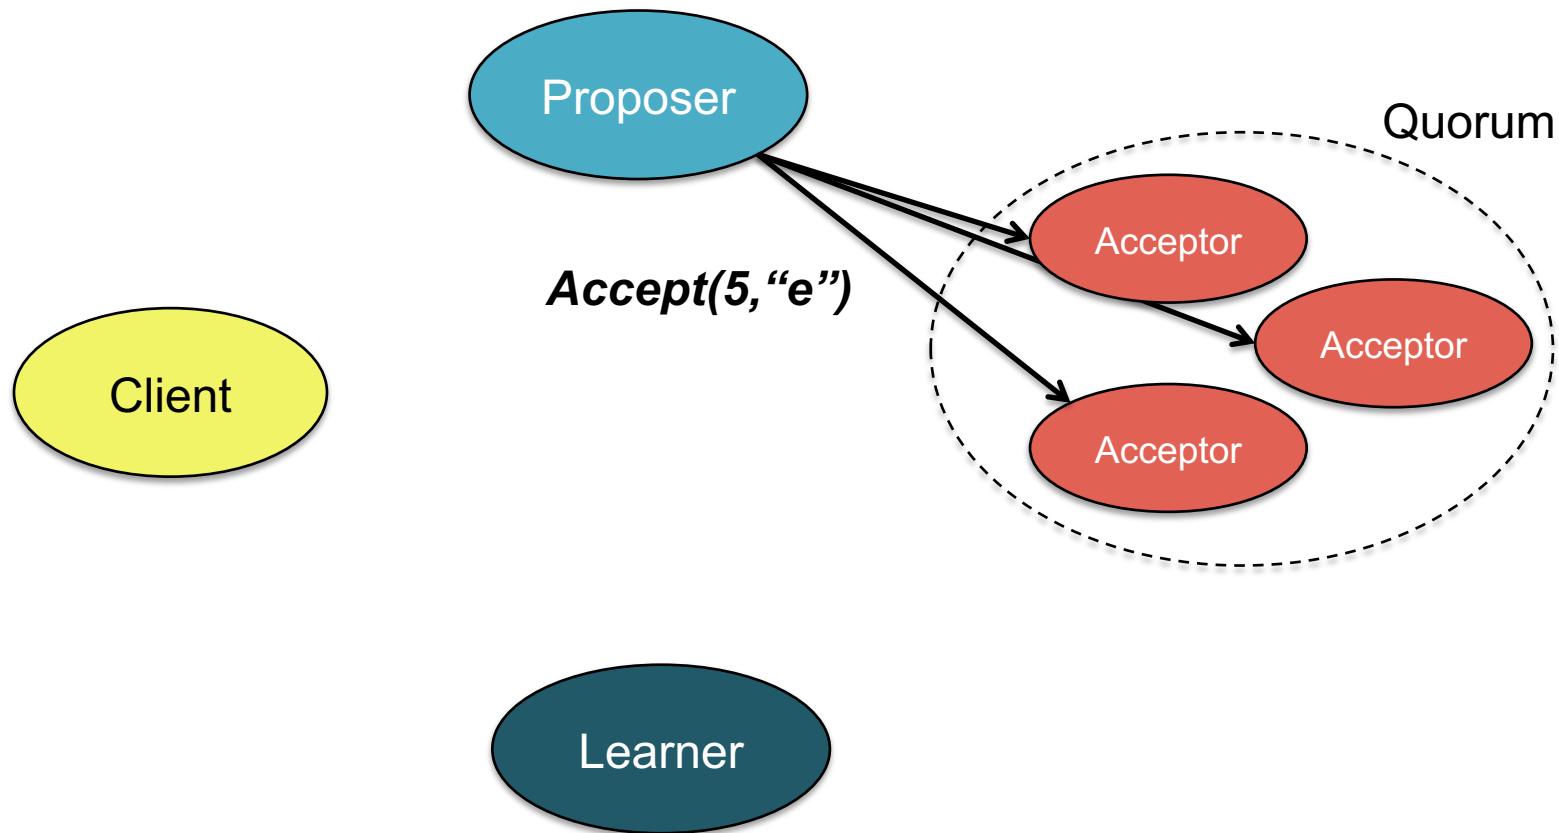
Paxos in action: Phase 1b – PROMISE

Acceptor: Suppose 5 is the highest sequence # any acceptor has seen
Each acceptor PROMISES not to accept any lower numbers



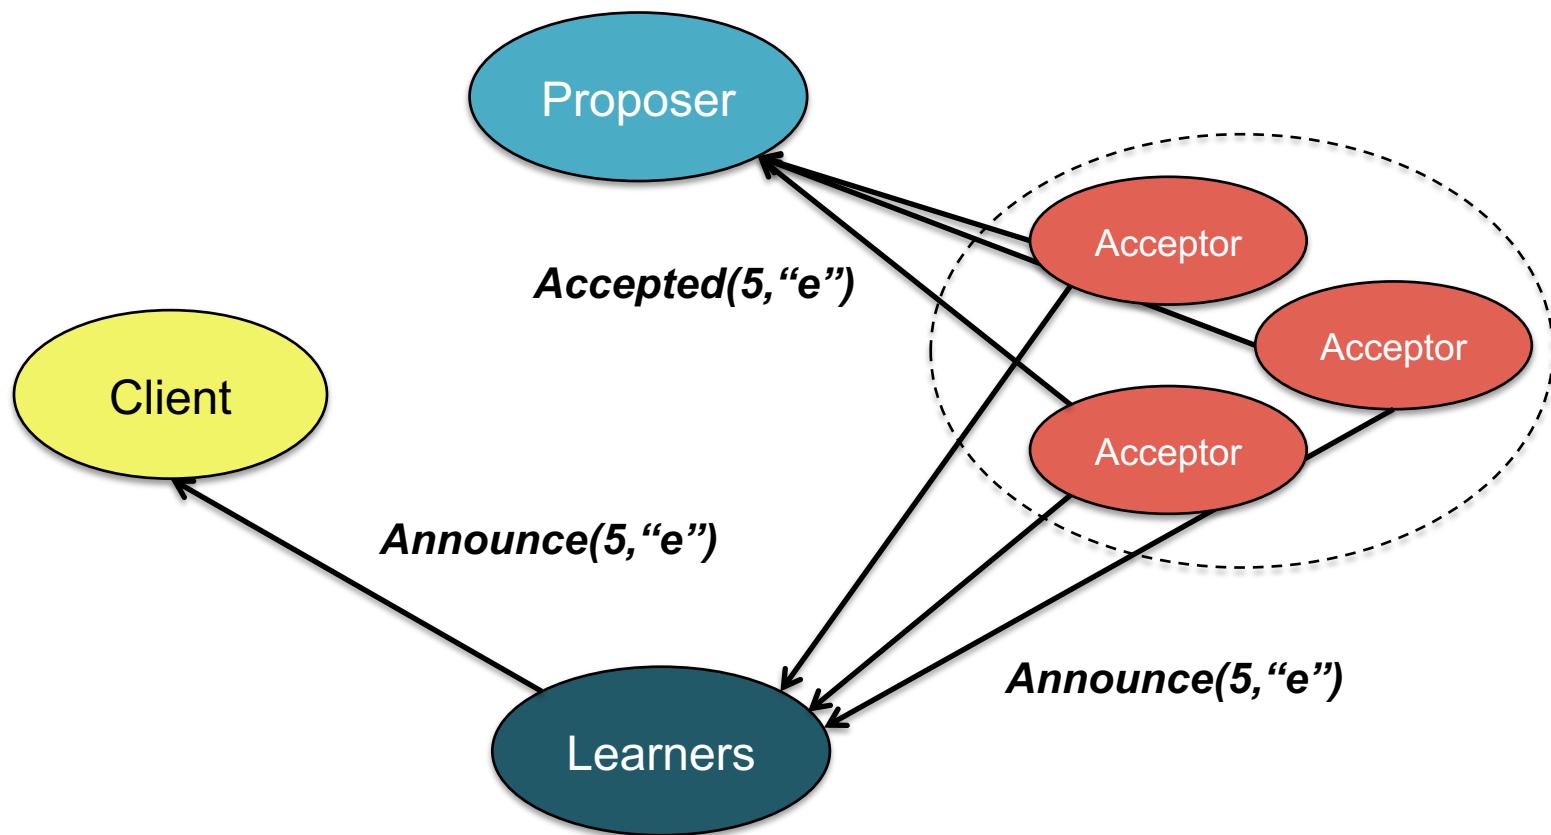
Paxos in action: Phase 2a – ACCEPT

Proposer: Proposer receives the promise from a majority of acceptors
Proposer must accept that $\langle \text{seq}, \text{value} \rangle$



Paxos in action: Phase 2b – ANNOUNCE

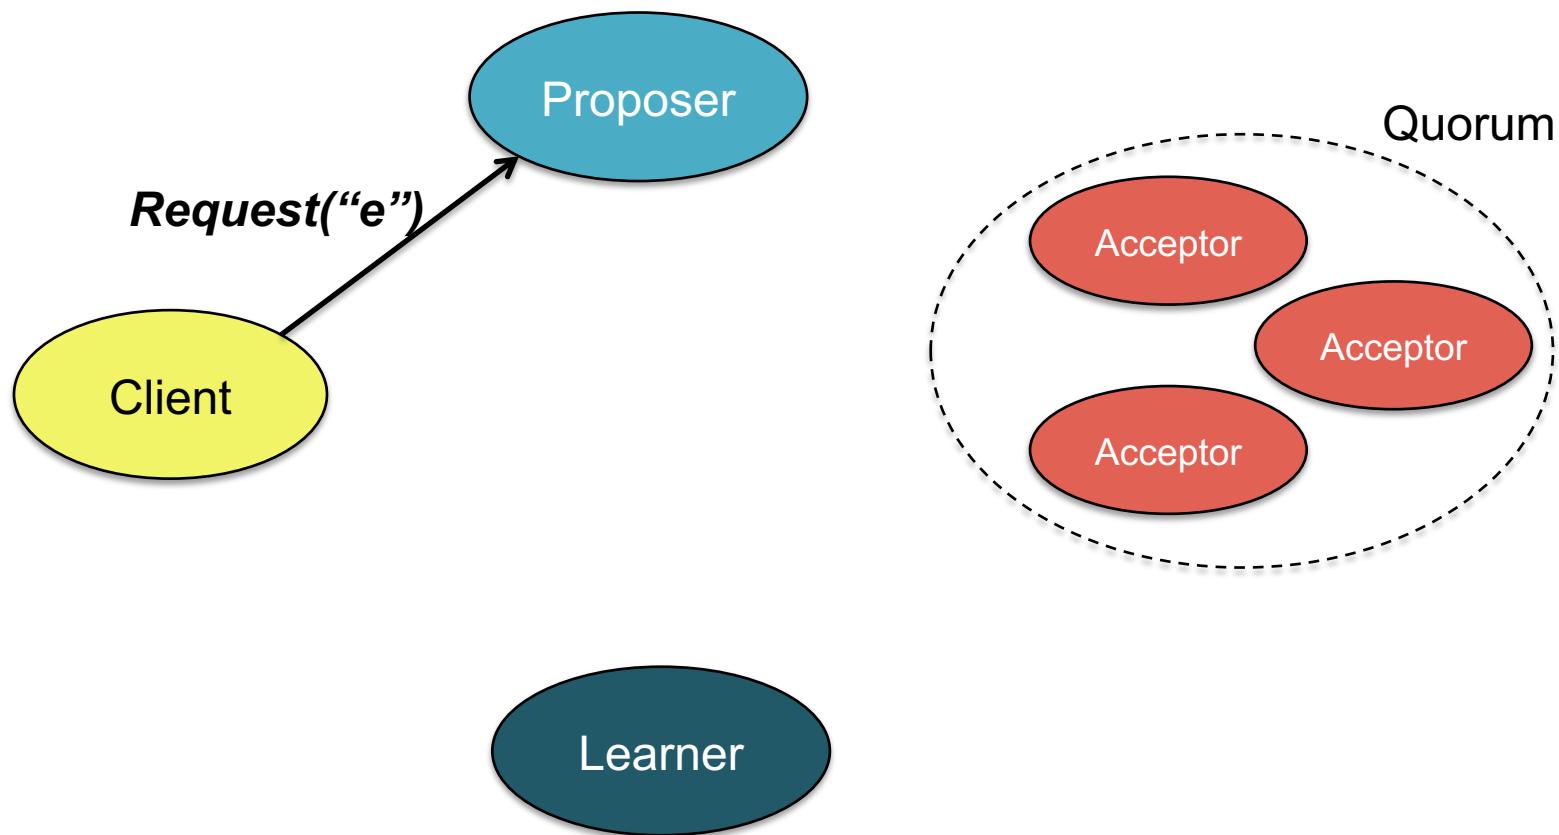
Acceptor: Acceptors state that they accepted the request



Paxos: A Simple Example – Higher Offer

Paxos in action: Phase 0

Client sends a request to a proposer

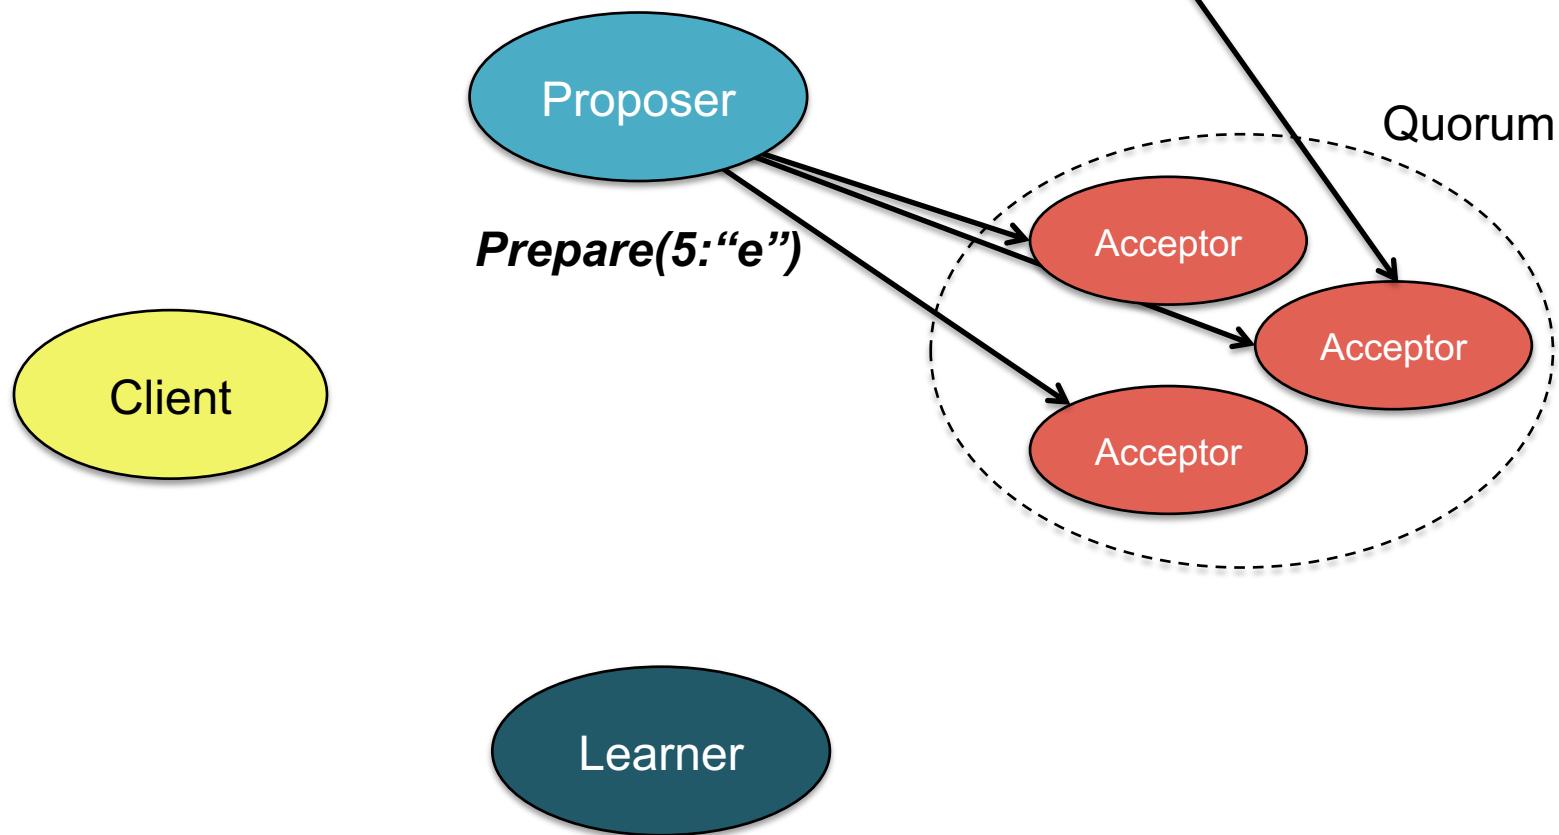


Paxos in action: Phase 1a – *PREPARE*

Proposer: picks a sequence number: 5
Send to Quorum of Acceptors

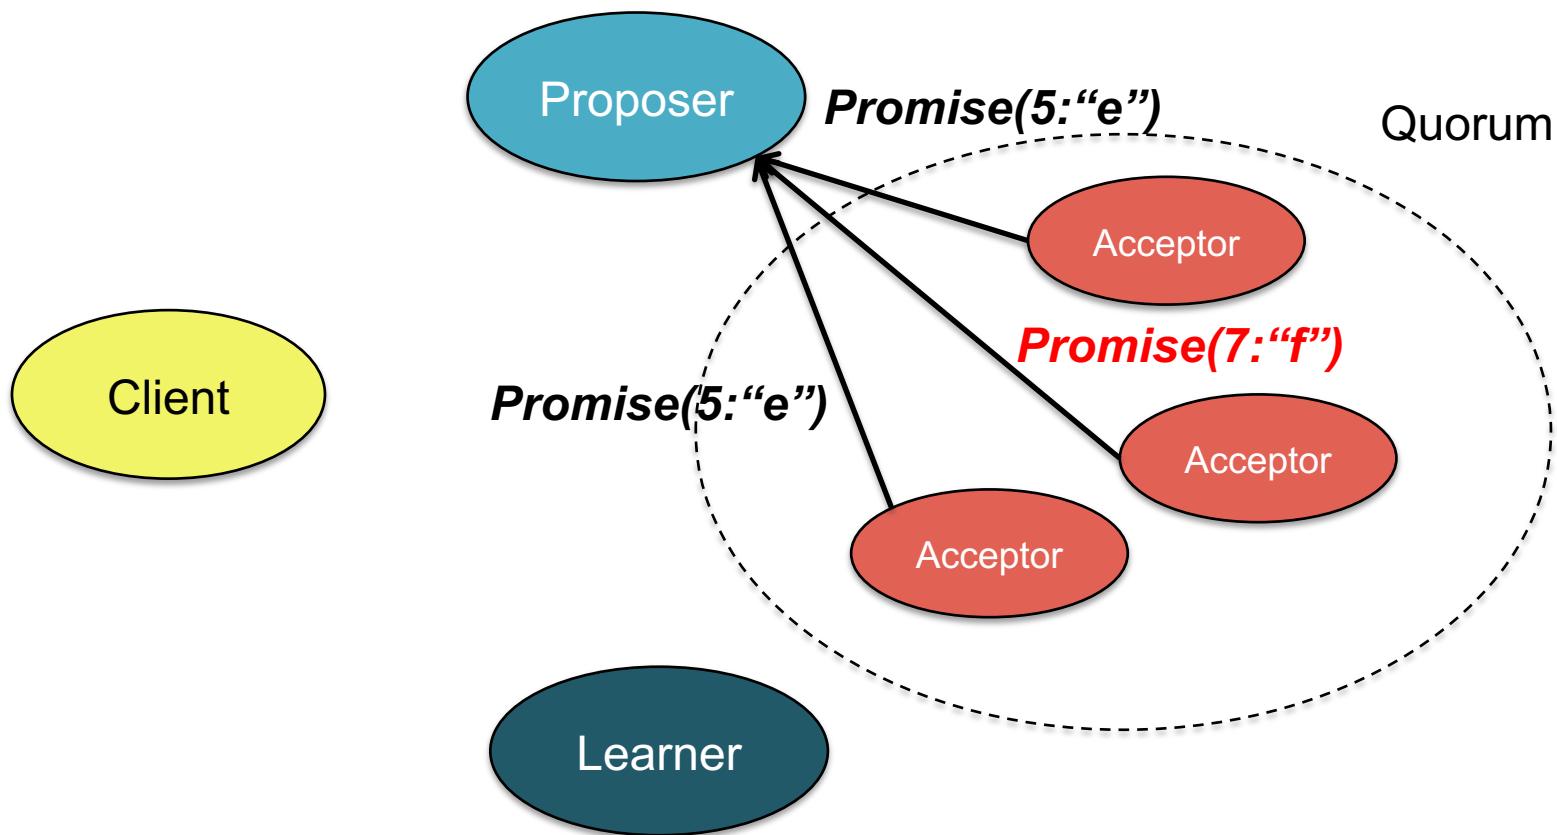
One acceptor receives a higher offer BEFORE it gets this PREPARE message

Prepare(7:“f”)



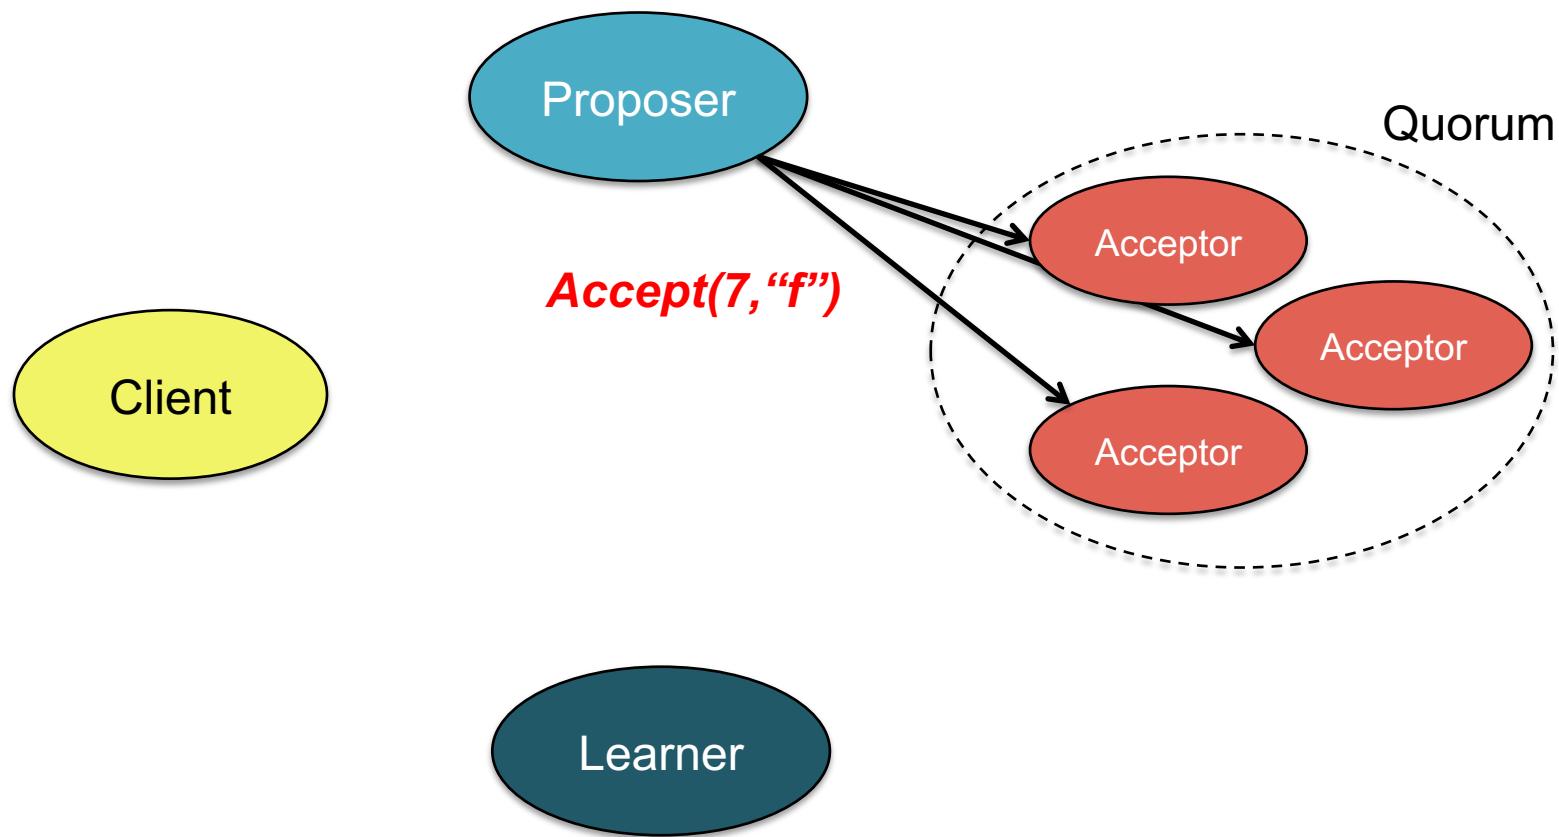
Paxos in action: Phase 1b – PROMISE

Acceptor: Suppose 5 is the highest sequence # any acceptor has seen
Each acceptor PROMISES not to accept any lower numbers



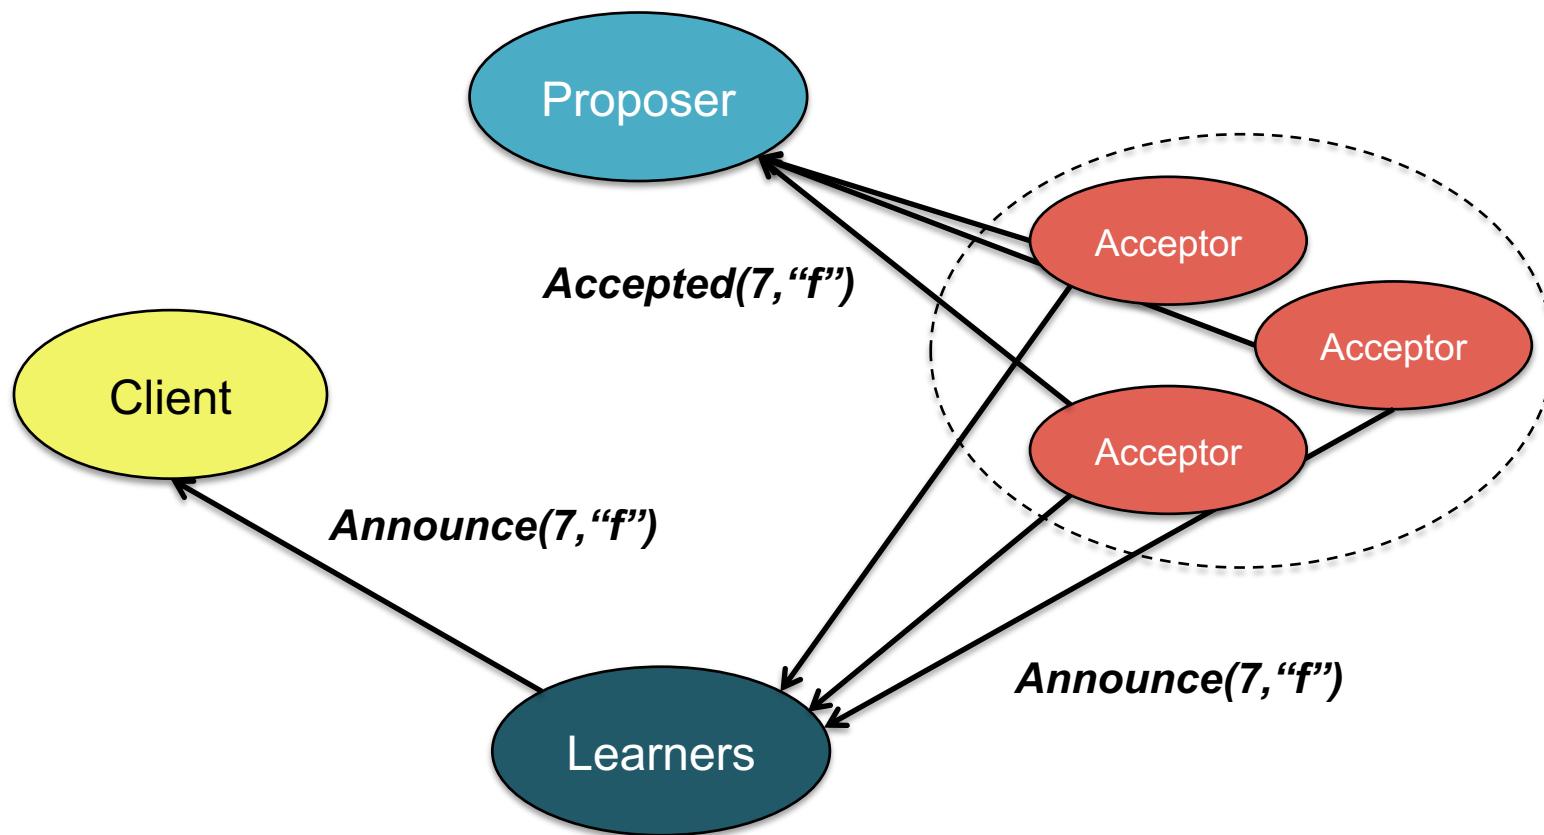
Paxos in action: Phase 2a – ACCEPT

Proposer: Proposer receives the higher # offer and MUST change its mind



Paxos in action: Phase 2b – ANNOUNCE

Acceptor: Acceptors state that they accepted the request



Paxos: Keep trying if you need to

- A proposal N may fail because
 - The acceptor may have made a new promise to ignore all proposals less than some value $M > N$
 - A proposer does not receive a quorum of responses: either *promise* (phase 1b) or *accept* (phase 2b)
- Algorithm then has to be restarted with a higher proposal #

Paxos summary

- Paxos allows us to ensure consistent (total) ordering over a set of events in a group of machines
 - Events = commands, actions, state updates
- Each machine will have the latest state or a previous version of the state
- Paxos used in:
 - Google Chubby lock manager / name server
 - Apache Zookeeper (clone of Google Chubby)
 - Cassandra lightweight transactions
 - Google Spanner, Megastore
 - Microsoft Autopilot cluster management service from Bing
 - VMware NSX Controller
 - Amazon Web Services, DynamoDB

Paxos summary

To make a change to the system:

- Tell the *proposer (leader)* the event/command you want to add
 - Note: these requests may occur concurrently
 - Leader = one elected proposer. Not necessary for Paxos algorithm but an optimization to ensure a single, increasing stream of proposal numbers. Cuts down on rejections and retries.
- The proposer picks its next highest event ID and **asks all the acceptors to reserve that event ID**
 - If any acceptor sees has seen a higher event ID, it rejects the proposal & returns that higher event ID
 - The proposer will have to try again with another event ID
- When the **majority of acceptors accept the proposal**, accepted events are sent to learners, which can act on them (e.g., update system state)
 - Fault tolerant: need $2k+1$ servers for k fault tolerance

Implementation

- Use only one proposer at a time – the leader
 - Other nodes can be active backups just in case the leader dies
 - No need to worry about sync of proposal # – those are local per proposer
 - Acts like a fault-tolerant coordinator
 - Avoids failed proposals due to higher numbers from other proposers
- Alternatively, embed proposer logic into client library
 - Too many clients issuing concurrent requests can cause a large # of retries
- Learners rarely needed
 - Acceptors are often running on the system that processes the request (e.g., data store, log, ...)
 - Just send an acknowledgement directly to the client.

The End