

Distributed Systems

16. Distributed Lookup

Paul Krzyzanowski

Rutgers University

Fall 2017

Distributed Lookup

- Look up (*key, value*)
- Cooperating set of nodes
- Ideally:
 - No central coordinator
 - Some nodes can be down

Approaches

1. Central coordinator
 - Napster
2. Flooding
 - Gnutella
3. Distributed hash tables
 - CAN, Chord, Amazon Dynamo, Tapestry, ...

1. Central Coordinator

- Example: Napster
 - Central directory
 - Identifies content (names) and the servers that host it
 - *lookup(name) → {list of servers}*
 - Download from any of available servers
 - Pick the best one by pinging and comparing response times
- Another example: GFS
 - Controlled environment compared to Napster
 - Content for a given key is broken into chunks
 - Master handles all queries ... but not the data

1. Central Coordinator - Napster

- **Pros**

- Super simple
- Search is handled by a single server (master)
- The directory server is a single point of control
 - Provides definitive answers to a query

- **Cons**

- Master has to maintain state of all peers
- Server gets all the queries
- The directory server is a single point of control
 - No directory, no service!

2. Query Flooding

- Example: Gnutella distributed file sharing
- Well-known nodes act as **anchors**
 - Nodes with files inform an anchor about their existence
 - Nodes select other nodes as peers


2. Query Flooding

- Send a query to peers if a file is not present locally
 - Each request contains:
 - Query key
 - Unique request ID
 - Time to Live (TTL, maximum hop count)
- Peer either responds or routes the query to its neighbors
 - Repeat until TTL = 0 or if the request ID has been processed
 - If found, send response (node address) to the requestor
 - **Back propagation**: response hops back to reach originator

Overlay network

An **overlay network** is a virtual network formed by **peer connections**

- Any node might know about a small set of machines
- “Neighbors” may not be physically close to you




Underlying IP Network

Overlay network


An **overlay network** is a virtual network formed by **peer connections**

- Any node might know about a small set of machines
- “Neighbors” may not be physically close to you




Overlay Network


Flooding Example: Overlay Network




Flooding Example: Query Flood



Flooding Example: Query response



Flooding Example: Download



What's wrong with flooding?

- Some nodes are not always up and some are slower than others
 - Gnutella & Kazaa dealt with this by classifying some nodes as special (“ultrapeers” in Gnutella, “supernodes” in Kazaa,)
- Poor use of network resources
- Potentially high latency
 - Requests get forwarded from one machine to another
 - Back propagation (e.g., in Gnutella's design), where the replies go through the same chain of machines used in the query, increases latency even more

3. Distributed Hash Tables

Locating content

- How do we locate distributed content?
 - A central server is the easiest

Napster	Central server
Gnutella & Kazaa	Network flooding Optimized to flood supernodes ... but it's still flooding
BitTorrent	Nothing! It's somebody else's problem

- Can we do better?

Hash tables

- Remember hash functions & hash tables?
 - Linear search: $O(N)$
 - Tree: $O(\log N)$
 - Hash table: $O(1)$

What's a hash function? (refresher)

- Hash function

- A function that takes a variable length input (e.g., a string) and generates a (usually smaller) fixed length result (e.g., an integer)
- Example: hash strings to a range 0-7:
 - $\text{hash}(\text{"Newark"}) \rightarrow 1$
 - $\text{hash}(\text{"Jersey City"}) \rightarrow 6$
 - $\text{hash}(\text{"Paterson"}) \rightarrow 2$

- Hash table


- Table of *(key, value)* tuples
- Look up a key:
 - Hash function maps *keys* to a range $0 \dots N-1$
table of N elements
`i = hash(key)`
`table[i]` contains the item
- No need to search through the table!

Considerations with hash tables (refresher)

- Picking a good hash function
 - We want uniform distribution of all values of *key* over the space $0 \dots N-1$
- Collisions
 - Multiple keys may hash to the same value
 - `hash("Paterson")` → 2
 - `hash("Edison")` → 2
 - `table[i]` is a **bucket** (**slot**) for all such (*key*, *value*) sets
 - Within `table[i]`, use a linked list or another layer of hashing
- Think about a hash table that grows or shrinks
 - If we add or remove buckets → need to rehash keys and move items

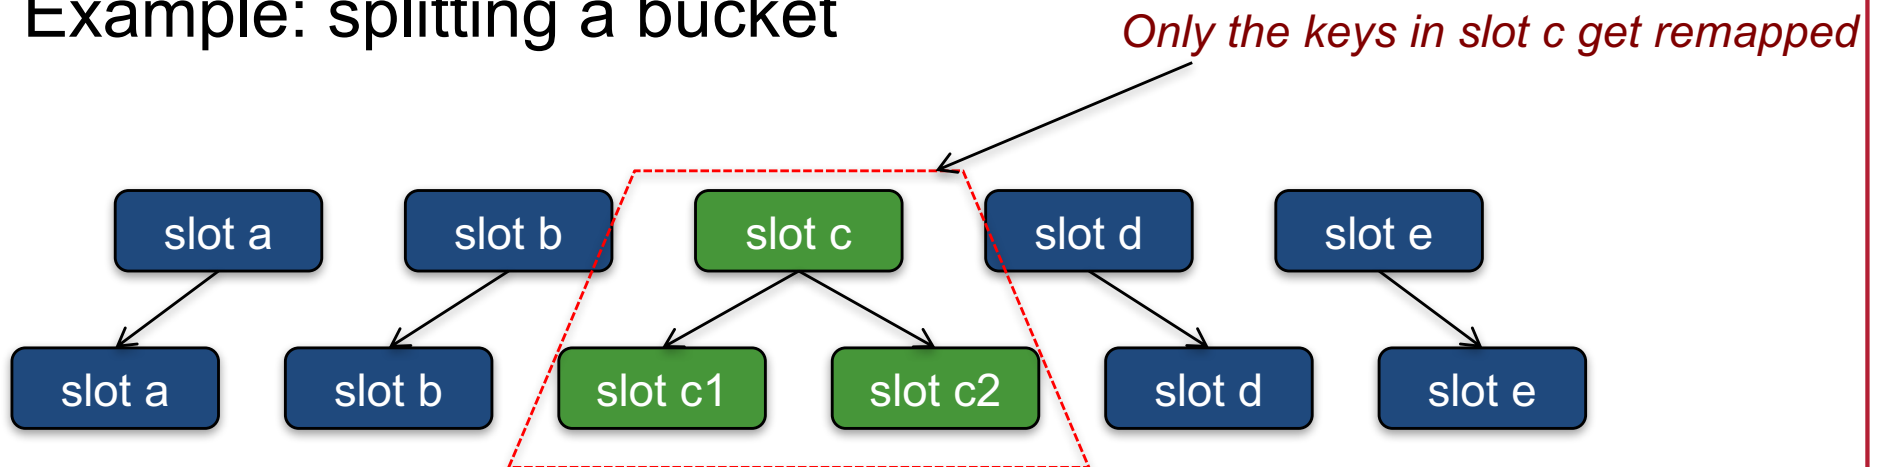
Distributed Hash Tables (DHT)

- Create a peer-to-peer version of a (*key*, *value*) data store
- How we want it to work
 1. A peer (*A*) queries the data store with a key
 2. The data store finds the peer (*B*) that has the value
 3. That peer (*B*) returns the (*key*, *value*) pair to the querying peer (*A*)
- Make it efficient!
 - A query should not generate a flood!



Consistent hashing

- Conventional hashing
 - Practically all keys have to be remapped if the table size changes
- Consistent hashing
 - Most keys will hash to the same value as before
 - On average, K/n keys will need to be remapped
 $K = \# \text{ keys}, n = \# \text{ of buckets}$
- Example: splitting a bucket



3. Distributed hashing

- Spread the hash table across multiple nodes
- Each node stores a portion of the key space

lookup(key) → node ID that holds (*key*, *value*)

lookup(node_ID, key) → value

Questions

How do we partition the data & do the lookup?

& keep the system decentralized?

& make the system scalable (lots of nodes with dynamic changes)?

& fault tolerant (replicated data)?


Distributed Hashing Case Study

CAN: Content Addressable Network

CAN design

- Create a logical grid
 - x-y in 2-D but not limited to 2-D
- Separate hash function per dimension
 - $h_x(\text{key})$, $h_y(\text{key})$
- A node
 - Is responsible for a range of values in each dimension
 - Knows its neighboring nodes

CAN *key*→*node* mapping: 2 nodes



$x = \text{hash}_x(\text{key})$


$y = \text{hash}_y(\text{key})$

if $x < (x_{\max}/2)$
 n_1 has (*key*, *value*)

if $x \geq (x_{\max}/2)$
 n_2 has (*key*, *value*)


n_2 is responsible for a **zone**
 $x=(x_{\max}/2 \dots x_{\max})$,
 $y=(0 \dots y_{\max})$

CAN partitioning



Any node can be split in two – either horizontally or vertically

CAN key→node mapping




$x = \text{hash}_x(\text{key})$

$y = \text{hash}_y(\text{key})$

```
if  $x < (x_{\max}/2)$  {  
  if  $y < (y_{\max}/2)$   
     $n_0$  has (key, value)  
  else  
     $n_1$  has (key, value)  
}
```

```
if  $x \geq (x_{\max}/2)$   
   $n_2$  has (key, value)
```

CAN partitioning




Any node can be split in two – either horizontally or vertically

Associated data has to be moved to the new node based on *hash(key)*

Neighbors need to be made aware of the new node

A node knows only of its **neighbors**


CAN neighbors



Neighbors refer to nodes that share **adjacent zones** in the **overlay network**

n_4 only needs to keep track of n_5 , n_7 , or n_8 as its right neighbor.

CAN routing



lookup(key) on a node that does not own the value

Compute $hash_x(key)$, $hash_y(key)$ and route request to a neighboring node

Ideally: route to minimize distance to destination

CAN


- Performance
 - For n nodes in d dimensions
 - # neighbors = $2d$
 - Average route for 2 dimensions = $O(\sqrt{n})$ hops
- To handle failures
 - Share knowledge of neighbor's neighbors
 - One of the node's neighbors takes over the failed zone

Distributed Hashing Case Study

Chord

Chord & consistent hashing

- A key is hashed to an m -bit value: $0 \dots (2^m-1)$
- A logical ring is constructed for the values $0 \dots (2^m-1)$
- Nodes are placed on the ring at *hash(IP address)*



Key assignment

- Example: $n=16$; system with 4 nodes (so far)
- Key, value data is stored at a **successor**
 - a node whose value is $\geq \text{hash}(\text{key})$


Node 14 is responsible for keys 11, 12, 13, 14

No nodes at these empty positions

Node 3 is responsible for keys 15, 0, 1, 2, 3

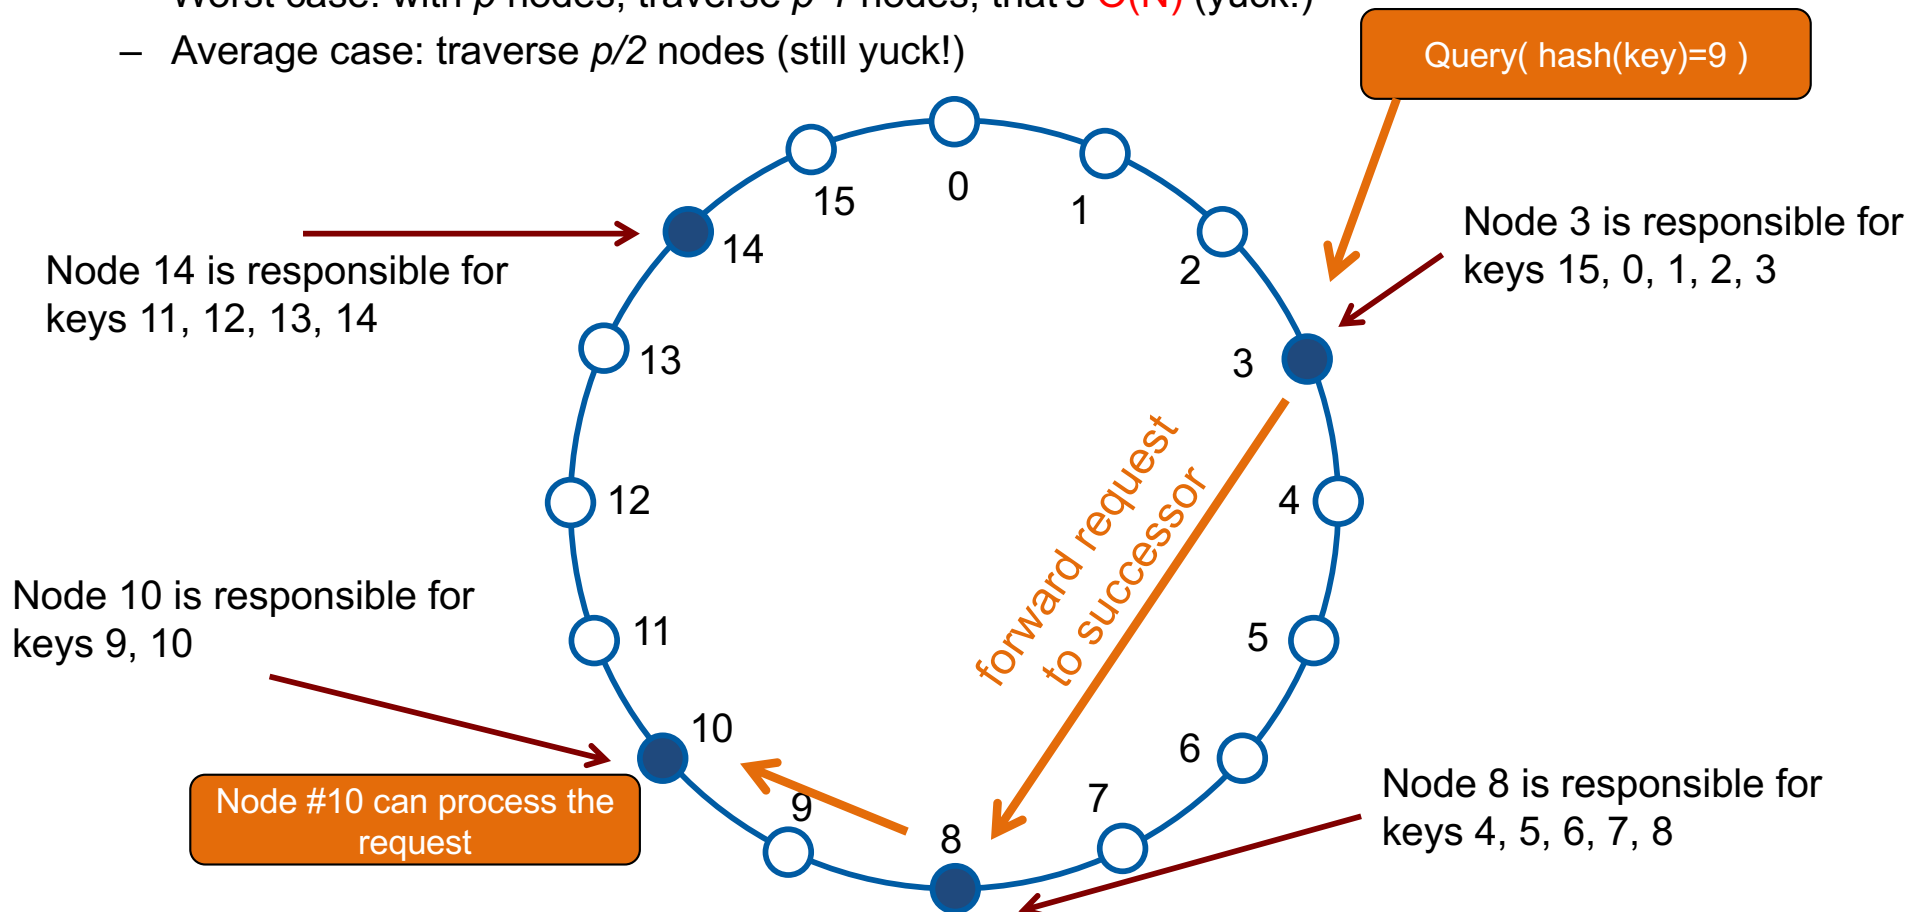
Node 10 is responsible for keys 9, 10

Node 8 is responsible for keys 4, 5, 6, 7, 8



Handling query requests

- Any peer can get a request (*insert* or *query*). If the *hash(key)* is not for its ranges of keys, it forwards the request to a successor.
 - The process continues until the responsible node is found
 - Worst case: with p nodes, traverse $p-1$ nodes; that's $O(N)$ (yuck!)
 - Average case: traverse $p/2$ nodes (still yuck!)
- Query(hash(key)=9)




Let's figure out three more things

1. Adding/removing nodes
2. Improving lookup time
3. Fault tolerance

Adding a node

- Some keys that were assigned to a node's successor now get assigned to the new node
- Data for those *(key, value)* pairs must be moved to the new node




Removing a node

- Keys are reassigned to the node's successor
- Data for those *(key, value)* pairs must be moved to the successor

Node 14 was responsible for keys 11, 12, 13, 14

Node 14 is now responsible for keys 9, 10, 11, 12, 13, 14



Fault tolerance

- Nodes might die
 - *(key, value)* data should be replicated
 - Create R replicas, storing each one at $R-1$ successor nodes in the ring
- Need to know multiple successors
 - A node needs to know how to find its successor's successor (or more)
 - Easy if it knows all nodes!
 - When a node is back up, it needs to check with successors for updates
 - Any changes need to be propagated to all replicas

Performance

- We're not thrilled about $O(N)$ lookup
- Simple approach for great performance
 - Have all nodes know about each other
 - When a peer gets a node, it searches its table of nodes for the node that owns those values
 - Gives us $O(1)$ performance
 - Add/remove node operations must inform everyone
 - Maybe not a good solution if we have millions of peers (huge tables)

Finger tables

- Compromise to avoid large tables at each node
 - Use **finger tables** to place an upper bound on the table size
- **Finger table** = partial list of nodes, progressively more distant
- At each node, i^{th} entry in finger table identifies node that succeeds it by at least 2^{i-1} in the circle
 - `finger_table[0]`: immediate (1st) successor
 - `finger_table[1]`: successor after that (2nd)
 - `finger_table[2]`: 4th successor
 - `finger_table[3]`: 8th successor
 - ...
- $O(\log N)$ nodes need to be contacted to find the node that owns a key
... not as cool as $O(1)$ but way better than $O(N)$

Improving performance even more

- Let's revisit $O(1)$ lookup
- Each node keeps track of all current nodes in the group
 - Is that really so bad?
 - We might have thousands of nodes ... so what?
- Any node will now know which node holds a *(key, value)*
- Add or remove a node: send updates to **all** other nodes

Distributed Hashing Case Study

Amazon Dynamo

Amazon Dynamo

- Not exposed as a web service
 - Used to power parts of Amazon Web Services (such as S3)
 - Highly available, key-value storage system
- In an infrastructure with millions of components, something is always failing!
 - Failure is the normal case
- A lot of services within Amazon only need primary-key access to data
 - Best seller lists, shopping carts, preferences, session management, sales rank, product catalog
 - No need for complex querying or management offered by an RDBMS
 - Full relational database is overkill: limits scale and availability
 - Still not efficient to scale or load balance RDBMS on a large scale

Core Assumptions & Design Decisions

- Two operations: **get(key)** and **put(key, data)**
 - Binary objects (data) identified by a unique key
 - Objects tend to be small (< 1MB)
- ACID gives poor availability
 - Use weaker consistency (C) for higher availability.
- Apps should be able to configure Dynamo for desired latency & throughput
 - Balance performance, cost, availability, durability guarantees.
- At least 99.9% of read/write operations must be performed within a few hundred milliseconds:
 - **Avoid routing requests through multiple nodes**
- Dynamo can be thought of as a **zero-hop DHT**

Core Assumptions & Design Decisions

- **Incremental scalability**
 - System should be able to grow by adding a storage host (node) at a time
- **Symmetry**
 - Every node has the same set of responsibilities
- **Decentralization**
 - Favor decentralized techniques over central coordinators
- **Heterogeneity (mix of slow and fast systems)**
 - Workload partitioning should be proportional to capabilities of servers

Consistency & Availability

- Strong consistency & high availability cannot be achieved simultaneously
- Optimistic replication techniques – eventually consistent model
 - propagate changes to replicas in the background
 - can lead to conflicting changes that have to be detected & resolved
- When do you resolve conflicts?
 - **During writes**: traditional approach – reject write if cannot reach all (or majority) of replicas – *but don't deal with conflicts*
 - **Resolve conflicts during reads**: Dynamo approach
 - Design for an "always writable" data store - highly available
 - read/write operations can continue even during network partitions
 - Rejecting customer updates won't be a good experience
 - A customer should always be able to add or remove items in a shopping cart

Consistency & Availability

- Who resolves conflicts?
 - **Choices**: the data store system or the application?
- **Data store**
 - Application-unaware, so choices limited
 - Simple policy, such as "last write wins"
- **Application**
 - App is aware of the meaning of the data
 - Can do application-aware conflict resolution
 - E.g., merge shopping cart versions to get a unified shopping cart.
- Fall back on "*last write wins*" if app doesn't want to bother

Reads & Writes

Two operations:

- **get(key)** returns
 1. **object** or list of objects with conflicting versions
 2. **context** (resultant version per object)
- **put(key, context, value)**
 - stores replicas
 - *context*: ignored by the application but includes version of object
 - key is hashed with MD5 to create a 128-bit identifier that is used to determine the storage nodes that serve the key
 - hash(key) identifies node*

Partitioning the data


- Break up database into chunks distributed over all nodes
 - Key to scalability
- Relies on **consistent hashing**
 - Regular hashing: change in # slots requires all keys to be remapped
 - Consistent hashing:
 - K/n keys need to be remapped, $K = \# \text{ keys}$, $n = \# \text{ slots}$
- **Logical ring of nodes: just like Chord**
 - Each node assigned a random value in the hash space: position in ring
 - Responsible for all hash values between its value and predecessor's value
 - Hash(key); then walk ring clockwise to find first node with position > hash
 - Adding/removing nodes affects only immediate neighbors

Partitioning: virtual nodes

- A node is assigned to **multiple points** in the ring
- Each point is a “**virtual node**”

Dynamo virtual nodes

- A physical node holds contents of multiple virtual nodes
- In this example: 2 physical nodes, 5 virtual nodes



Partitioning: virtual nodes

Advantage: **balanced load distribution**

- If a node becomes unavailable, load is evenly dispersed among available nodes
- If a node is added, it accepts an equivalent amount of load from other available nodes
- # of virtual nodes per system can be based on the capacity of that node
 - Makes it easy to support changing technology and addition of new, faster systems

Replication

- Data replicated on N hosts (N is configurable)
 - Key is assigned a **coordinator** node (via hashing) = main node
 - Coordinator is in charge of replication
- Coordinator replicates keys at the $N-1$ clockwise successor nodes in the ring

Dynamo Replication


Coordinator replicates keys at the $N-1$ clockwise successor nodes in the ring

Example: $N=3$

Node 14 holds replicas for
Nodes 1 and 3

Node 10 holds replicas for
Node 14 and 1

Node 8 holds replicas for
Nodes 10 and 14



Versioning

- Not all updates may arrive at all replicas
 - Clients may modify or read stale data
- Application-based reconciliation
 - Each modification of data is treated as a new version
- Vector clocks are used for versioning
 - Capture causality between different versions of the same object
 - Vector clock is a set of (node, counter) pairs
 - Returned as a context from a get() operation

Availability

- **Configurable values**
 - R : minimum # of nodes that must participate in a successful read operation
 - W : minimum # of nodes that must participate in a successful write operation
- **Metadata hints to remember original destination**
 - If a node was unreachable, the replica is sent to another node in the ring
 - Metadata sent with the data contains a hint stating the original desired destination
 - Periodically, a node checks if the originally targeted node is alive
 - if so, it will transfer the object and may delete it locally to keep # of replicas in the system consistent
- **Data center failure**
 - System must handle the failure of a data center
 - Each object is replicated across multiple data centers

Storage Nodes

Each node has three components

1. *Request coordination*

- Coordinator executes read/write requests on behalf of requesting clients
- State machine contains all logic for identifying nodes responsible for a key, sending requests, waiting for responses, retries, processing retries, packaging response
- Each state machine instance handles one request

2. Membership and failure detection

3. Local persistent storage

- Different storage engines may be used depending on application needs
 - Berkeley Database (BDB) Transactional Data Store (most popular)
 - BDB Java Edition
 - MySQL (for large objects)
 - in-memory buffer with persistent backing store

Amazon S3 (Simple Storage Service)

Commercial service that implements many of Dynamo's features

- Storage via web services interfaces (REST, SOAP, BitTorrent)
 - Stores more than 449 billion objects
 - 99.9% uptime guarantee (43 minutes downtime per month)
 - Proprietary design
 - Stores arbitrary objects up to 5 TB in size
- Objects organized into buckets and within a bucket identified by a unique user-assigned key
- Buckets & objects can be created, listed, and retrieved via REST or SOAP
 - <http://s3.amazonaws.com/bucket/key>
- objects can be downloaded via HTTP GET or BitTorrent protocol
 - S3 acts as a seed host and any BitTorrent client can retrieve the file
 - reduces bandwidth costs
- S3 can also host static websites

The end