

Java Concurrency

1. [Java Concurrency / Multithreading Tutorial](#)
2. [Multithreading Benefits](#)
3. [Multithreading Costs](#)
4. **[Concurrency Models](#)**
5. [Same-threading](#)
6. [Concurrency vs. Parallelism](#)
7. [Creating and Starting Java Threads](#)
8. [Race Conditions and Critical Sections](#)
9. [Thread Safety and Shared Resources](#)
10. [Thread Safety and Immutability](#)
11. [Java Memory Model](#)
12. [Java Synchronized Blocks](#)
13. [Java Volatile Keyword](#)
14. [Java ThreadLocal](#)
15. [Thread Signaling](#)
16. [Deadlock](#)
17. [Deadlock Prevention](#)
18. [Starvation and Fairness](#)
19. [Nested Monitor Lockout](#)
20. [Slipped Conditions](#)
21. [Locks in Java](#)
22. [Read / Write Locks in Java](#)
23. [Reentrance Lockout](#)
24. [Semaphores](#)
25. [Blocking Queues](#)
26. [Thread Pools](#)
27. [Compare and Swap](#)
28. [Anatomy of a Synchronizer](#)
29. [Non-blocking Algorithms](#)
30. [Amdahl's Law](#)
31. [Java Concurrency References](#)

Concurrency Models

- [Concurrency Models and Distributed System Similarities](#)
- [Parallel Workers](#)
- [Parallel Workers Advantages](#)
- [Parallel Workers Disadvantages](#)
 - [Shared State Can Get Complex](#)
 - [Stateless Workers](#)
 - [Job Ordering is Nondeterministic](#)
- [Assembly Line](#)
 - [Reactive, Event Driven Systems](#)
 - [Actors vs. Channels](#)
- [Assembly Line Advantages](#)
 - [No Shared State](#)
 - [Stateful Workers](#)
 - [Better Hardware Conformity](#)
 - [Job Ordering is Possible](#)
- [Assembly Line Disadvantages](#)
- [Functional Parallelism](#)
- [Which Concurrency Model is Best?](#)

Jakob Jenkov
Last update: 2015-04-13



Concurrent systems can be implemented using different concurrency models. A *concurrency model* specifies how threads in the the system collaborate to complete the jobs they are are given. Differ concurrency models split the jobs in different ways, and the threads may communicate and collab different ways. This concurrency model tutorial will dive a bit deeper into the most popular concurr models in use at the time of writing (2015).

Concurrency Models and Distributed System Similarities


The concurrency models described in this text are similar to different architectures used in distribu systems. In a concurrent system different threads communicate with each other. In a distributed s different processes communicate with each other (possibly on different computers). Threads and processes are quite similar to each other in nature. That is why the different concurrency models look similar to different distributed system architectures.

Of course distributed systems have the extra challenge that the network may fail, or a remote con process is down etc. But a concurrent system running on a big server may experience similar probl a CPU fails, a network card fails, a disk fails etc. The probability of failure may be lower, but it can theoretically still happen.

Because concurrency models are similar to distributed system architectures, they can often borro from each other. For instance, models for distributing work among workers (threads) are often sim models of **load balancing in distributed systems**. The same is true of error handling techniques logging, fail-over, idempotency of jobs etc.

Parallel Workers

The first concurrency model is what I call the *parallel worker* model. Incoming jobs are assigned to different workers. Here is a diagram illustrating the parallel worker concurrency model:



In the parallel worker concurrency model a delegator distributes the incoming jobs to different workers. Each worker completes the full job. The workers work in parallel, running in different threads, and on different CPUs.

If the parallel worker model was implemented in a car factory, each car would be produced by one worker. The worker would get the specification of the car to build, and would build everything from start to finish.

The parallel worker concurrency model is the most commonly used concurrency model in Java applications (although that is changing). Many of the concurrency utilities in the [java.util.concurrent](#) package are designed for use with this model. You can also see traces of this model in the design of Java Enterprise Edition application servers.

Parallel Workers Advantages

The advantage of the parallel worker concurrency model is that it is easy to understand. To increase parallelization of the application you just add more workers.

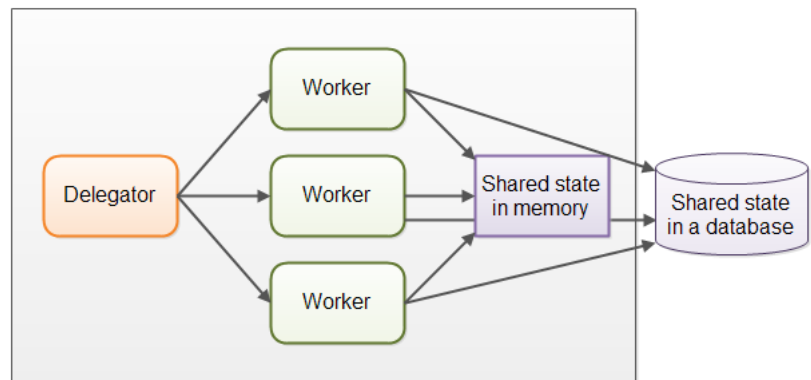
For instance, if you were implementing a web crawler, you could crawl a certain amount of pages with different numbers of workers and see which number gives the shortest total crawl time (meaning the highest performance). Since web crawling is an IO intensive job you will probably end up with a few threads per CPU / core in your computer. One thread per CPU would be too little, since it would be a lot of the time while waiting for data to download.

Parallel Workers Disadvantages

The parallel worker concurrency model has some disadvantages lurking under the simple surface. I will explain the most obvious disadvantages in the following sections.

Shared State Can Get Complex

In reality the parallel worker concurrency model is a bit more complex than illustrated above. The workers often need access to some kind of shared data, either in memory or in a shared database. The following diagram shows how this complicates the parallel worker concurrency model:



Some of this shared state is in communication mechanisms like job queues. But some of this shared state is business data, data caches, connection pools to the database etc.

As soon as shared state sneaks into the parallel worker concurrency model it starts getting complicated. The threads need to access the shared data in a way that makes sure that changes by one thread are visible to the others (pushed to main memory and not just stuck in the CPU cache of the CPU executing the thread). Threads need to avoid [race conditions](#), [deadlock](#) and many other shared state concurrency problems.

Additionally, part of the parallelization is lost when threads are waiting for each other when accessing shared data structures. Many concurrent data structures are blocking, meaning one or a limited set of threads can access them at any given time. This may lead to contention on these shared data structures. High contention will essentially lead to a degree of serialization of execution of the part of the code that accesses the shared data structures.

Modern [non-blocking concurrency algorithms](#) may decrease contention and increase performance, but non-blocking algorithms are hard to implement.

Persistent data structures are another alternative. A persistent data structure always preserves the previous version of itself when modified. Thus, if multiple threads point to the same persistent data structure and one thread modifies it, the modifying thread gets a reference to the new structure. All other threads keep a reference to the old structure which is still unchanged and thus consistent. The Scala programming language contains several persistent data structures.

While persistent data structures are an elegant solution to concurrent modification of shared data structures, persistent data structures tend not to perform that well.

For instance, a persistent list will add all new elements to the head of the list, and return a reference to the newly added element (which then points to the rest of the list). All other threads still keep a reference to the previously first element in the list, and to these threads the list appears unchanged. They cannot see the newly added element.

over the computer's memory. Modern CPUs are much faster at accessing data sequentially, so on hardware you will get a lot higher performance out of a list implemented on top of an array. An array stores data sequentially. The CPU caches can load bigger chunks of the array into the cache at a time and have the CPU access the data directly in the CPU cache once loaded. This is not really possible with a linked list where elements are scattered all over the RAM.

Stateless Workers

Shared state can be modified by other threads in the system. Therefore workers must re-read the state every time it needs it, to make sure it is working on the latest copy. This is true no matter whether shared state is kept in memory or in an external database. A worker that does not keep state internally (but re-reads it every time it is needed) is called *stateless*.

Re-reading data every time you need it can get slow. Especially if the state is stored in an external database.


Job Ordering is Nondeterministic

Another disadvantage of the parallel worker model is that the job execution order is nondeterministic. There is no way to guarantee which jobs are executed first or last. Job A may be given to a worker, yet job B may be executed before job A.

The nondeterministic nature of the parallel worker model makes it hard to reason about the state of the system at any given point in time. It also makes it harder (if not impossible) to guarantee that one happens before another.

Assembly Line

The second concurrency model is what I call the *assembly line* concurrency model. I chose that name to fit with the "parallel worker" metaphor from earlier. Other developers use other names (e.g. reactive systems, or event driven systems) depending on the platform / community. Here is a diagram illustrating the assembly line concurrency model:




The workers are organized like workers at an assembly line in a factory. Each worker only performs a part of the full job. When that part is finished the worker forwards the job to the next worker.


Each worker is running in its own thread, and shares no state with other workers. This is also sometimes referred to as a *shared nothing* concurrency model.

Systems using the assembly line concurrency model are usually designed to use non-blocking IO. Non-blocking IO means that when a worker starts an IO operation (e.g. reading a file or data from a network connection) the worker does not wait for the IO call to finish. IO operations are slow, so waiting for IO operations to complete is a waste of CPU time. The CPU could be doing something else in the meanwhile. When the IO operation finishes, the result of the IO operation (e.g. data read or status written) is passed on to another worker.

With non-blocking IO, the IO operations determine the boundary between workers. A worker does as much as it can until it has to start an IO operation. Then it gives up control over the job. When the operation finishes, the next worker in the assembly line continues working on the job, until that too starts an IO operation etc.



In reality, the jobs may not flow along a single assembly line. Since most systems can perform multiple jobs at once, jobs flow from worker to worker depending on the job that needs to be done. In reality there could be multiple different virtual assembly lines going on at the same time. This is how job flow in an assembly line system might look in reality:



The diagram illustrates a delegator and its interaction with workers. On the left, an orange rounded rectangle labeled "Delegator" has three arrows pointing to three green rounded rectangles labeled "Worker". These workers are arranged in a vertical column. From each "Worker", an arrow points to the next "Worker" in the column. Additionally, there are arrows from the top and bottom "Worker" to a third "Worker" on the right, indicating a parallel or redundant path.

Reactive, Event Driven Systems

- **Vert.x**
- Akka
- Node.JS (JavaScript)

Actors vs. Channels

```

graph LR
    A1[Actor] --> A2[Actor]
    A1 --> A3[Actor]
    A2 --> A4[Actor]
    A3 --> A4
    A3 --> A5[Actor]
    A3 --> A6[Actor]
    A3 --> A7[Actor]
  
```

```

graph LR
    W1[Worker] --> C1[Channel 1]
    W1 --> C2[Channel 2]
    C1 --> W2[Worker]
    C2 --> W3[Worker]
    W2 --> C3[Channel 3]
    W2 --> C4[Channel 4]
    C3 --> W4[Worker]
    C4 --> W5[Worker]
    C4 --> W6[Worker]
  
```

Assembly Line Advantages

Stateful Workers

--	--	--	--	--	--	--	--	--

external storage systems. A stateful worker can therefore often be faster than a stateless worker.

Better Hardware Conformity


Singlethreaded code has the advantage that it often conforms better with how the underlying hardware works. First of all, you can usually create more optimized data structures and algorithms when you assume the code is executed in single threaded mode.

Second, singlethreaded stateful workers can cache data in memory as mentioned above. When data is cached in memory there is also a higher probability that this data is also cached in the CPU cache of the CPU executing the thread. This makes accessing cached data even faster.

I refer to it as *hardware conformity* when code is written in a way that naturally benefits from how the underlying hardware works. Some developers call this *mechanical sympathy*. I prefer the term hardware conformity because computers have very few mechanical parts, and the word "sympathy" in this context is used as a metaphor for "matching better" which I believe the word "conform" conveys reasonably well. Anyways, this is nitpicking. Use whatever term you prefer.

Job Ordering is Possible

It is possible to implement a concurrent system according to the assembly line concurrency mode, a way that guarantees job ordering. Job ordering makes it much easier to reason about the state of the system at any given point in time. Furthermore, you could write all incoming jobs to a log. This log can then be used to rebuild the state of the system from scratch in case any part of the system fails. Tasks are written to the log in a certain order, and this order becomes the guaranteed job order. Here is how such a design could look:



Implementing a guaranteed job order is not necessarily easy, but it is often possible. If you can, it simplifies tasks like backup, restoring data, replicating data etc. as this can all be done via the log.

Assembly Line Disadvantages

The main disadvantage of the assembly line concurrency model is that the execution of a job is spread out over multiple workers, and thus over multiple classes in your project. Thus it becomes harder to see exactly what code is being executed for a given job.

It may also be harder to write the code. Worker code is sometimes written as callback handlers. Code with many nested callback handlers may result in what some developers call *callback hell*. Callback hell simply means that it gets hard to track what the code is really doing across all the callbacks, and as making sure that each callback has access to the data it needs.

With the parallel worker concurrency model this tends to be easier. You can open the worker code and read the code executed pretty much from start to finish. Of course parallel worker code may also be spread over many different classes, but the execution sequence is often easier to read from the code.

Functional Parallelism

Functional parallelism is a third concurrency model which is being talked about a lot these days (2012).

The basic idea of functional parallelism is that you implement your program using function calls. Functions can be seen as "agents" or "actors" that send messages to each other, just like in the assembly line concurrency model (AKA reactive or event driven systems). When one function calls another, that is similar to sending a message.

All parameters passed to the function are copied, so no entity outside the receiving function can manipulate the data. This copying is essential to avoiding race conditions on the shared data. This makes the function execution similar to an atomic operation. Each function call can be executed independently of any other function call.

When each function call can be executed independently, each function call can be executed on separate CPUs. That means, that an algorithm implemented functionally can be executed in parallel, on multiple CPUs.

With Java 7 we got the `java.util.concurrent` package which contains the `ForkAndJoinPool` which can help you implement something similar to functional parallelism. With Java 8 we got parallel `streams` which help you parallelize the iteration of large collections. Keep in mind that there are developers who are critical of the `ForkAndJoinPool` (you can find a link to criticism in my `ForkAndJoinPool` tutorial).

The hard part about functional parallelism is knowing which function calls to parallelize. Coordinating function calls across CPUs comes with an overhead. The unit of work completed by a function needs to be of a certain size to be worth this overhead. If the function calls are very small, attempting to parallelize them may actually be slower than a singlethreaded, single CPU execution.

From my understanding (which is not perfect at all) you can implement an algorithm using an reactive

parallelize (in my opinion).

Additionally, splitting a task over multiple CPUs with the overhead the coordination of that incurs, makes sense if that task is currently the only task being executed by the the program. However, if system is concurrently executing multiple other tasks (like e.g. web servers, database servers and other systems do), there is no point in trying to parallelize a single task. The other CPUs in the core are anyways going to be busy working on other tasks, so there is not reason to try to disturb them slower, functionally parallel task. You are most likely better off with an assembly line (reactive) concurrency model, because it has less overhead (executes sequentially in singlethreaded mode), conforms better with how the underlying hardware works.

Which Concurrency Model is Best?

So, which concurrency model is better?

As is often the case, the answer is that it depends on what your system is supposed to do. If your system is naturally parallel, independent and with no shared state necessary, you might be able to implement a system using the parallel worker model.

Many jobs are not naturally parallel and independent though. For these kinds of systems I believe the assembly line concurrency model has more advantages than disadvantages, and more advantages than the parallel worker model.

You don't even have to code all that assembly line infrastructure yourself. Modern platforms like [Vert.x](#) has implemented a lot of that for you. Personally I will be exploring designs running on top of platforms like Vert.x for my next projects. Java EE just doesn't have the edge anymore, I feel.

Next: [Same-threading](#)



Tweet

Jakob Jenkov

